

# Computer Vision

## Assignment 4

Sebastian Bergner

July 7, 2023

### Task

Construct a basic object recognition system (bag of words) and compare to a simple Convolutional Neural Network (CNN).

1. **GOAL:** Develop a system that can detect leafs or aircraft signals (see below) present in images. The object recognition system will be given an image and its output is to provide if contains the object to be recognized.
  - a. The first system will follow the process explained in lecture 10 (slide 66), namely:




Figure 1: BoW Pipeline

- b. The second system will make use of Neural Network, namely a CNN.
2. The Features:

You will use these features:

    - a. Shi-Tomasi corner detector
    - b. HoG (Histogram of Gradients)

These are popular feature descriptors available online and in OpenCV.

The Classifier: Use a Support Vector Machine (SVM). Look for a library (one is given at the end of this document) and learn how to use it.

### 3. THE DATASET AND DATA PREPARATION:

- a. Choose one task from:
  - [Leaf dataset](#)
  - [Aircraft signal dataset](#)
- b. Read the images and assign their corresponding labels.
  - Consider that the images are at different sizes, you must transform all images to the same size.
  - Colored pictures may have to be transformed to black and white (bag of features).

- Usually images are not normalized regarding contrast, if needed this can be done easily with just one line of code such as (this is Matlab):  

```
img2 = uint8((double(img)-double(min(min(img))))*(255/
double(max(max(img))-min(min(img)))));
```

#### 4. BAG OF FEATURES. OBJECT REPRESENTATION: feature extraction using Shi- Tomasi/HoG.

- After the keypoint localization, you have to convert them into their feature vectors.
- Those feature vectors must be clustered into a histogram, for this the most used method is k-means (available in OpenCV: MiniBatchKMeans).

#### 5. LEARNING: Training a classifier. Download the code for a classifier. You will use an SVM (e.g. library libsvm for Python or any other one). For comparison you will adapt the code of CNN explained in class, but replace VGG for Resnet-18. The CNN must also be adapted to work with your dataset. For the CNN you do not need to create a histogram, just feed the images to the CNN and record the accuracy and confusion matrix. Remember to use data augmentation.

#### 6. RECOGNITION: Using classifier in test data. We have now a trained system that can distinguish an image (leaf type, aircraft sign).

#### 7. RESULTS

At this point we have a classification system that can recognize what is on the picture. We have to measure how good is our system, for this we compare the true labels of the test set with the ones obtained in step 5. The higher the percentage the better, if you have 2 classes and this percentage is below 50% it means our system is ever worse than using chance (a coin with tails and heads), if it is a bit over 50% it means it is ok, 70-80% good enough and over 90% quite good, if over 95% it is great. Other cases: if you have 10 classes, you should get over 10%.

Create first a confusion matrix, from this, you can build a table reporting true positives, false positives, true negatives and false negatives.

Compare the three features for the two classifiers (Bag of words with SVM and CNN): Shi-tomasi, HoG to see which one performs better at the task. This should be reported in graphs, tables or ROC curves

## Preliminary

This document and the code has been done on my own. I shared the code and the images for the BoW/SVM part with Ori Sofer, therefore there will be similarities.

## Introduction on object recognition and steps to it

A very big and also important part of computer vision is object recognition. Object recognition can also be sectioned into object recognition, where we want to identify the class of an object. And into instance recognition where we want to identify one instance of a class, for example a human face should be used for unlocking a computer, of course we don't want anyone to be able to unlock the computer. In this assignment general object recognition is pursued. There is also a significant difference between recognizing one object class (or instance) or multiple objects. There are several approaches to object recognition. The simplest being a Bag-of-Words approach. This technique originates in the area of natural language processing. The steps for BoW:

1. Detect key points.
2. Generate feature vectors of the detected key points
3. Using k-means clustering find clusters and count the number of occurrences (histogram computation).
4. With the help of a neural network or a support vector machine, try to discriminate the classes from each other.

This usually works pretty good and is fast and simple compared to other methods. One of these other methods is a parts-based representation approach.

This approach tries to cope with a problem that arises with the BoW approach. That is, location of the found features and their correspondence isn't taken into account which means that there could be parts (or rather features) which can be found all over the image but the resulting class has a completely different structure. With the parts-based approach, objects are build of parts that have a certain location in respect to each other. But this is not easy as the relational position cannot be fixed as for example faces differ from each person to another. Also partial occlusion has to be dealt with which is also a hard problem.

Another way to get to object recognition is using a convolutional neural network. This is a very sophisticated approach that can work very accurate. A CNN consists of convolutional layers as the name suggests. These layers apply a trainable filter to the images in a convolutional fashion, shown in equation 1. In this equation  $g$  is the convolved image,  $f$  is the input image and  $h$  is the filter.

$$g(x, y) = \sum_j \sum_k f(x - j, y - k) * h(j, k) \quad (1)$$

The result of the first layer is similar to applying a gabor filter, therefore is able to detect edges. With more convolutional layers the extractable information is much more complex for example structural information can be extracted. After the convolutional layer normaly a fully connected feed forward neural network is used to discriminate the classes. A softmax function is applied to the output to have exactly one class as output.

## Purpose/Hypothesis

The purpose of all of this is to evaluate for one the performance of two different Bag-of-Words approaches. First one using Shi-Tomasi edge detector to find key points and another one using a Histogram-of-Gradients approach. And another approach using a CNN. The results should then be evaluated and interpreted. I hypothesize that the CNN will be performing better than the SVM approaches as it is much more sophisticated.

# Approach and Implementation

## Data

I decided to use the leafs dataset provided further above in the instructions. The plant classes are divided into healthy and diseased, I used both for training and testing of the object classification approaches. There is a strong bias in the data for example to plant class Jamun (P5) as there are 624 images (in healthy and diseased) opposing to Bael (P4) which has only 118 images, where all of them are diseased. To prevent a bias also in the machine learning model data augmentation has been used. I generated so many images that there were 1000 per class. To have more variance of the leafs in the images I used at least one of the following transformations

- flip
- rotate by 0 to 360°(with 90°steps)
- darken or brighten
- changing the contrast
- increase or decrease the image saturation
- or scale by 10% and crop it

This worked out pretty good. Two classes Basil (P8) and the previously mentioned Bael (P4) class contained just healthy and diseased images respectively. For these classes I've decided to just generate 1000 images within the categories that were available.

## SVM/BoW

### General

All of the aforementioned steps have to be run through and are practically the same for both Shi-Tomasi and HoG. To use a zip as input and not having to unpack the entire zip I've used the zipfile library in python. The code has been reused to a great part from last semesters project from the Visual Computing PS. The steps can be roughly described as:

1. Load files contained in zip.
2. Encode the classes to integers from 0 to 11.
3. Load all the images into memory.
4. Now extract feature vectors from all images.
5. Calculate number of classes \* 10 clusters of the feature vector space.
6. Calculate distance of each feature vector of each image to closest cluster.
7. Split files into train and test data (80% for training).
8. Train SVM using the test data.
9. Evaluate the performance of the trained SVM.
10. Save the SVM parameters.

### Shi-Tomasi

As Shi-Tomasi is just a corner detection algorithm we have to use another algorithm that helps us generating a feature vector for each found key point. I've decided to use SIFT for this as I am already familiar with it. As SIFT uses instances of the class cv.KeyPoint instead of simple image coordinates, because it can hold other useful information like size and magnitude. Therefore I had to wrap it. All further steps are the same

as for HoG.

## **HoG**

HoG is pretty straight forward compared to Shi-Tomasi. At first we have to resize the images to a lower resolution as the memory demand of HoG is drastically greater than the other approach. After that we create a HoGDescriptor instance which we then apply to each image. All further steps are again the same as described in the General section above.

## **CNN**

For the CNN I've used large parts of the provided Jupyter Notebook. But before training was even possible I had to write a custom dataset class to access the data of the zip and get the correct class names. This worked in a similar way to the BoW approach dataloader. Then the train- and test-dataloaders are set up with again the splitted files (this is being randomized), the dataloaders were set up with a batch size of 10 images. Now the ResNet18 model is loaded, and we freeze all layers. I added a few layers to the fully connected part of the network. Also I used kaiming initialization to hopefully start from better values for training. After that the criterion and the the optimizer are defined. I used Adam with a learning rate of 0.001 and a weight decay rate of 0.0001. Then the neural network is trained for 20 epochs, evaluated and saved afterwards.

## Results

### BoW - Shi-Tomasi

For Shi-Tomasi I tried to change various parameters like maxCorners, qualityLevel, minDistance and blockSize also I used a gridsearch approach to optimize the SVM training.

The best results were achieved by using a high number of maxCorners  $\geq 800$  a low quality level  $\leq 0.00001$  as well as a low minDistance  $\leq 3$  between found corners and a block size of 15 as parameters for Shi-Tomasi.

The best SVM parameters were the kernel CHI2, C didn't seem to affect the result as much best were between 1,000 – 10,000. Also  $\gamma$  didn't seem to have a large impact therefore I used the default value of 1. The termination criteria has been set to 5,000,000 iterations which seemed to be large enough, but also required quite a long time for fitting the SVM to the data.

Also the input images have been resized to 1024x1024 as otherwise it uses too much memory. I've also applied a Gaussian filter of size 7x7 to remove some of that noise.

### Accuracy

The combined Shi-Tomasi+SIFT SVM achieved a respectable accuracy of 85.875%.

### Confusion Matrix




Figure 2: Confusion matrix Shi-Tomasi+SIFT.

It is visible from the matrix that most predictions are pretty accurate due to the highest values being on the diagonal. There are some outliers like for example class 10 being detected as class 6 or 7. I missed to replace the class numbers with the labels and just saw this as of writing (same for HoG).

### BoW - HoG

For HoG I wasn't able to find such good suiting parameters as for the Shi-Tomasi+SIFT SVM. Nevertheless the best combination for HoG was using a winSize of 128x128 a blockSize of 16x16 a blockStride of 4x4 (anything less wouldn't work as it required loads of main memory) cellSize 8x8 (also the same here, less would not be possible) and nbins of size 9.

The SVM parameters I used were the same as for the Shi-Tomasi implementation. I tried many different values and the variance wasn't as large as for the Shi-Tomasi. So different kernels were much closer to the best performance. To be precise the kernels CHI2, RBF and SIGMOID performed very similar.

### Accuracy

The HoG SVM achieved the worst performance of the bunch with an accuracy of 65.625%.

### Confusion Matrix




Figure 3: Confusion Matrix

Compared to the first confusion matrix [2](#) here are much more outliers with the maximum being 38 for the true class 6 and predicted 2. This is 18.4% of the entire class not counting all other outliers, that would sum up to over 32%. It is quite obvious that this classifier does not perform good.

## CNN

This CNN is build upon ResNet-18 which looks like the following:




Figure 4: ResNet-18 Source: [Researchgate](#)

To get ResNet working with our leaf dataset I replaced the default fully connected layers at the end and froze the convolutional layers as they were pretrained and hopefully would be able to detect all relevant information. The output of the convolutional layer part is a vector of length 512, therefore the fc part must have such an input. And as we have 12 classes the output must be of length 12 too. I decided to add a few layers, 6 to be precise, in between to cope with all the information better (structure can be seen below in the Appendix).

The CNN performed good after just 10 epochs with about 84% accuracy, but I decided to train for 20 epochs which yielded a very good 95% accuracy.



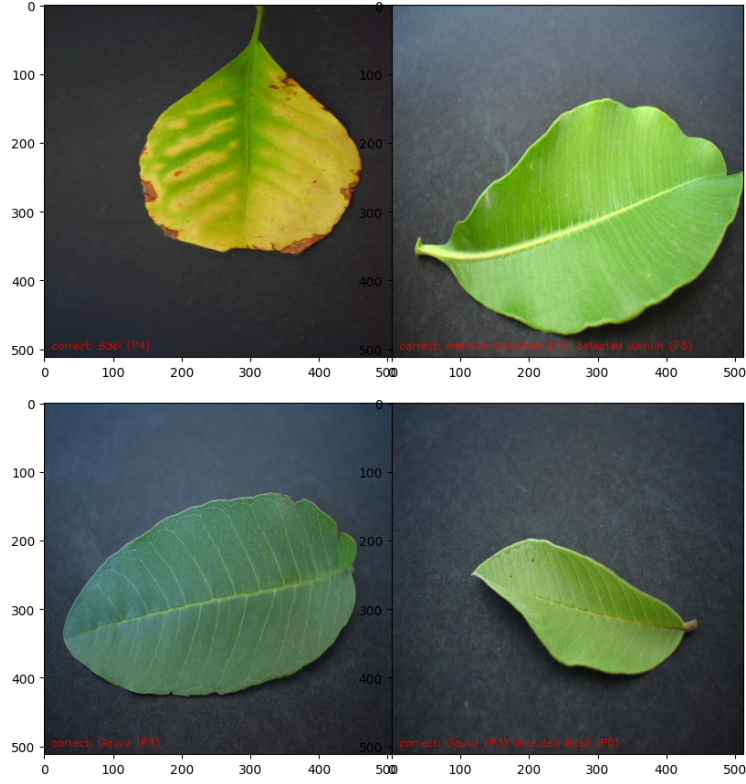


Figure 5: Example images of correct classified images and false ones on the right (with the wrong class and the correct one).

### Confusion Matrix

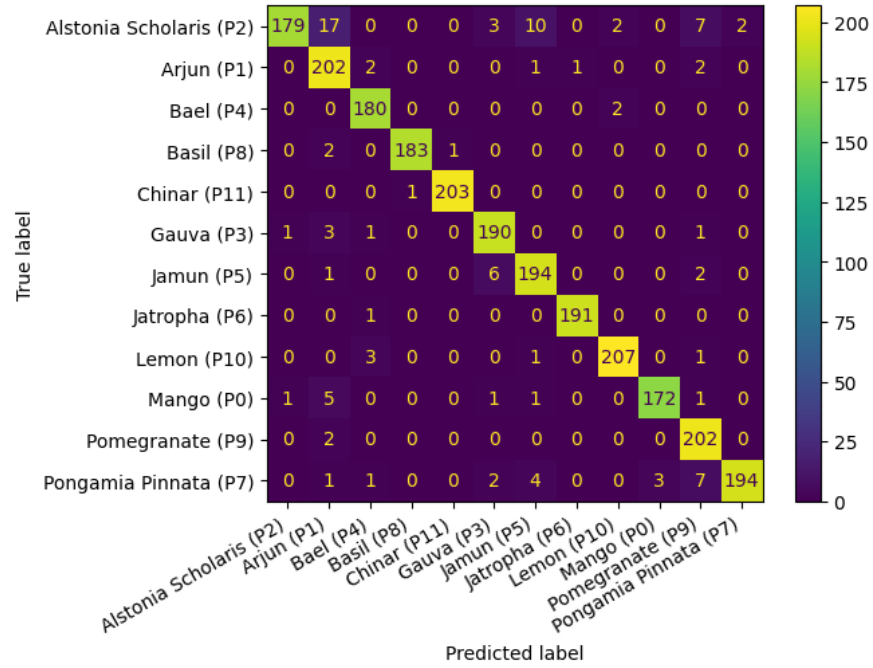


Figure 6: Confusion matrix of the CNN classification.

In the confusion matrix we can observe looking at the diagonal that most of the images have a very good classification. The worst performance is on the class *Alstonia Scholaris* (P2) which has been misclassified the most often (as *Arjun* (P1) and *Jamun* (P5)).

**ROC Curve** The ROC curve has 'only' been represented using 400 datapoints as it requires very much memory. I tried to use swap as I was only able to use 150 points at first before 32 GB of memory were full. 400 still used about 60 gigabytes of memory.

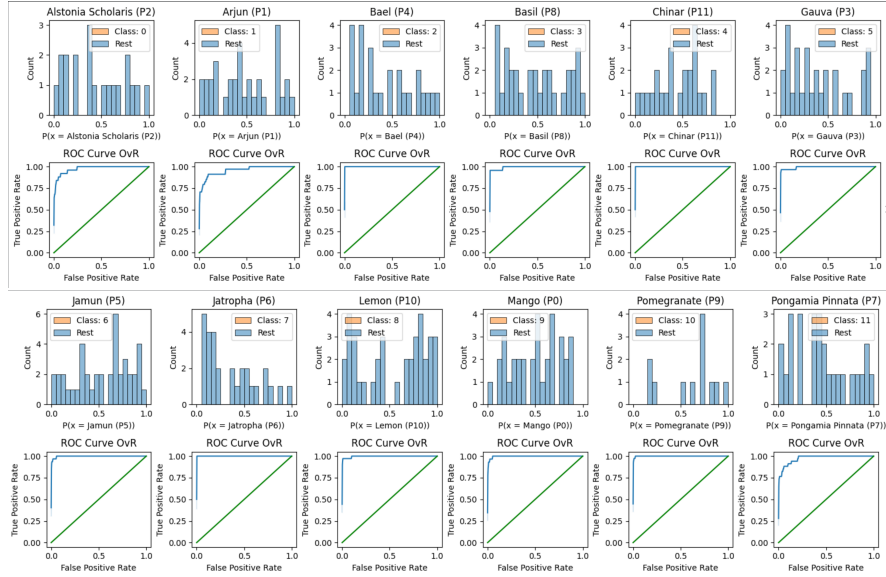


Figure 7: ROC curves for each class (One-vs-Rest).

From the ROC curves can be seen that for most of the classes the performance is good or even very good and for some classes, e.g. *Alstonia Scholaris* (P2) the area under the curve isn't as much. This data has to be interpreted with caution as just a part of the data has been used, therefore it isn't representative.

Class	TN	FP	FN	TP
<i>Alstonia Scholaris</i>	79326	69474	474	10726
<i>Arjun</i>	78863	69537	937	10663
<i>Bael</i>	78932	64268	868	15932
<i>Basil</i>	79180	67220	620	12980
<i>Chinar</i>	79365	68635	435	11565
<i>Gauva</i>	79164	69236	636	10964
<i>Jamun</i>	79007	64993	793	15207
<i>Jatropha</i>	79019	64981	781	15219
<i>Lemon</i>	79500	70500	300	9700
<i>Mango</i>	79003	65397	797	14803
<i>Pomegranate</i>	79302	67898	498	12302
<i>Pongamia Pinnata</i>	79048	68152	752	12048

## Discussion

What you learnt about this assignment and how the method can be improved or extended.

I learned a bit more about using OpenCV (especially at the augmentation part) and their implementation of LibSVM. Also I learned a bit more about using ResNet with PyTorch and some quirks with Torch as I used it on two systems my computer with a GPU and my laptop without a GPU and especially the juggling around with the memory location of objects.

I liked this assignment, there was just a bit of trouble with visualizing the ROC curves that's also a reason why I just did it for the CNN. Maybe a bit more time would be nice, as this entire exercise was joined by some exams.

## Appendix

The code has been handed in seperately as it is a Jupyter Notebook for one and all of the BoW is distributed over several files.

## CNN - Structure

Layer number	Layer (type)	Output Shape	Param #
1	Conv2d-1	[-1, 64, 16, 16]	9408
2	BatchNorm2d-2	[-1, 64, 16, 16]	128
3	ReLU-3	[-1, 64, 16, 16]	0
4	MaxPool2d-4	[-1, 64, 8, 8]	0
5	Conv2d-5	[-1, 64, 8, 8]	36864
6	BatchNorm2d-6	[-1, 64, 8, 8]	128
7	ReLU-7	[-1, 64, 8, 8]	0
8	Conv2d-8	[-1, 64, 8, 8]	36864
9	BatchNorm2d-9	[-1, 64, 8, 8]	128
10	ReLU-10	[-1, 64, 8, 8]	0
11	BasicBlock-11	[-1, 64, 8, 8]	0
12	Conv2d-12	[-1, 64, 8, 8]	36864
13	BatchNorm2d-13	[-1, 64, 8, 8]	128
14	ReLU-14	[-1, 64, 8, 8]	0
15	Conv2d-15	[-1, 64, 8, 8]	36864
16	BatchNorm2d-16	[-1, 64, 8, 8]	128
17	ReLU-17	[-1, 64, 8, 8]	0
18	BasicBlock-18	[-1, 64, 8, 8]	0
19	Conv2d-19	[-1, 128, 4, 4]	73728
20	BatchNorm2d-20	[-1, 128, 4, 4]	256
21	ReLU-21	[-1, 128, 4, 4]	0
22	Conv2d-22	[-1, 128, 4, 4]	147456
23	BatchNorm2d-23	[-1, 128, 4, 4]	256
24	Conv2d-24	[-1, 128, 4, 4]	8192
25	BatchNorm2d-25	[-1, 128, 4, 4]	256
26	ReLU-26	[-1, 128, 4, 4]	0
27	BasicBlock-27	[-1, 128, 4, 4]	0
28	Conv2d-28	[-1, 128, 4, 4]	147456
29	BatchNorm2d-29	[-1, 128, 4, 4]	256
30	ReLU-30	[-1, 128, 4, 4]	0
31	Conv2d-31	[-1, 128, 4, 4]	147456
32	BatchNorm2d-32	[-1, 128, 4, 4]	256
33	ReLU-33	[-1, 128, 4, 4]	0
34	BasicBlock-34	[-1, 128, 4, 4]	0
35	Conv2d-35	[-1, 256, 2, 2]	294912
36	BatchNorm2d-36	[-1, 256, 2, 2]	512
37	ReLU-37	[-1, 256, 2, 2]	0
38	Conv2d-38	[-1, 256, 2, 2]	589824
39	BatchNorm2d-39	[-1, 256, 2, 2]	512
40	Conv2d-40	[-1, 256, 2, 2]	32768
41	BatchNorm2d-41	[-1, 256, 2, 2]	512
42	ReLU-42	[-1, 256, 2, 2]	0
43	BasicBlock-43	[-1, 256, 2, 2]	0
44	Conv2d-44	[-1, 256, 2, 2]	589824
45	BatchNorm2d-45	[-1, 256, 2, 2]	512
46	ReLU-46	[-1, 256, 2, 2]	0
47	Conv2d-47	[-1, 256, 2, 2]	589824
48	BatchNorm2d-48	[-1, 256, 2, 2]	512
49	ReLU-49	[-1, 256, 2, 2]	0
50	BasicBlock-50	[-1, 256, 2, 2]	0

Layer number	Layer (type)	Output Shape	Param #
51	Conv2d-51	[-1, 512, 1, 1]	1179648
52	BatchNorm2d-52	[-1, 512, 1, 1]	1024
53	ReLU-53	[-1, 512, 1, 1]	0
54	Conv2d-54	[-1, 512, 1, 1]	2359296
55	BatchNorm2d-55	[-1, 512, 1, 1]	1024
56	Conv2d-56	[-1, 512, 1, 1]	131072
57	BatchNorm2d-57	[-1, 512, 1, 1]	1024
58	ReLU-58	[-1, 512, 1, 1]	0
59	BasicBlock-59	[-1, 512, 1, 1]	0
60	Conv2d-60	[-1, 512, 1, 1]	2359296
61	BatchNorm2d-61	[-1, 512, 1, 1]	1024
62	ReLU-62	[-1, 512, 1, 1]	0
63	Conv2d-63	[-1, 512, 1, 1]	2359296
64	BatchNorm2d-64	[-1, 512, 1, 1]	1024
65	ReLU-65	[-1, 512, 1, 1]	0
66	BasicBlock-66	[-1, 512, 1, 1]	0
67	AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
68	Linear-68	[-1, 400]	205200
69	LeakyReLU-69	[-1, 400]	0
70	Linear-70	[-1, 300]	120300
71	LeakyReLU-71	[-1, 300]	0
72	Linear-72	[-1, 200]	60200
73	LeakyReLU-73	[-1, 200]	0
74	Linear-74	[-1, 100]	20100
75	LeakyReLU-75	[-1, 100]	0
76	Linear-76	[-1, 50]	5050
77	LeakyReLU-77	[-1, 50]	0
78	Linear-78	[-1, 12]	612