



universität
innsbruck



Lecture 4. Feature detection and matching

703142. Computer Vision

Assoz.Prof. Antonio Rodríguez-Sánchez, PhD.

Outline

- Introduction
- Point and patches
- Edges
- Corners

Introduction

- What features to use for matching?



Introduction

- What features to use for matching?



Introduction



(a)



(b)



(c)



(d)

Introduction

- Feature detection and matching are essential in many computer vision applications
- We want to establish a set of *correspondences*
 - A 3D model can be constructed
 - An in-between view can be generated



Introduction

- Feature detection and matching are essential in many computer vision applications
- We want to establish a set of *correspondences*
 - A 3D model can be constructed
 - An in-between view can be generated



What features do we select?

Introduction

- Information content in images can be easily localized
 - Abrupt changes in image radiance: edges, corners, ...
 - Configurations of intensity changes corresponding to simple patterns: extended lines, circles
- Image **features** are local, meaningful and detectable parts of an image
 - Local implies limited spatial support 
 - Meaningful implies that they can be of use of subsequent operations 
 - Detectable implies that we can develop an algorithm for extracting the position and description of these entities given image data

Outline

- Introduction
- Point and patches
- Edges
- Corners

Points and patches

- *Keypoints* are used to perform object **instance**  and **category** recognition 
 - They permit matching even in the presence of
 - Occlusion
 - Large scale
 - Orientation changes
 - Its origin comes from stereo and are used presently for image-stitching applications

Points and patches

- Two main approaches to finding feature points and their correspondences
 - Find features in one image that can be accurately tracked using local search techniques
 - i.e. correlation
 - Least squares
 - Detect features and then match features based on their local appearance



Points and patches

4 stages:

Feature detectors

Feature descriptors

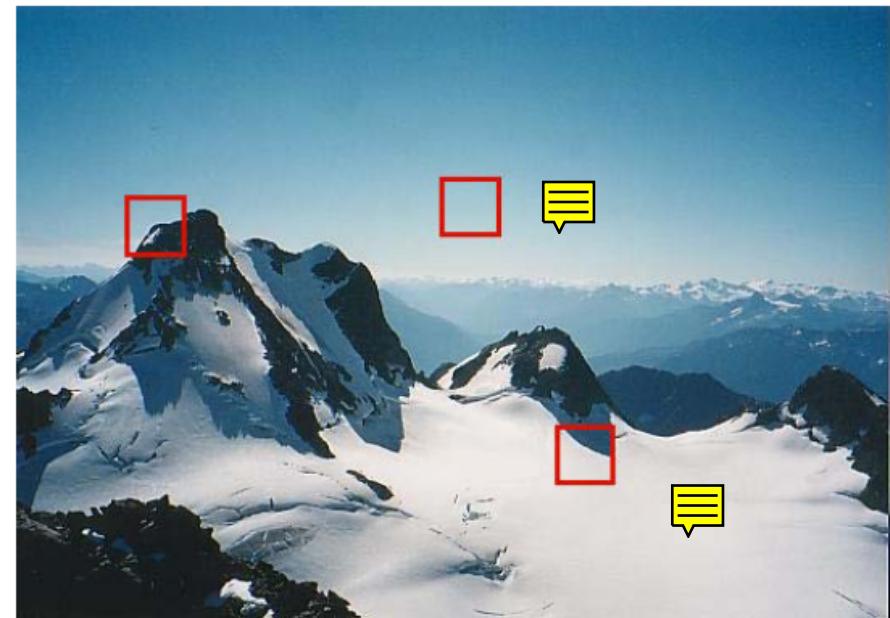
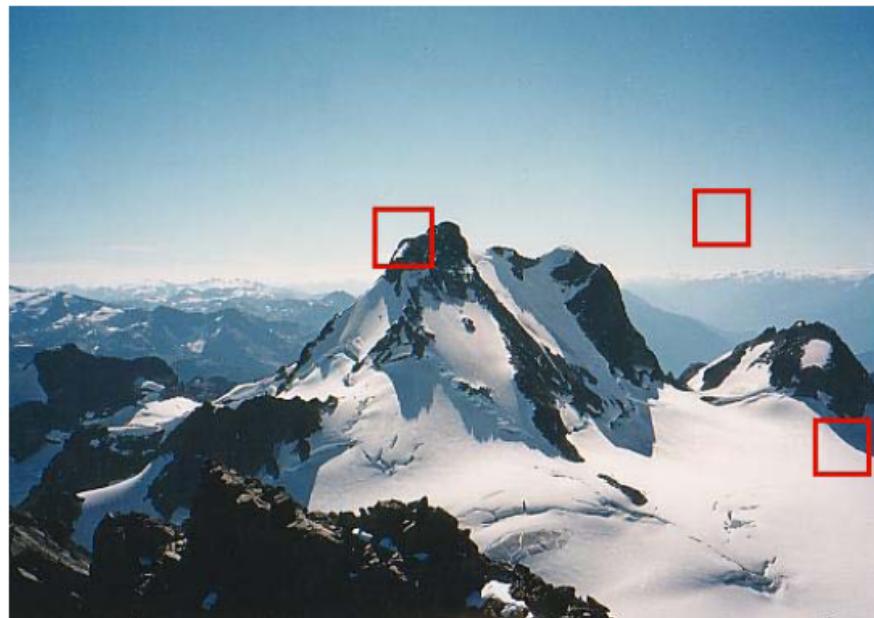
Feature matching

Feature tracking

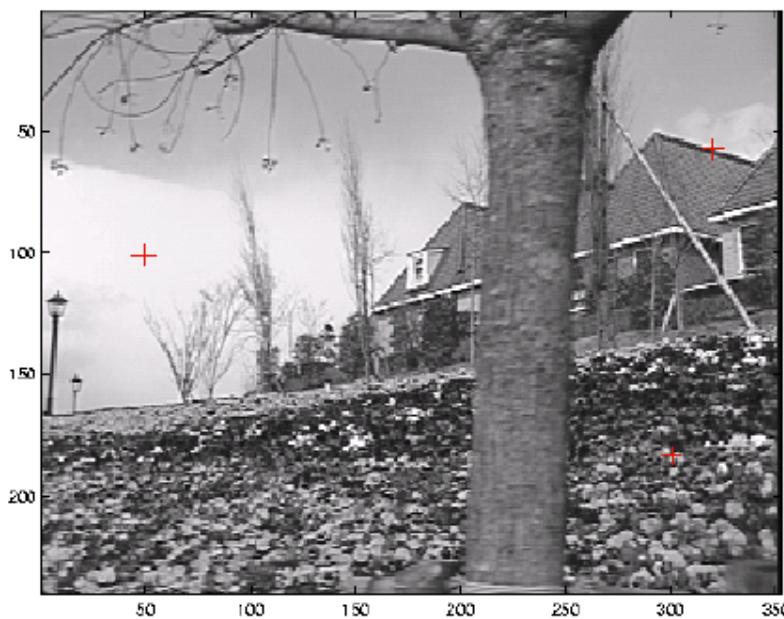


Feature detectors

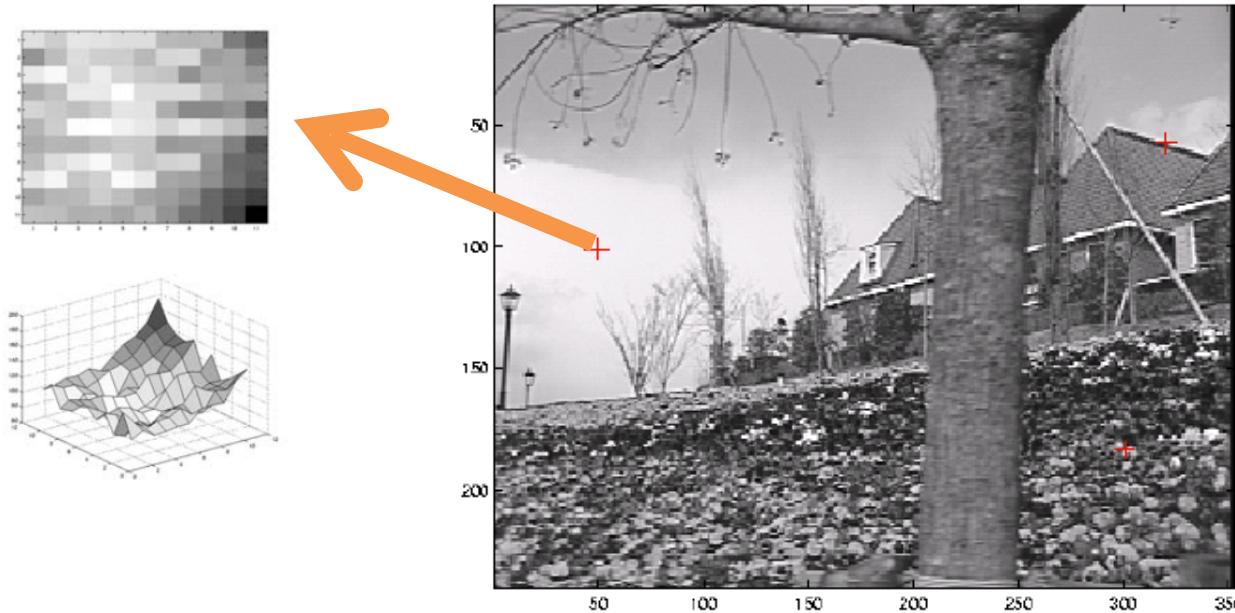
- What are the *good* features?



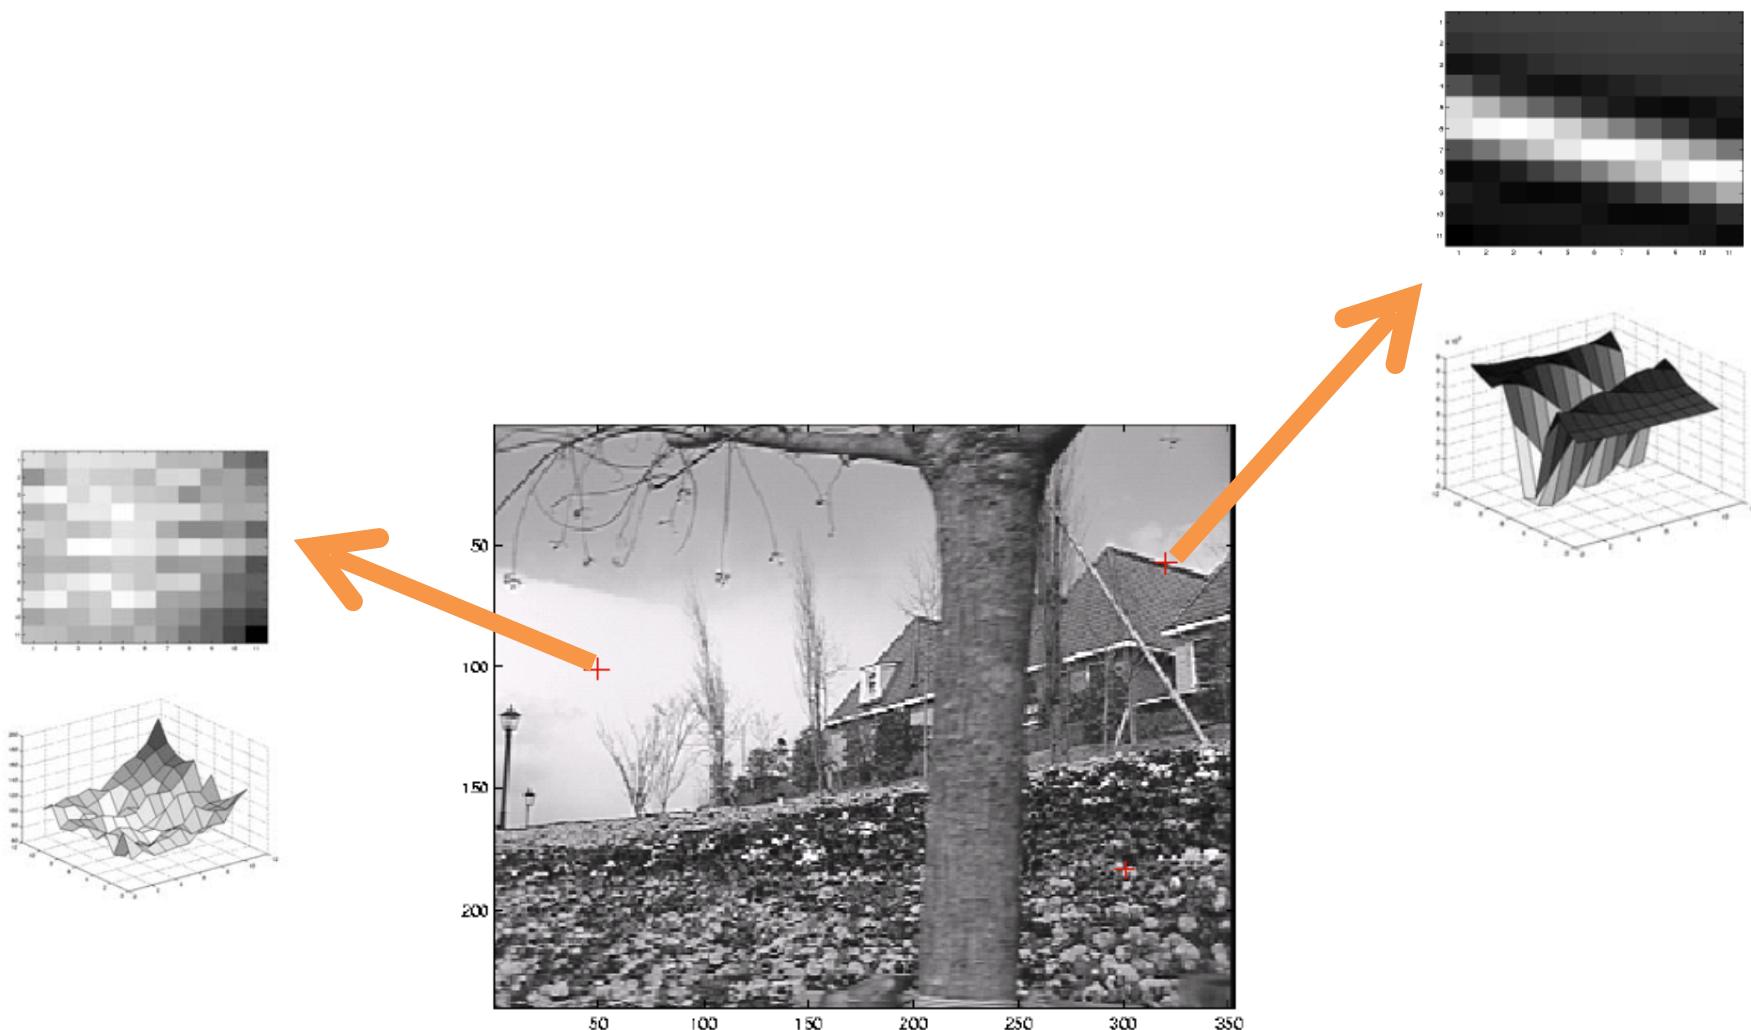
Feature descriptors



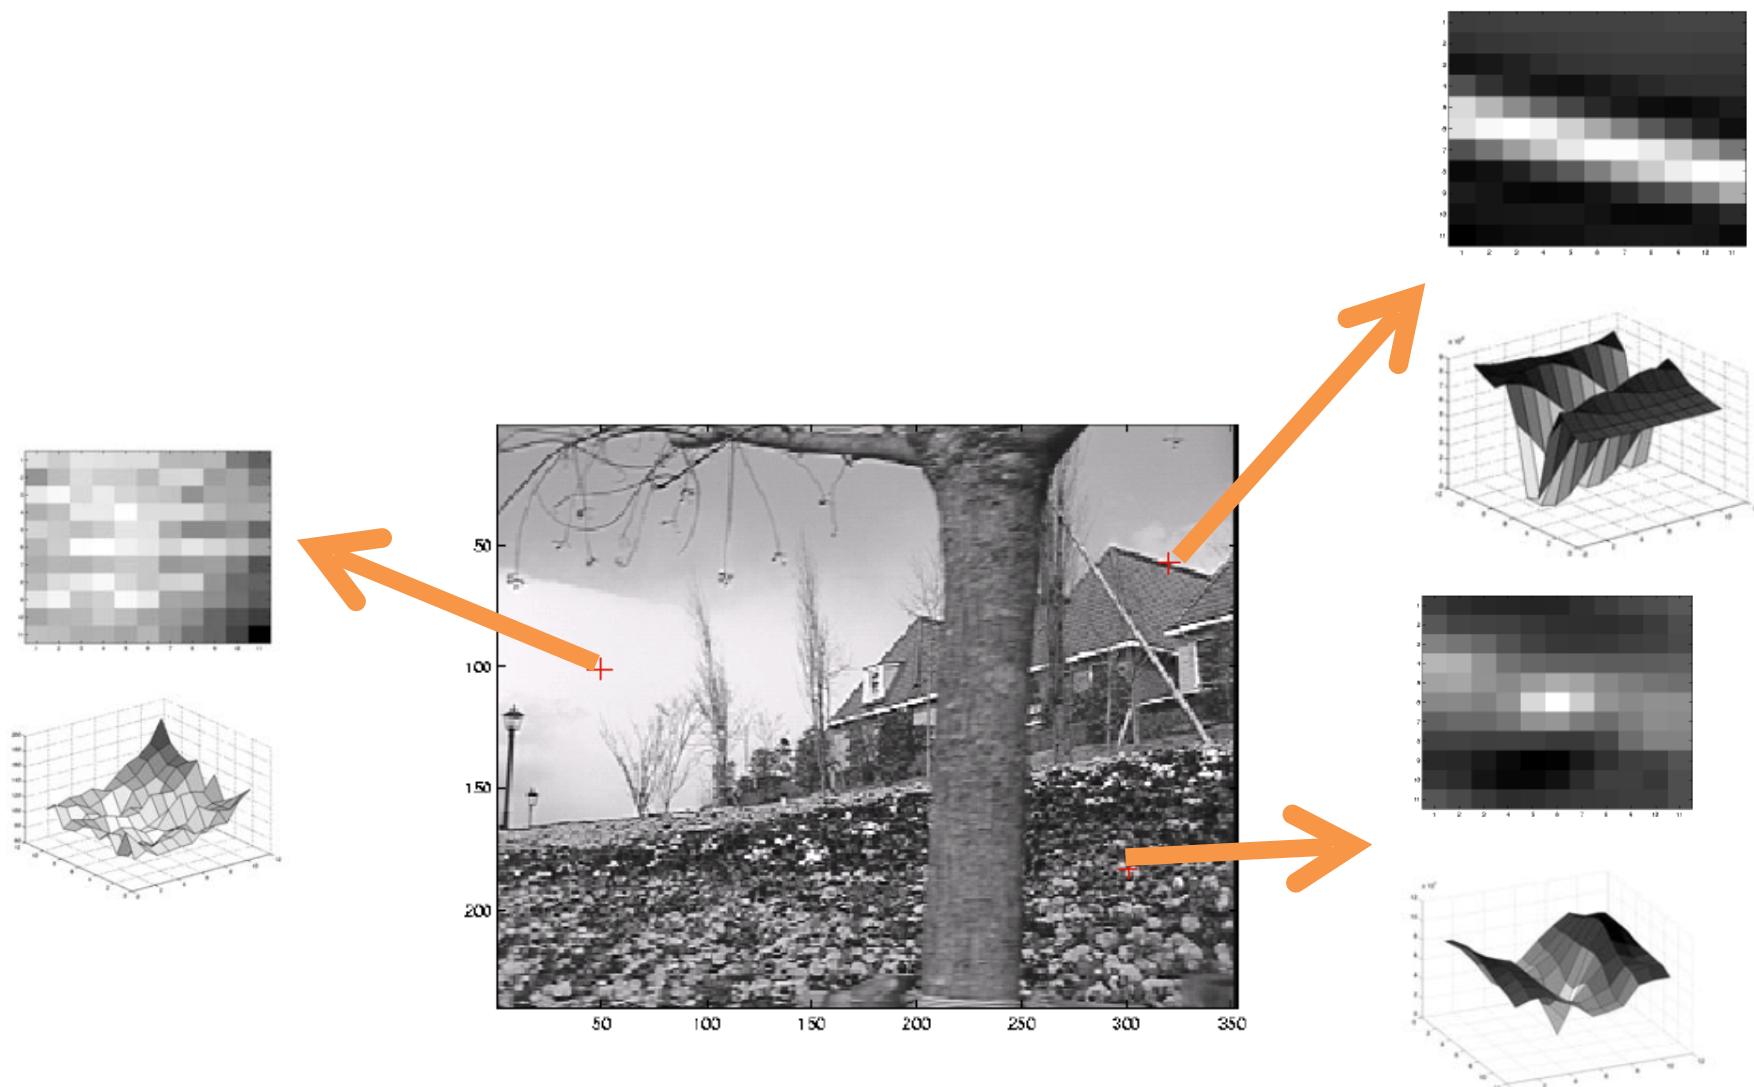
Feature descriptors



Feature descriptors



Feature descriptors

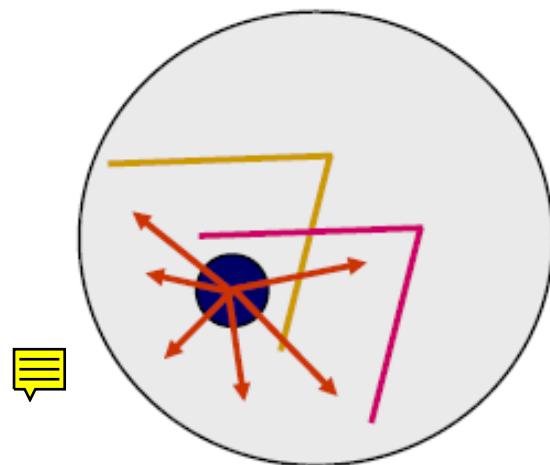


Feature detectors

- What are the *good* features?

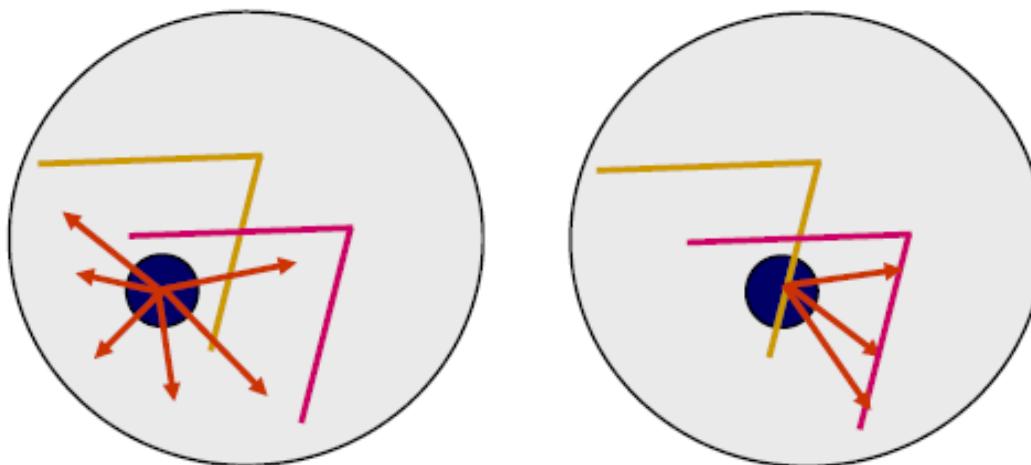
Feature detectors

- What are the *good* features?
 - Textureless patches are nearly impossible to localize



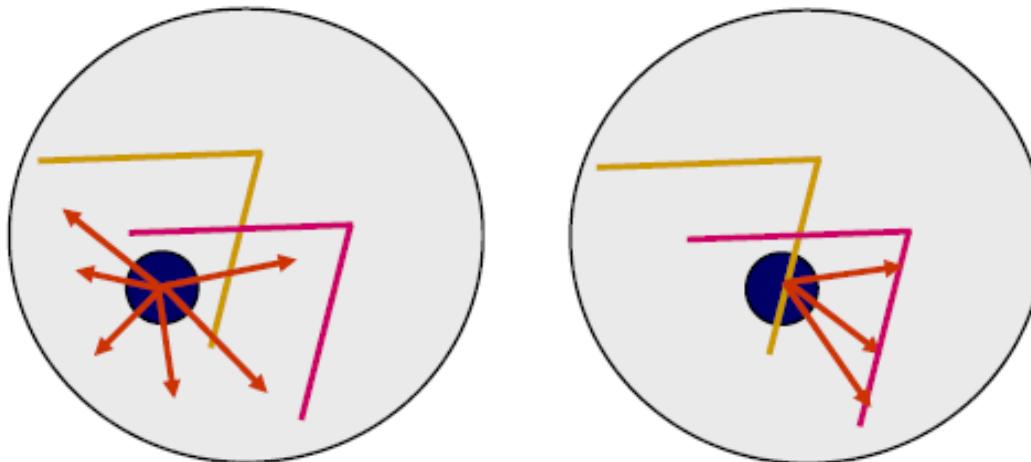
Feature detectors

- What are the *good* features?
 - Textureless patches are nearly impossible to localize
 - Patches with large contrast changes are easier to localize



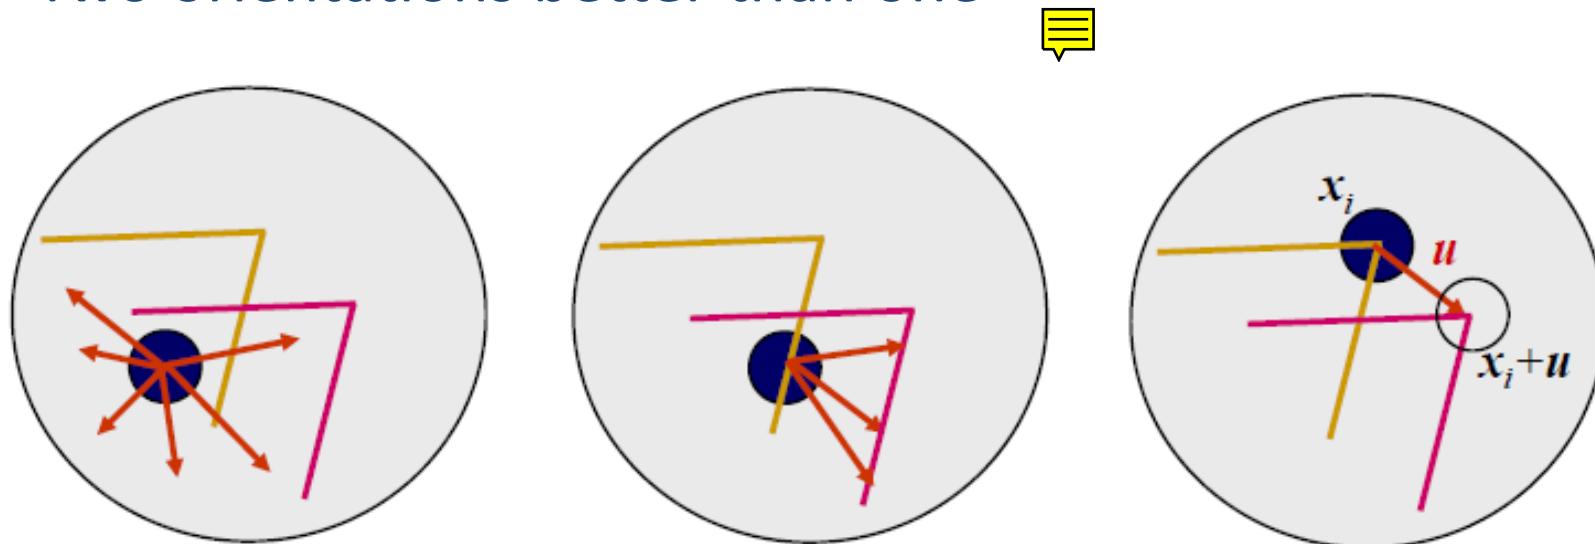
Feature detectors

- What are the *good* features?
 - Textureless patches are nearly impossible to localize
 - Patches with large contrast changes are easier to localize
 - But lines suffer of the *aperture problem*



Feature detectors

- What are the *good* features?
 - Textureless patches are nearly impossible to localize
 - Patches with large contrast changes are easier to localize
 - But lines suffer of the *aperture problem*
 - Two orientations better than one



Feature detectors

- Local appearance of features will change in
 - Orientation
 - Scale
 - Affine deformations
 - Intensity
 - ...

Feature detectors

- Interesting characteristics
 - Repeatability
 - Frequency with which keypoints are found
 - Scale invariance
 - Extract features at a variety of scales (i.e. pyramid) or
 - Extract features that are stable in scale
 - Rotation invariance and orientation estimation
 - Design descriptors rotationally invariant or
 - Estimate a dominant orientation at each keypoint
 - Once the local orientation and scale keypoint have been estimated, a scaled and oriented patch around the detected point can be extracted
 - We can use steerable filters
 - Affine invariance (deformations)



Feature detectors

- Scale invariance
 - 1. Extract features at a variety of scales



E.g. MOPS descriptors
extracted at five pyramid
levels

- 2. Extract features stable at location and scale (e.g. SIFT)

Feature detectors

- Rotation invariance and orientation estimation
 1. Design descriptors that are rotationally invariant
 2. Estimate a dominant orientation

Average gradient within a region around the keypoint

Histogram of orientation



E.g. MOPS descriptors
extracted at five pyramid
levels

Feature detectors

- Affine invariance
 - Fit an ellipse to the autocorrelation or Hessian matrix

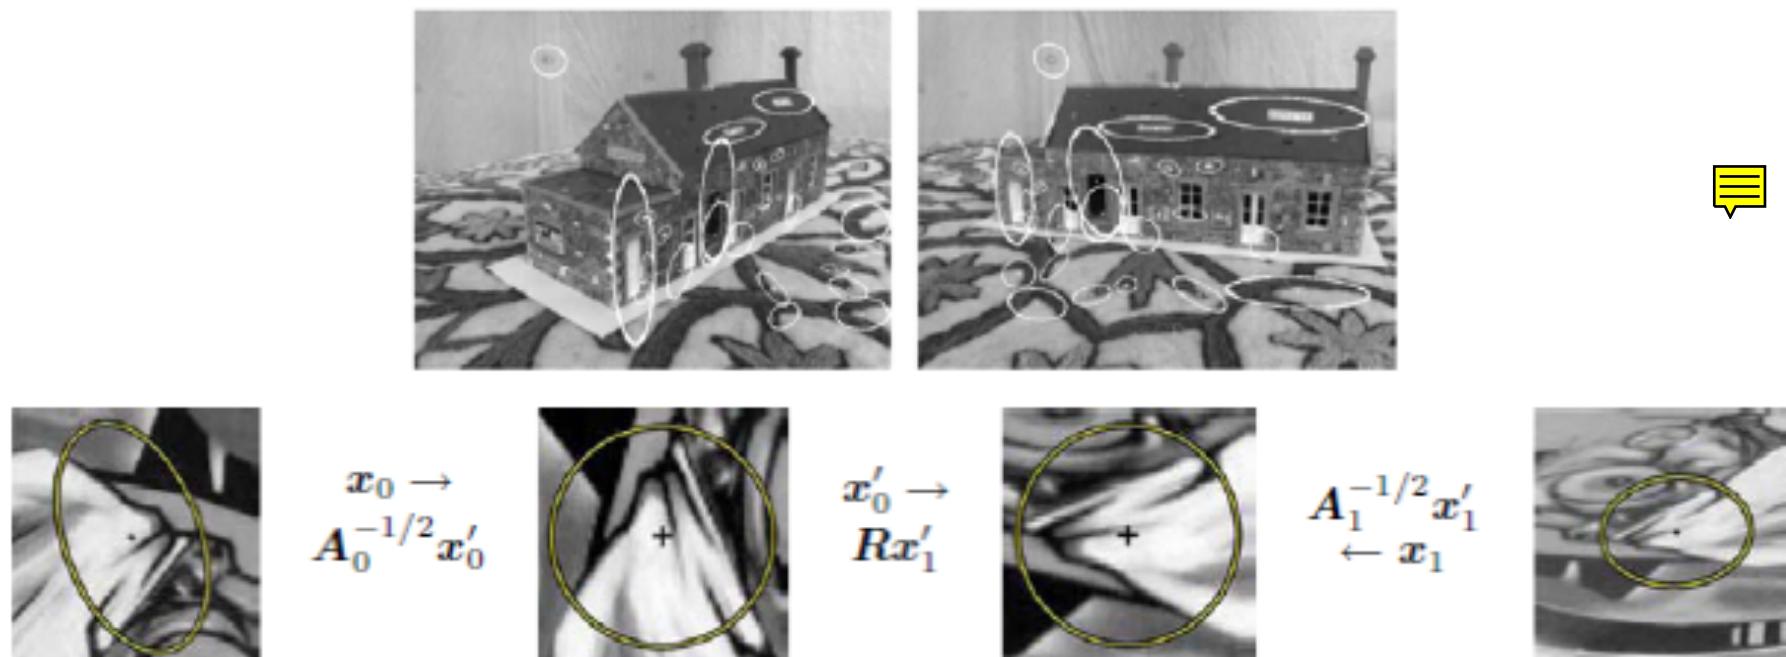


Figure 4.14 Affine normalization using the second moment matrices, as described by Mikolajczyk, Tuytelaars, Schmid *et al.* (2005) © 2005 Springer. After image coordinates are transformed using the matrices $A_0^{-1/2}$ and $A_1^{-1/2}$, they are related by a pure rotation R , which can be estimated using a dominant orientation technique.

Feature descriptors

- Affine invariance
 - Bias and gain normalization (MOPS)
 - Brown, Szeliski and Winder (2005)
 - Scale invariant feature transform (SIFT)
 - PCA-SIFT
 - Gradient location-orientation histogram (GLOH)
 - Mikolajczyk and Schmid (2005)
 - Steerable filters
 - Freeman and Adelson (1991)

Feature matching

- After detecting features, we must match them
 - We need to define a metric

Feature matching

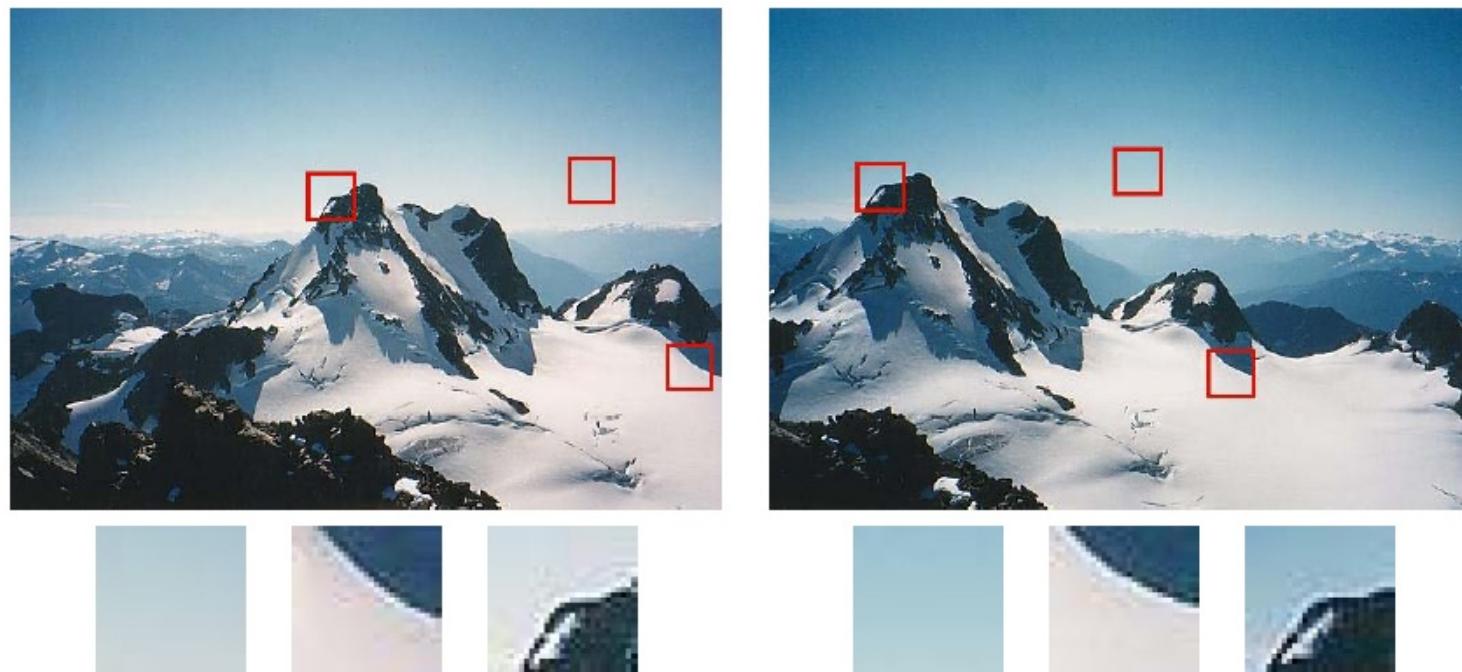
- Matching criterion
 - E.g. Weighted summed square difference

$$E_{\text{WSSD}}(\mathbf{u}) = \sum_i w(x_i) [I_1(x_i + \mathbf{u}) - I_0(x_i)]^2,$$

Feature matching

- Matching criterion
 - E.g. Weighted summed square difference

$$E_{\text{WSSD}}(\mathbf{u}) = \sum_i w(x_i) [I_1(x_i + \mathbf{u}) - I_0(x_i)]^2,$$



Feature matching

- Matching criterion
 - E.g. Weighted summed square difference

$$E_{\text{WSSD}}(u) = \sum_i w(x_i)[I_1(x_i + u) - I_0(x_i)]^2,$$

- But there are many other measures, sometimes depending on the task
- Indexing to make matching faster

Feature matching

- Matching criterion
 - We need also some measure of how “good” is our method for matching

Feature matching

- Matching criterion
 - We need also some measure of how “good” is our method for matching
 - **TP**: True positives, number of correct matches
 - **FN**: False negatives, matches not correctly detected
 - **FP**: False positives, proposed matches that are incorrect
 - **TN**: True negatives, not-matches that were correctly detected
 - We usually **create a *confusion matrix*.**

Feature matching

- Matching criterion
 - We usually create a *confusion matrix*.

	Cylinder	Circular	Box	Paper	Wrench
Cylinder	80%	0	0	20%	0
Circular	0	100%	0	0	0
Box	0	0	68.68%	31.43%	0
Paper	6.67%	0	16.67%	70%	6.67%
Wrench	0	0	6.67%	0	93.33%

Feature matching

- Matching criterion
 - We usually create a *confusion matrix*.

	Cylinder	Circular	Box	Paper	Wrench
Cylinder	80%	0	0	20%	0
Circular	0	100%	0	0	0
Box	0	0	68.68%	31.43%	0
Paper	6.67%	0	16.67%	70%	6.67%
Wrench	0	0	6.67%	0	93.33%

– TP=82.4%

Feature matching

- Matching criterion

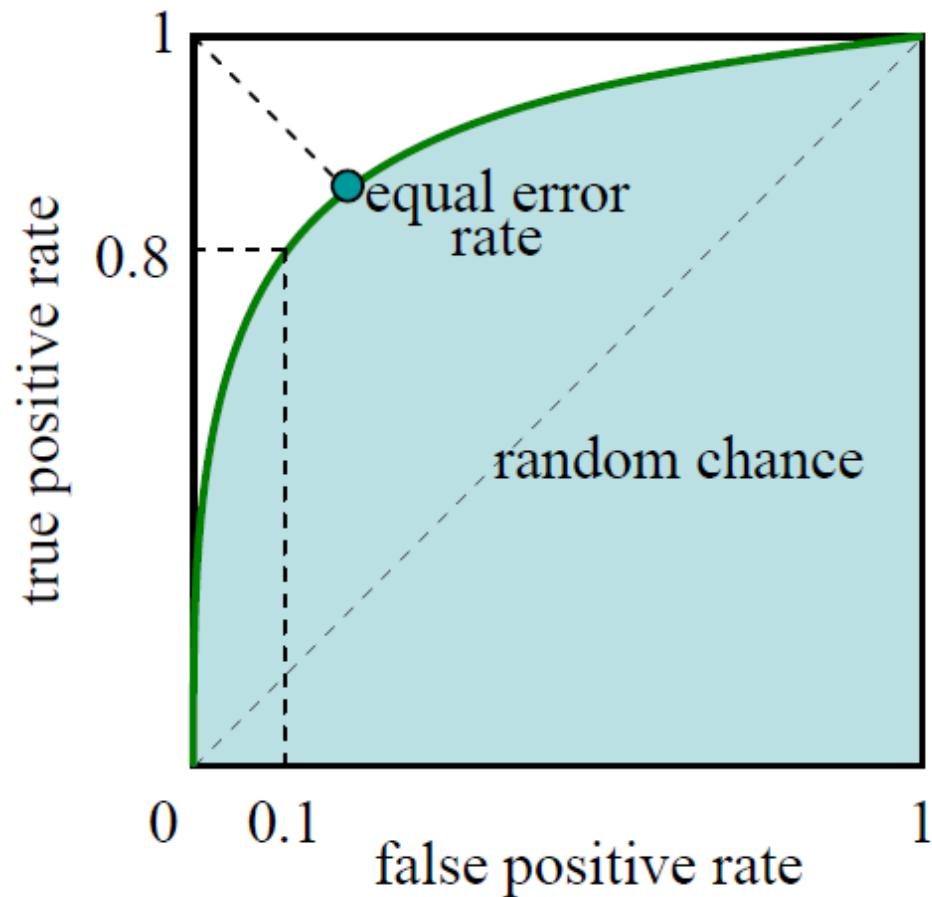
	True matches	True non-matches	
Predicted matches	TP = 18	FP = 4	P' = 22
Predicted non-matches	FN = 2	TN = 76	N' = 78
	P = 20	N = 80	Total = 100
TPR = 0.90		FPR = 0.05	

– True Positive Rate (TPR) $TPR = \frac{TP}{TP + FN}$

– False Positive Rate (FPR) $FPR = \frac{FP}{FP + TN}$

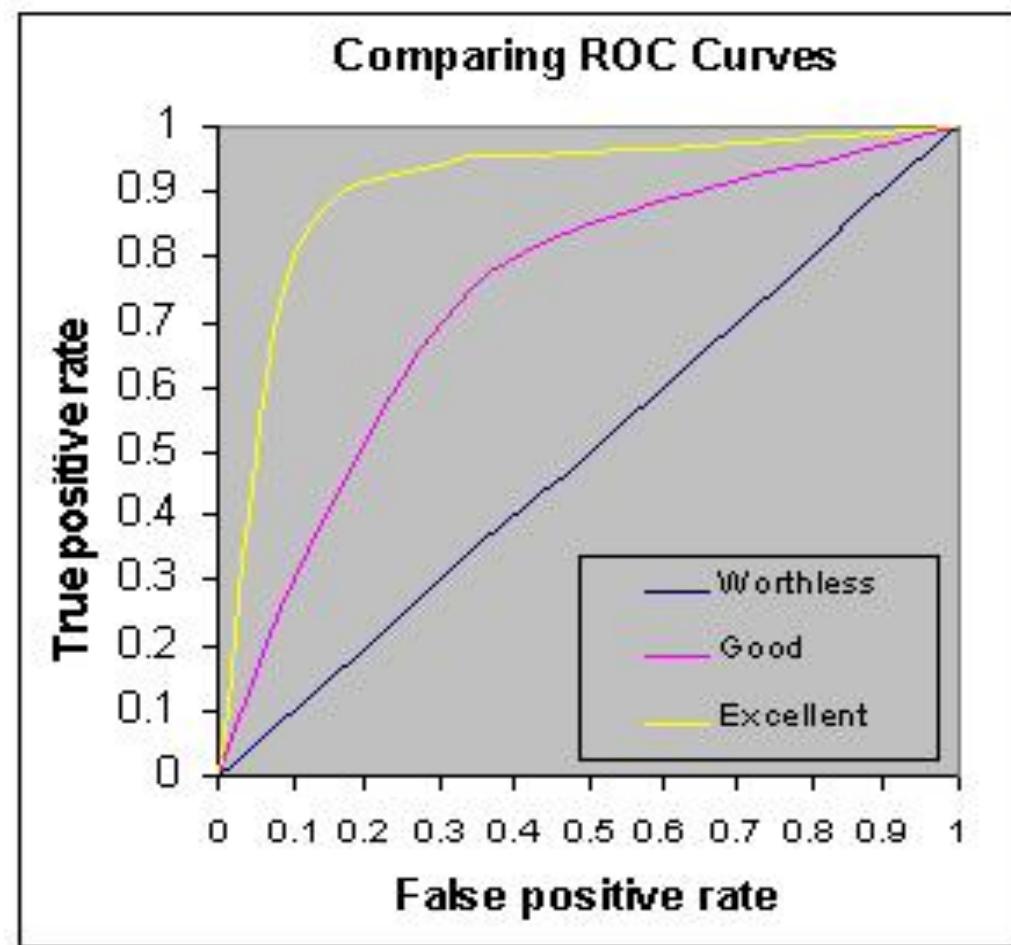
Feature matching

- Matching criterion
 - ROC curve



Feature matching

- Matching criterion
 - ROC curve



Feature matching

- Matching criterion

	True matches	True non-matches	
Predicted matches	TP = 18	FP = 4	P' = 22
Predicted non-matches	FN = 2	TN = 76	N' = 78
	P = 20	N = 80	Total = 100
TPR = 0.90		FPR = 0.05	

— TPR

$$TPR = \frac{TP}{TP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$



— FPR

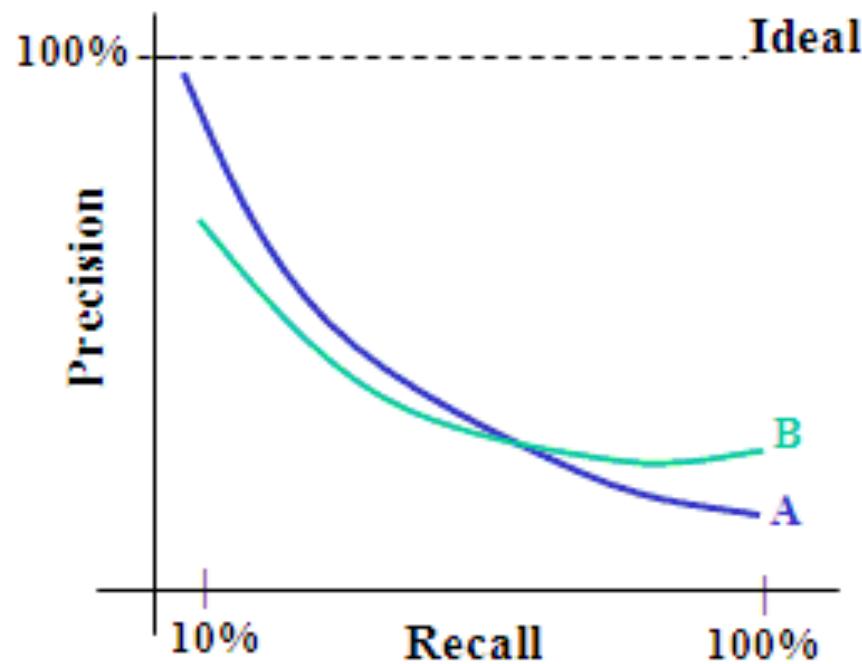
$$FPR = \frac{FP}{FP + TN}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$\text{Accuracy} = \frac{TP + TN}{TP + FN + FP + FN}$$

Feature matching

- Matching criterion



Feature tracking

- Feature tracking
 - Find a set of likely feature locations and then search for locations in subsequent images
 - Detect then track: used in video applications



Feature tracking

- Feature tracking
 - Find a set of likely feature locations and then search for locations in subsequent images
 - Detect then track: used in video applications
 - Example: Facial puppetry

Case study: SIFT

- David G. Lowe, "**Object recognition from local scale-invariant features,**" *International Conference on Computer Vision*, Corfu, Greece (September 1999), pp. 1150-1157
- David G. Lowe, "**Local feature view clustering for 3D object recognition,**" *IEEE Conference on Computer Vision and Pattern Recognition*, Kauai, Hawaii (December 2001), pp. 682-688
- David G. Lowe, "**Distinctive image features from scale-invariant keypoints,**" *International Journal of Computer Vision*, 60, 2 (2004), pp. 91-110
- <http://www.cs.ubc.ca/~lowe/keypoints/>

Case study: SIFT

1. Scale-space extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

Case study: SIFT

- Scale-space extrema detection
 - Identify locations and scales that can be repeatably assigned under differing views of the same objects
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k

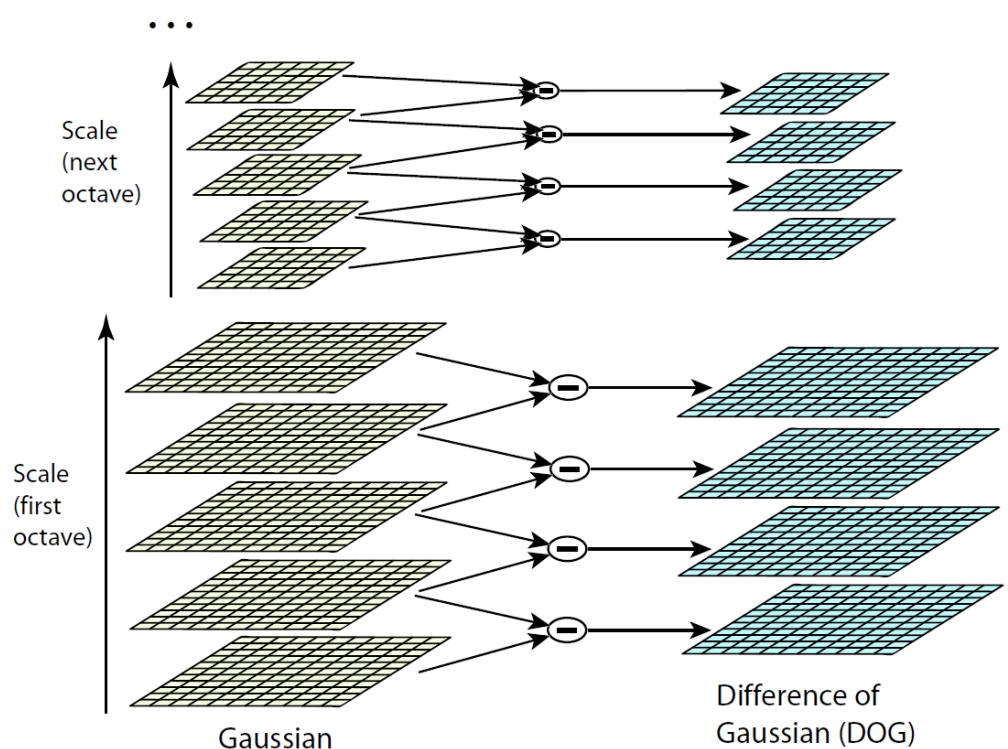
$$\begin{aligned} D(x, y, \sigma) &= (G(x, y, k\sigma) - G(x, y, \sigma)) * I(x, y) \\ &= L(x, y, k\sigma) - L(x, y, \sigma). \end{aligned}$$

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

Case study: SIFT

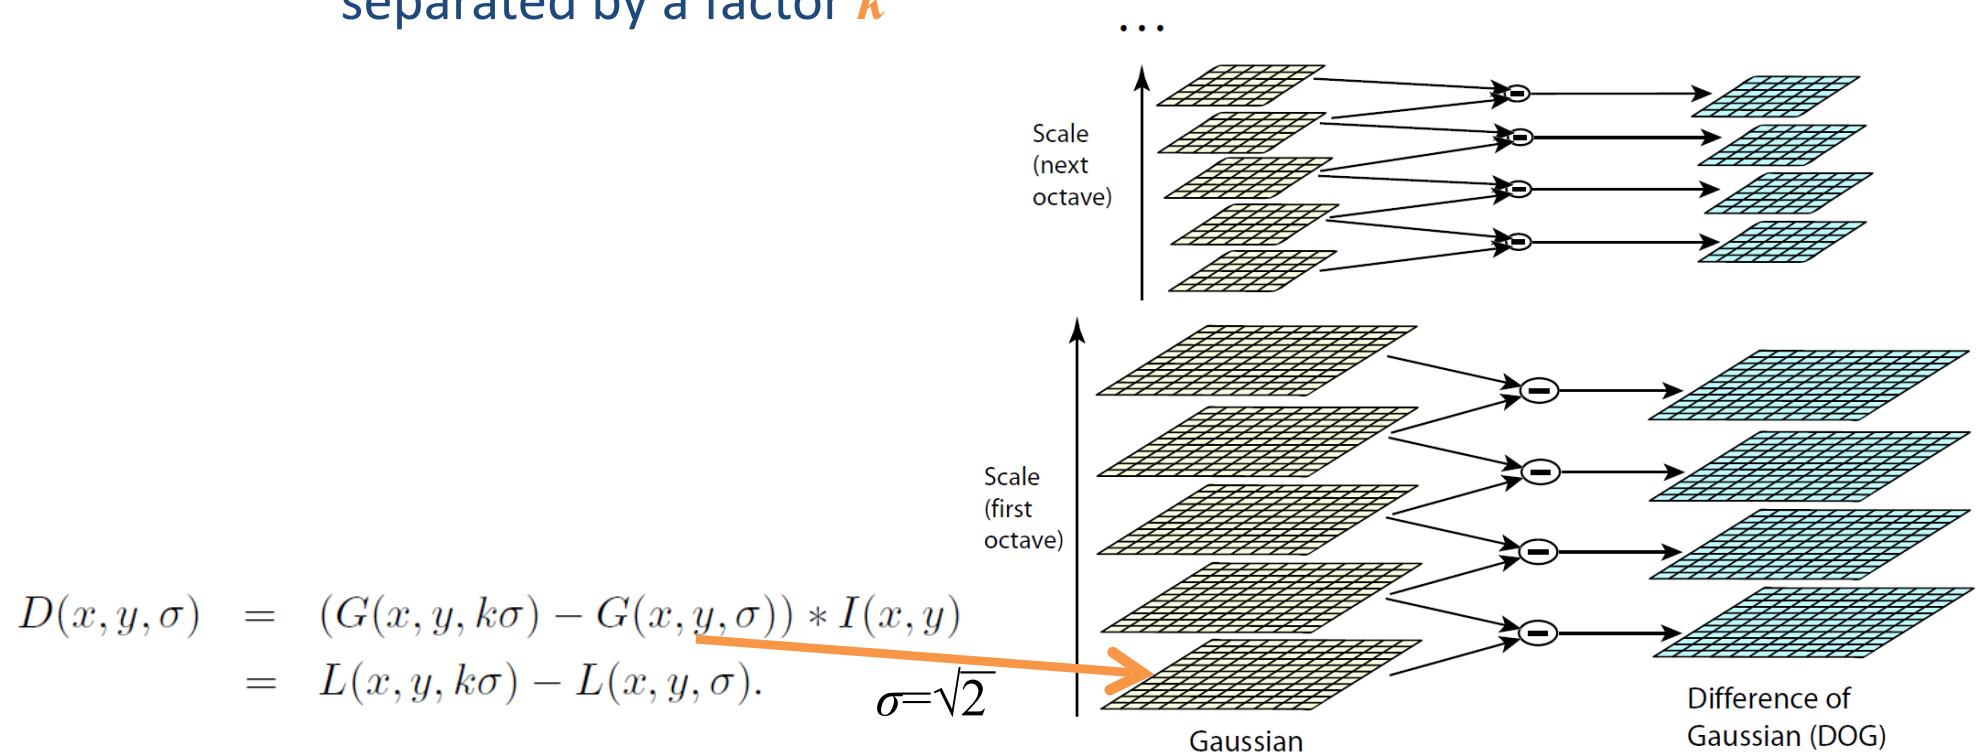
- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k

$$\begin{aligned} D(x, y, \sigma) &= (G(x, y, k\sigma) - G(x, y, \sigma)) * I(x, y) \\ &= L(x, y, k\sigma) - L(x, y, \sigma). \end{aligned}$$



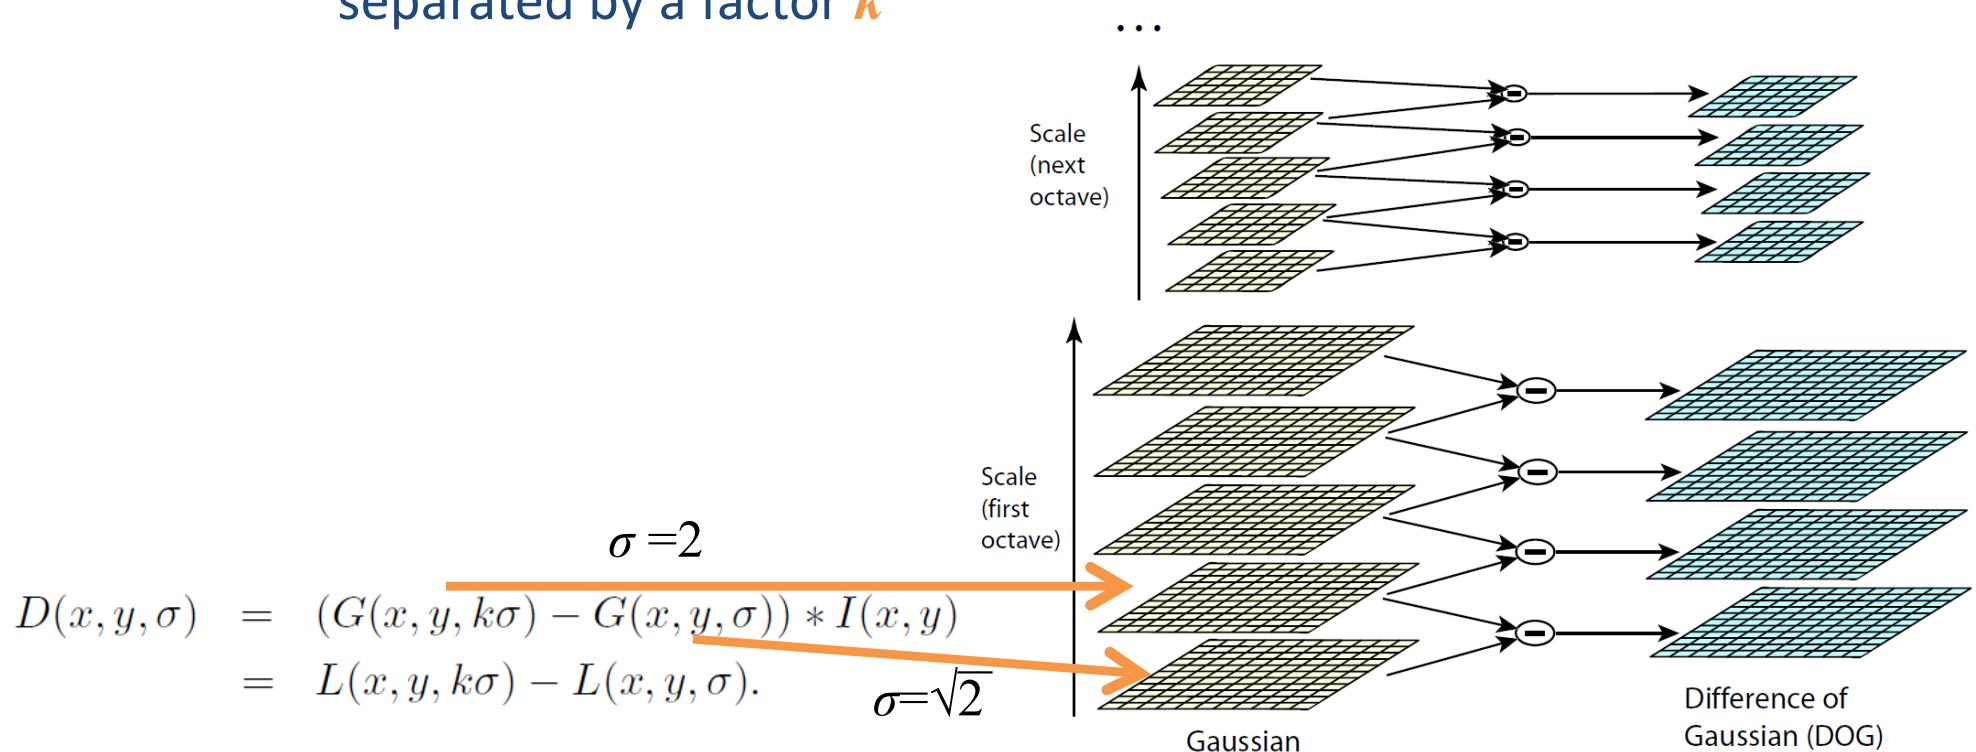
Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k



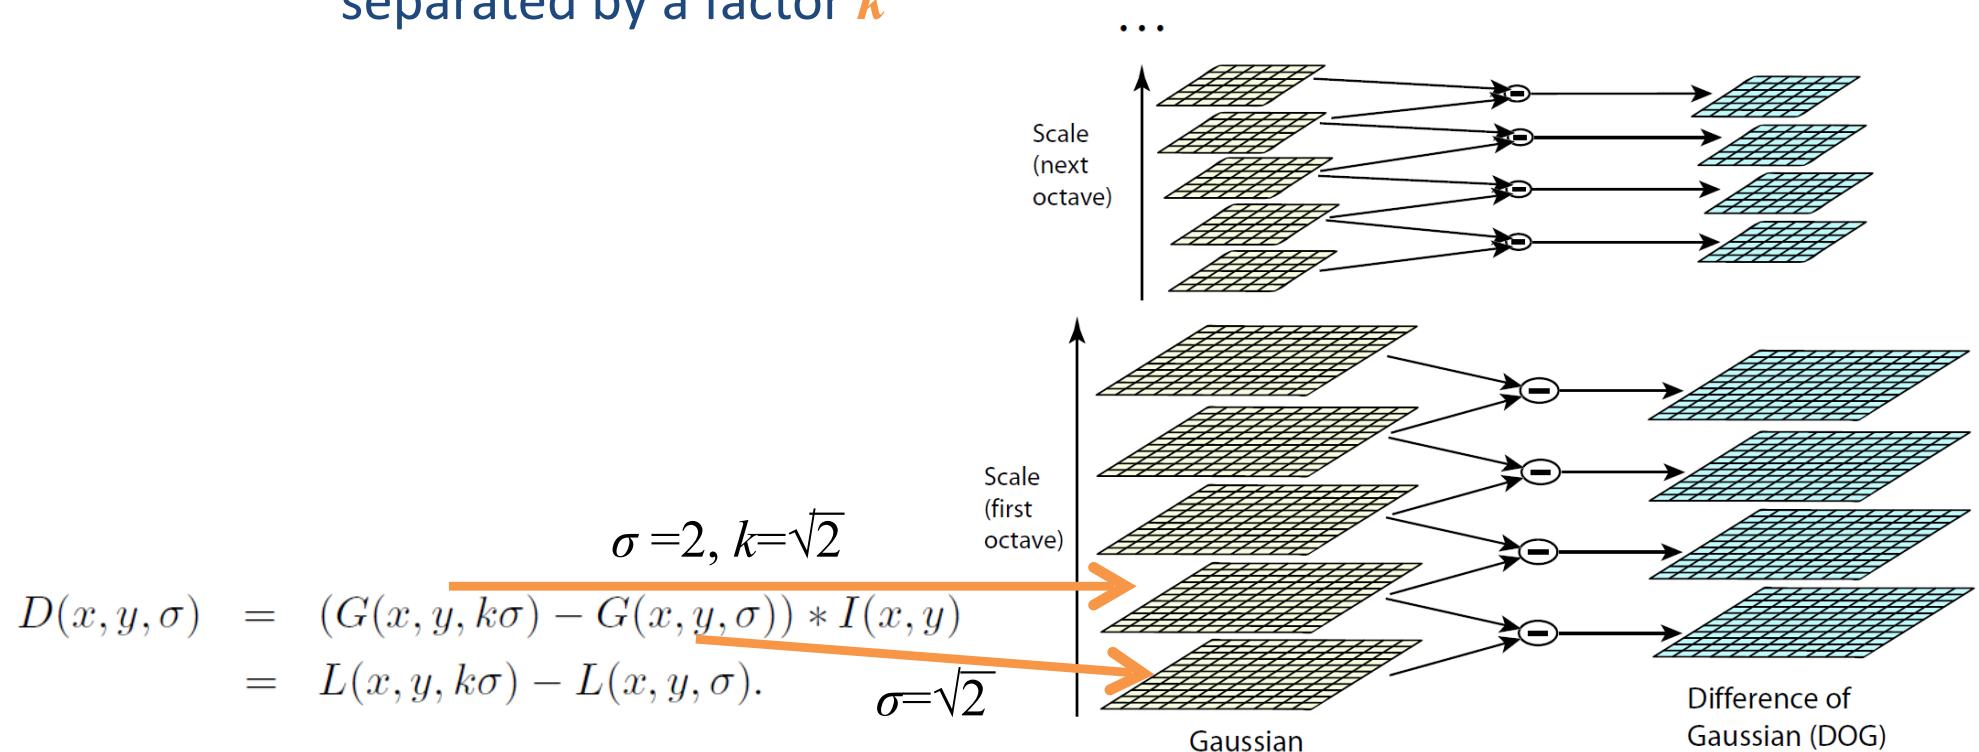
Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k



Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k



Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k
 - Note: This is similar to normalizing the Laplacian by σ^2

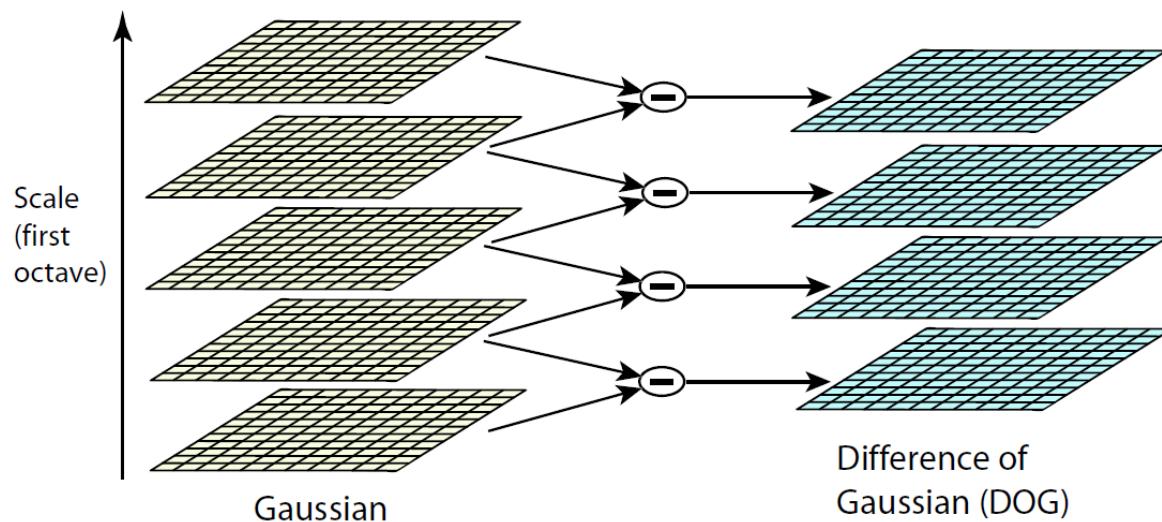
$$\frac{\partial G}{\partial \sigma} = \sigma \nabla^2 G. \rightarrow G(x, y, k\sigma) - G(x, y, \sigma) \approx (k - 1)\sigma^2 \nabla^2 G.$$

Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k
 - Then, maxima and minima are detected by comparing a pixel with its 26 neighbors in 3x3 regions at the current and adjacent scales

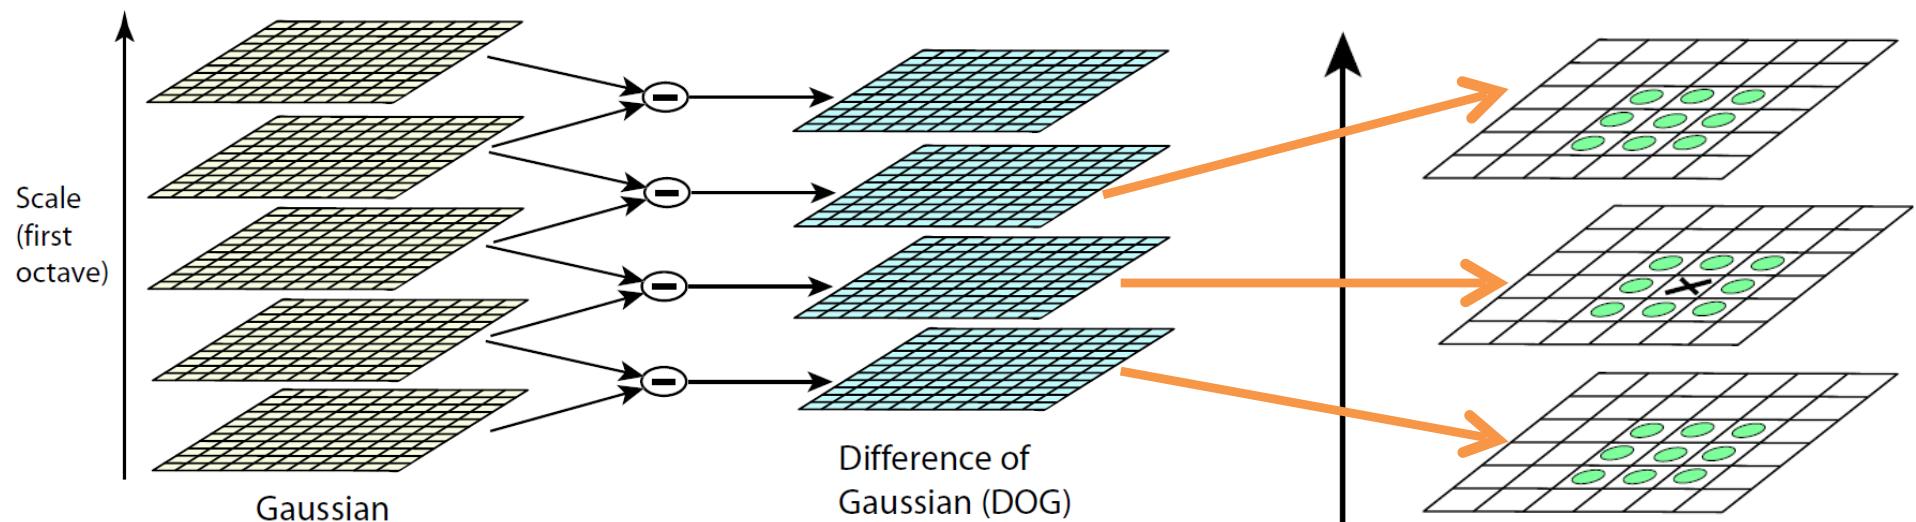
Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k
 - Then, maxima and minima are detected by comparing a pixel with its 26 neighbors in 3x3 regions at the current and adjacent scales



Case study: SIFT

- Scale-space extrema detection
 - Search feature across all possible scales
 - Using difference of Gaussian function of two nearby scales separated by a factor k
 - Then, maxima and minima are detected by comparing a pixel with its 26 neighbors in 3x3 regions at the current and adjacent scales



Case study: SIFT

1. Scale-space extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

Case study: SIFT

- Keypoint localization
 - Keypoint candidates selected in the previous step are fitted to the nearby data for location scale and ratio of principal curvatures
 - This allows to reject points of low contrast (sensitive to noise) or are poorly localized along the edge

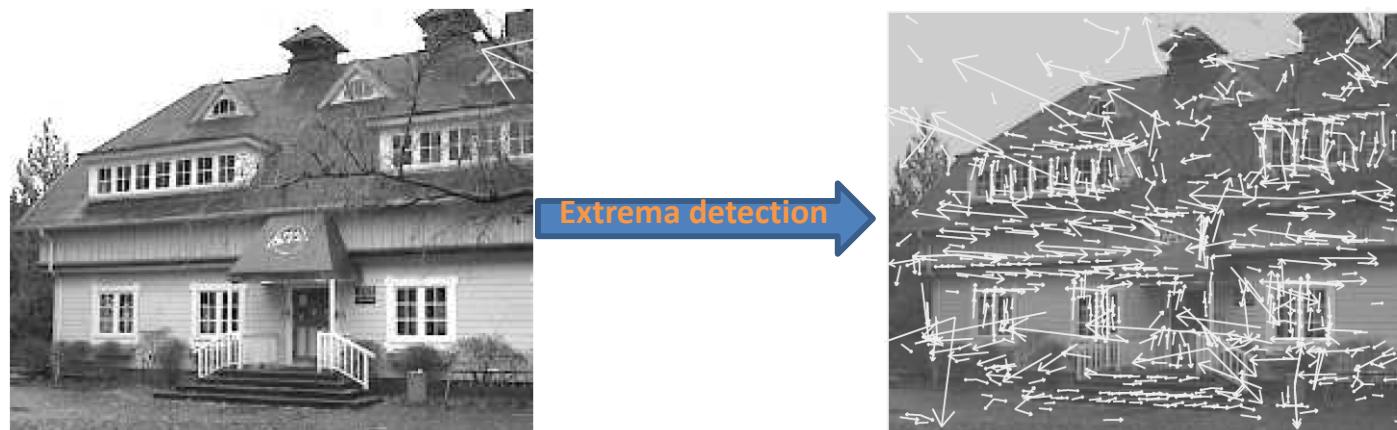
Case study: SIFT

- Keypoint localization



Case study: SIFT

- Keypoint localization



Case study: SIFT

- Keypoint localization



Extrema detection



Threshold on minimum contrast



Case study: SIFT

- Keypoint localization



Extrema detection



Threshold on minimum contrast



Additional thresholds



Case study: SIFT

1. Scale-space extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

Case study: SIFT

- Orientation assignment
 - Each keypoint descriptor is represented relative to a consistent orientation based on local image properties
 - For?

Case study: SIFT

- Orientation assignment
 - Each keypoint descriptor is represented relative to a consistent orientation based on local image properties
 - For? Rotation invariance !

Case study: SIFT

- Orientation assignment
 - Each keypoint descriptor is represented relative to a consistent orientation based on local image properties
 - For? Rotation invariance !
 - For each image sample, $L(x,y)$ at the scale of the keypoint, the gradient magnitude $m(x,y)$ and orientation $\theta(x,y)$ is precomputed using pixel differences

$$m(x,y) = \sqrt{(L(x+1,y) - L(x-1,y))^2 + (L(x,y+1) - L(x,y-1))^2}$$

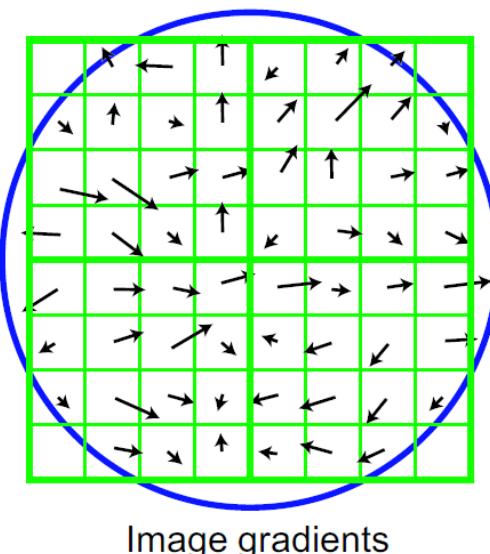
$$\theta(x,y) = \tan^{-1}((L(x,y+1) - L(x,y-1))/(L(x+1,y) - L(x-1,y)))$$

Case study: SIFT

- Orientation assignment
 - Each keypoint descriptor is represented relative to a consistent orientation based on local image properties
 - For? Rotation invariance !
 - For each image sample, $L(x,y)$ at the scale of the keypoint, the gradient magnitude $m(x,y)$ and orientation $\theta(x,y)$ is precomputed using pixel differences.
 - Then, an orientation histogram is formed within a region around a keypoint.
 - Each sample added to the histogram is weighted by its gradient magnitude.
 - Peaks in the orientation histogram correspond to dominant directions of local gradients.
 - The orientations corresponding to the highest peak and local peaks that are within 80% of the highest peaks are assigned to the keypoint

Case study: SIFT

- Orientation assignment
 - Each keypoint descriptor is represented relative to a consistent orientation based on local image properties
 - For? Rotation invariance !
 - For each image sample, $L(x,y)$ at the scale of the keypoint, the gradient magnitude $m(x,y)$ and orientation $\theta(x,y)$ is precomputed using pixel differences.



Case study: SIFT

1. Scale-space extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

Case study: SIFT

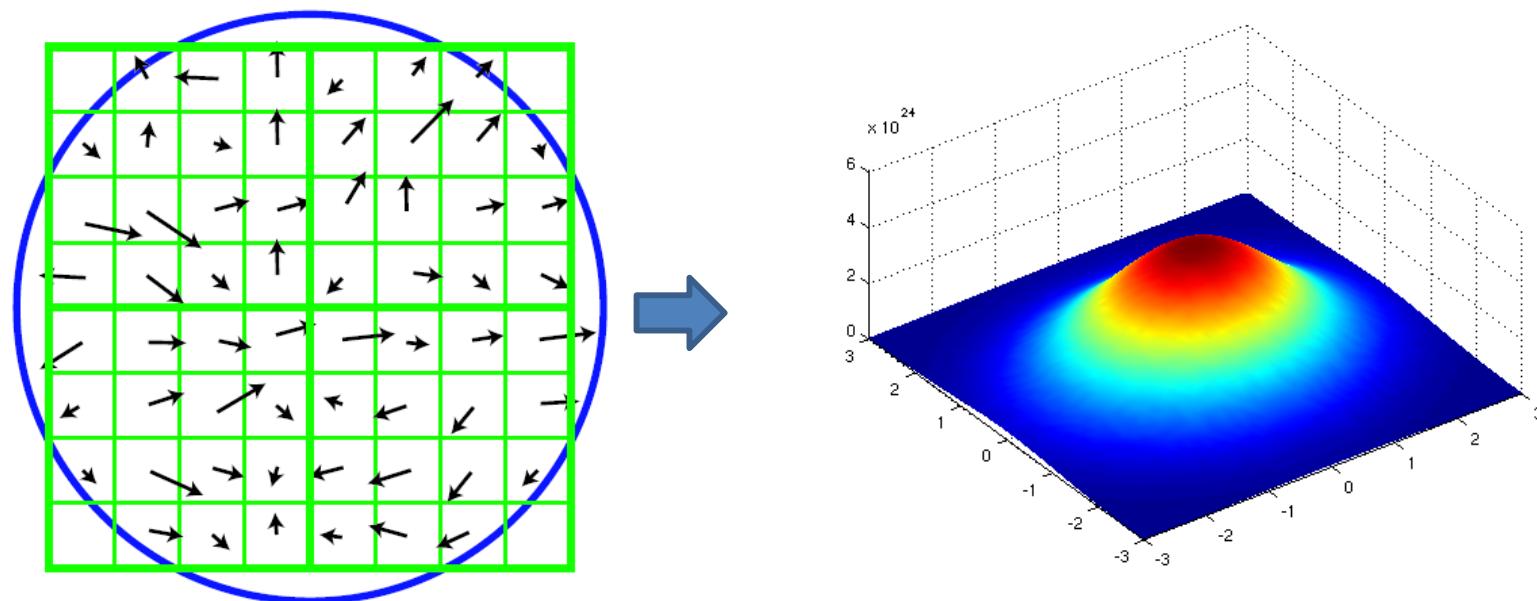
- Keypoint descriptor
 - So far, we have for each keypoint
 - Location
 - Scale and
 - Orientation

Case study: SIFT

- Keypoint descriptor
 - So far, we have for each keypoint
 - Location
 - Scale and
 - Orientation
 - We also need, invariance to
 - Illumination
 - 3D viewpoint

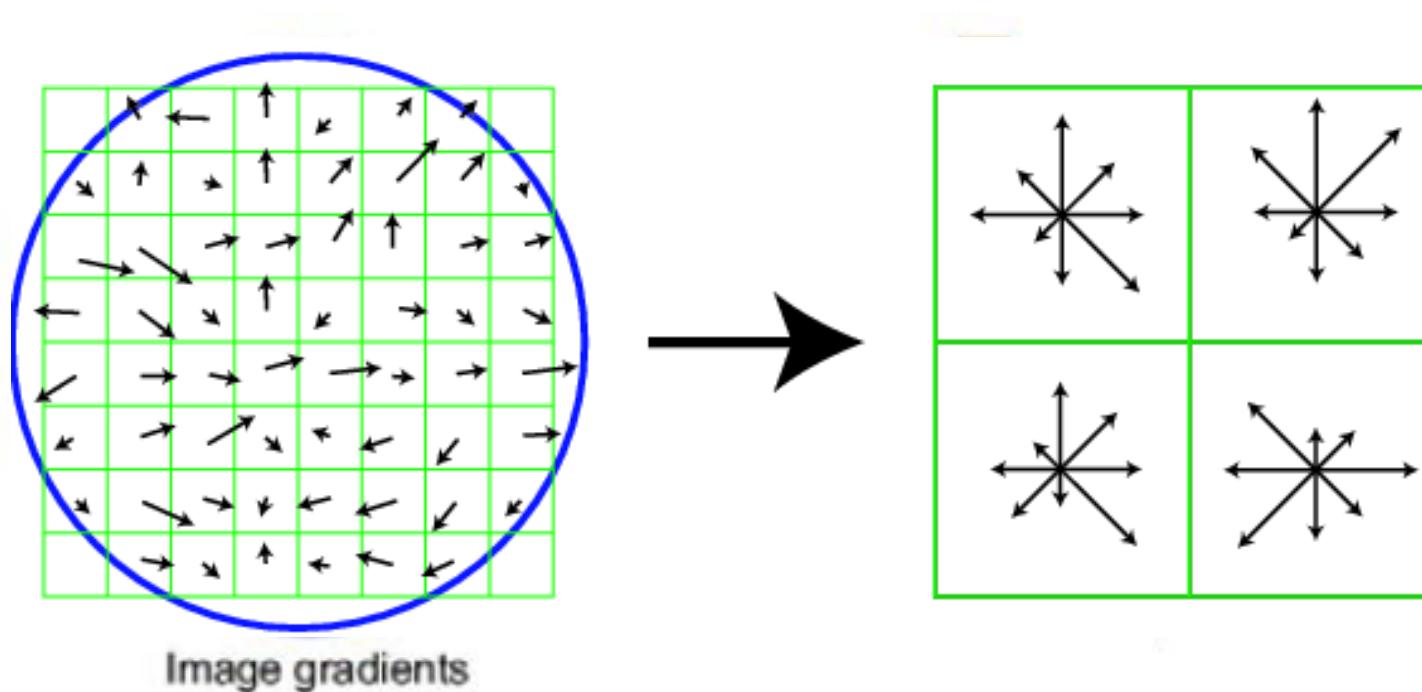
Case study: SIFT

- Keypoint descriptor
 - First, window is Gaussian weighted



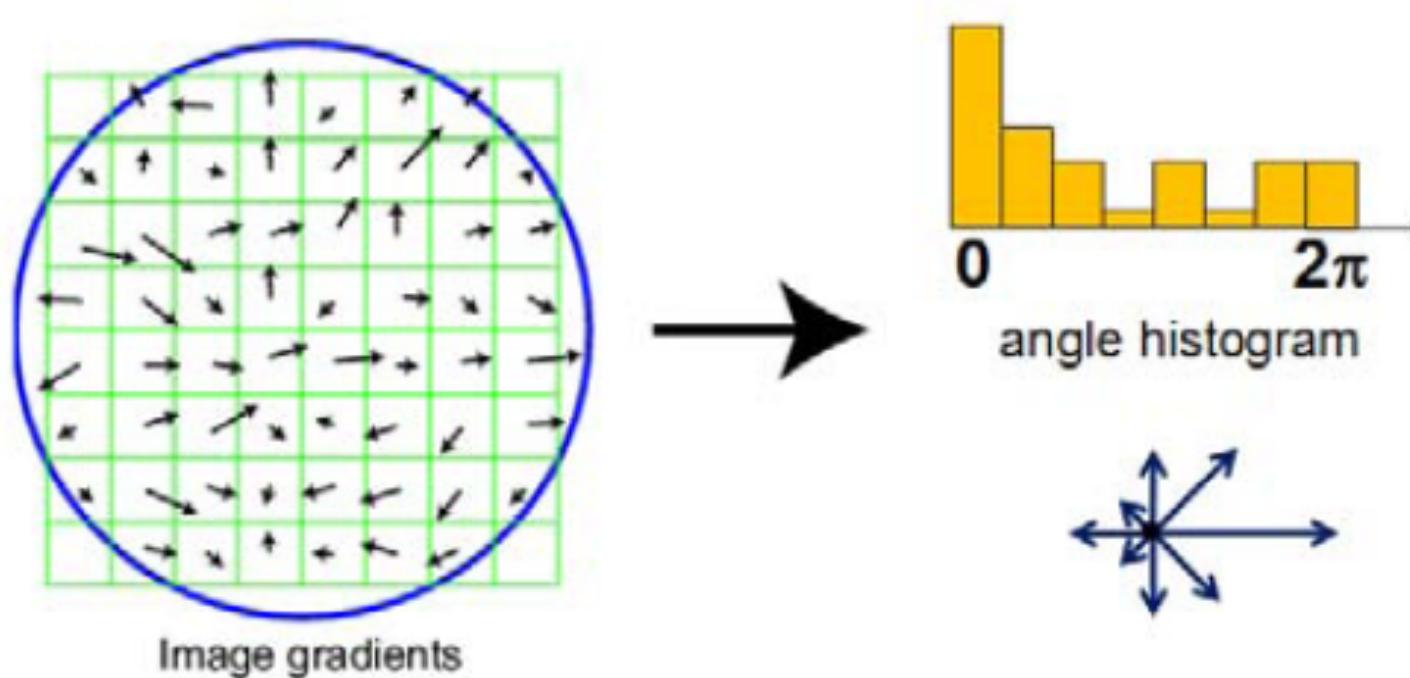
Case study: SIFT

- Keypoint descriptor
 - First, window is Gaussian weighted
 - Then, a histogram is computed in subregions using trilinear interpolations



Case study: SIFT

- Keypoint descriptor
 - First, window is Gaussian weighted
 - Then, a histogram is computed in subregions using trilinear interpolations



Case study: SIFT

- Results in object recognition



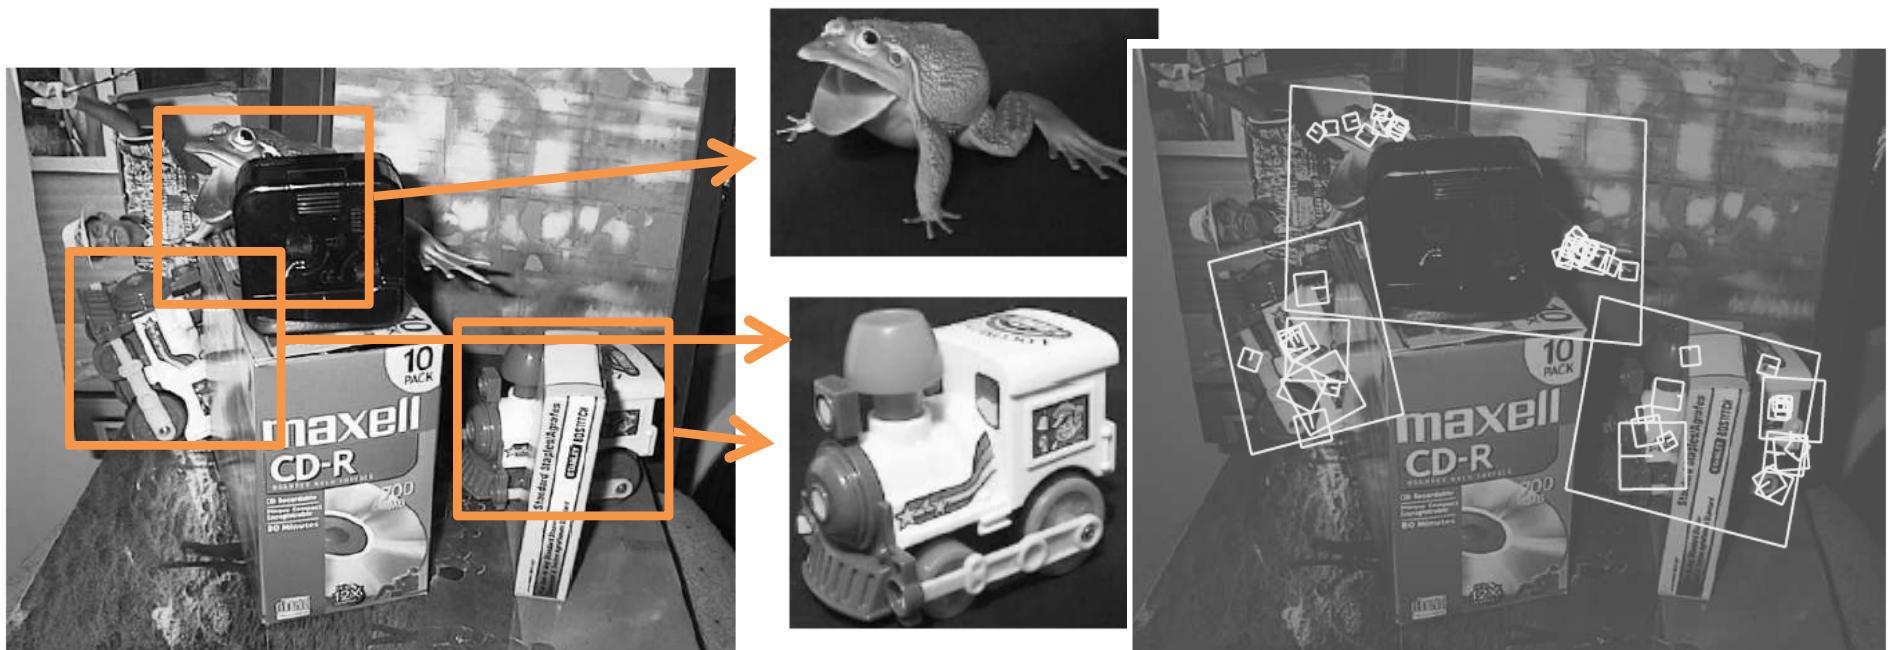
Case study: SIFT

- Results in object recognition



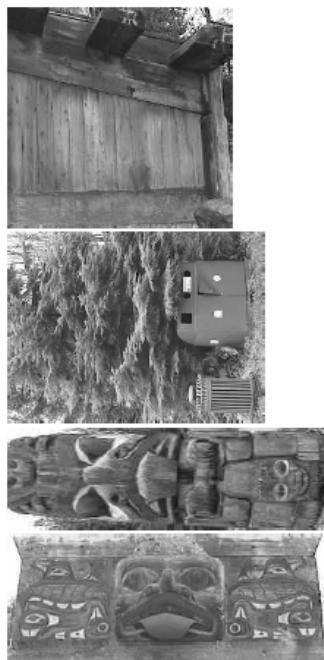
Case study: SIFT

- Results in object recognition



Case study: SIFT

- Results in object recognition



Case study: SIFT

- Results in object recognition



Outline

- Introduction
- Point and patches
- Edges
- Corners

Edge detection

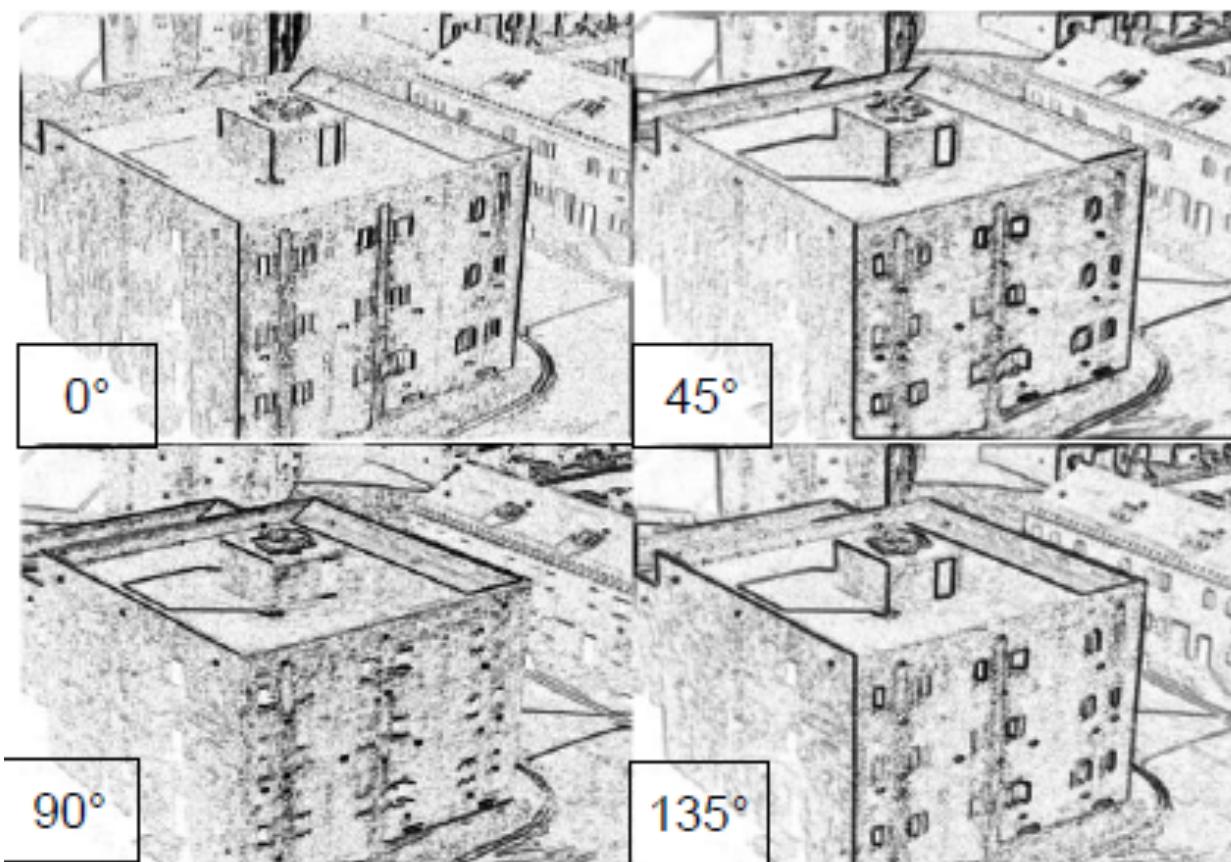
- Sharp changes in image brightness are interesting.
 - Object boundaries generate sharp changes in brightness.
 - Reflectance changes often generate sharp changes in brightness.
 - Sharp changes in surface orientation is also associated with sharp changes in image brightness.
 - Note : Shadows also generate sharp changes in brightness.
- Points in the image where brightness changes sharply are often called edges or edge points.
 - We want those that can be associated with boundaries.
 - But it is difficult to tell a meaningful edge from a nuisance edge.

Edge detection

- What is an edge
 - A border between two regions, each of which have approximately uniform brightness.
 - Edges often arise as the result of occluding contours in an image
 - The two image regions correspond to two different surfaces.



Edge detection



Edge detection

- What is an edge
 - A border between two regions, each of which have approximately uniform brightness.
 - Edges often arise as the result of occluding contours in an image
 - The two image regions correspond to two different surfaces.



Edge detection

- What is an edge
 - A border between two regions, each of which have approximately uniform brightness.
 - Edges often arise as the result of occluding contours in an image.
 - The two image regions correspond to two different surfaces.
 - But again there are other sources of image edges.
 - Abrupt changes in surface orientation.



Edge detection

- What is an edge
 - A border between two regions, each of which have approximately uniform brightness.
 - Edges often arise as the result of occluding contours in an image.
 - The two image regions correspond to two different surfaces.
 - But again there are other sources of image edges.
 - Abrupt changes in surface orientation.
 - Discontinuities in surface reflectance.



Edge detection

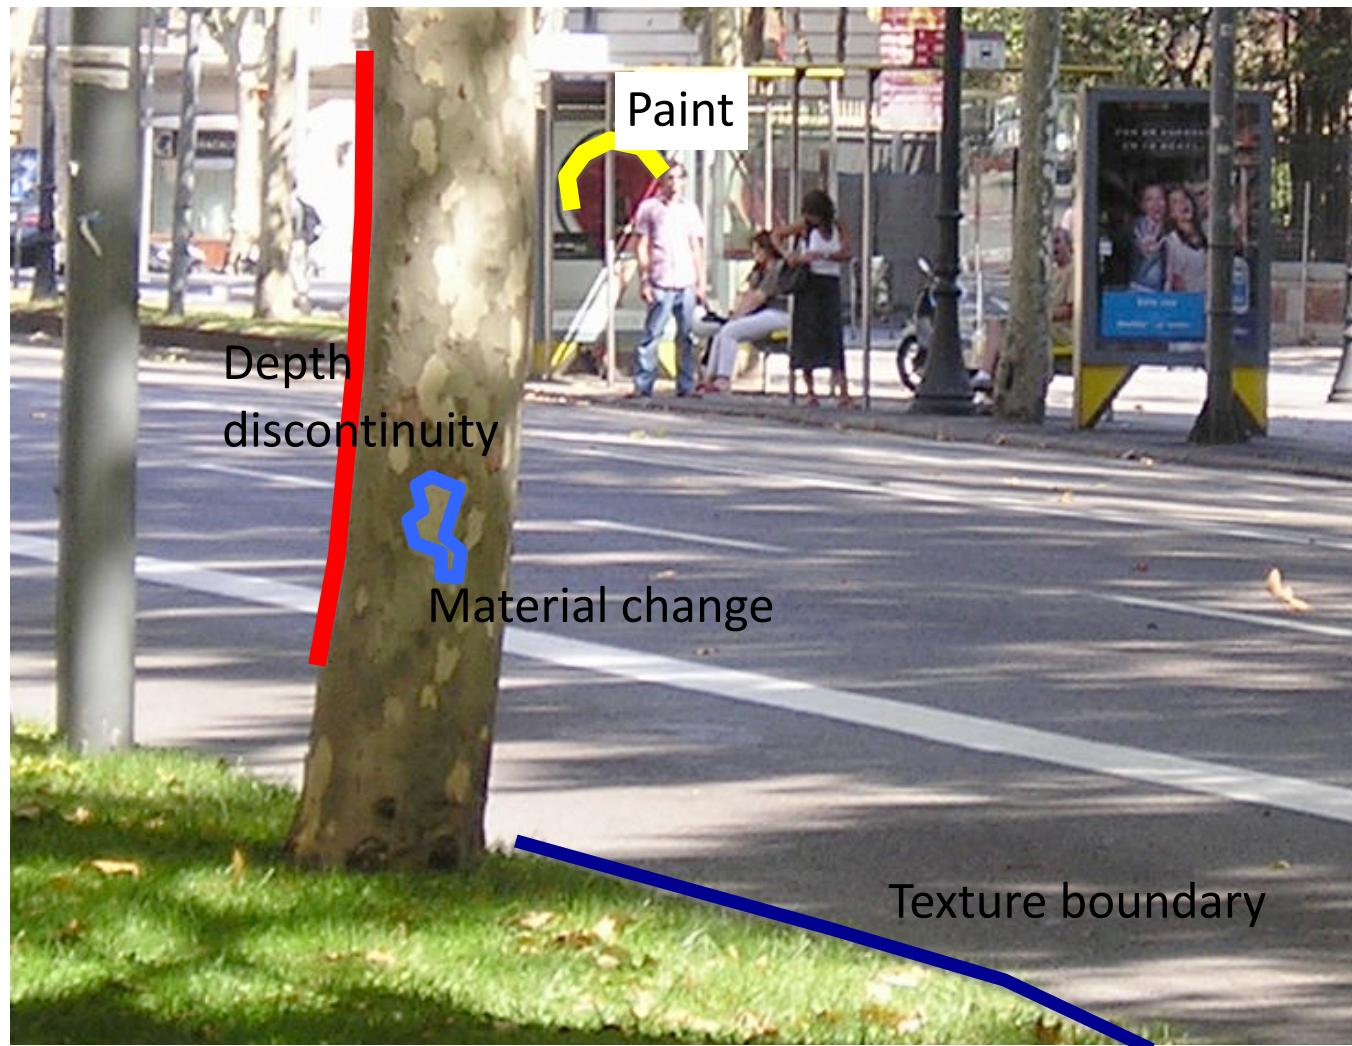
- What is an edge



Source: Antonio Torralba slides

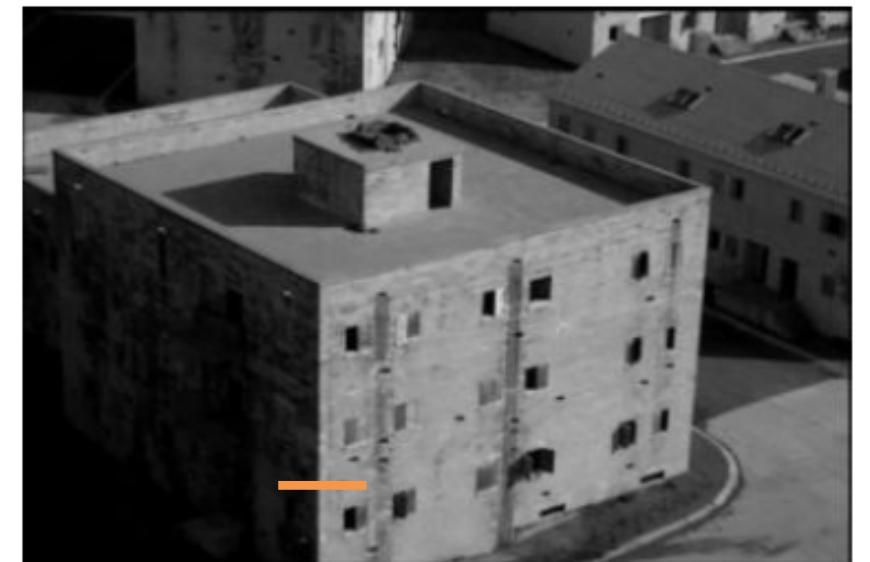
Edge detection

- What is an edge



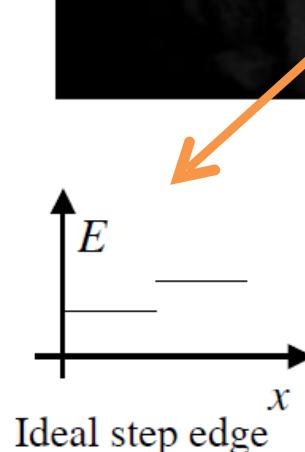
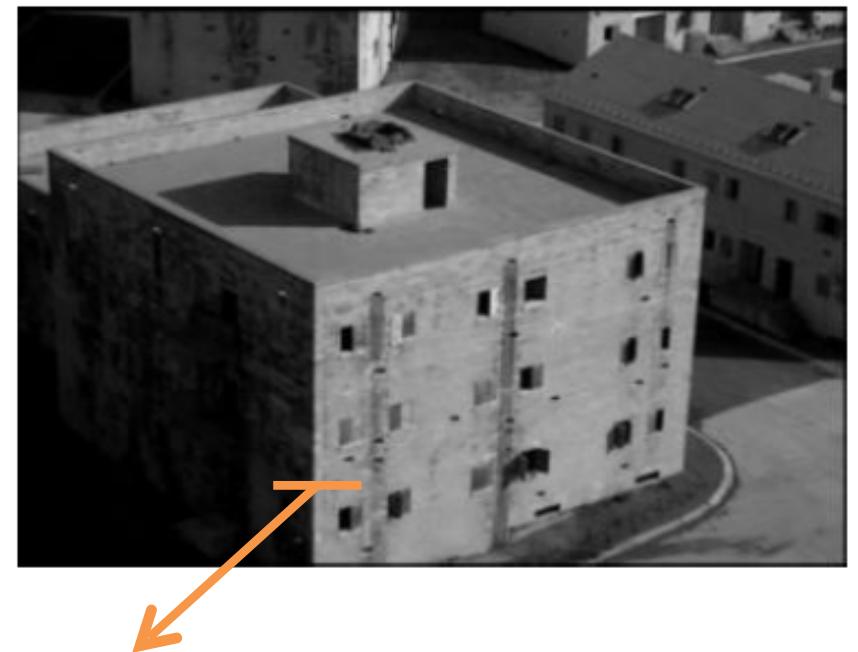
Edge detection

- Simple model
 - Consider a one-dimensional slice through an image in the vicinity of an abrupt change in image brightness.



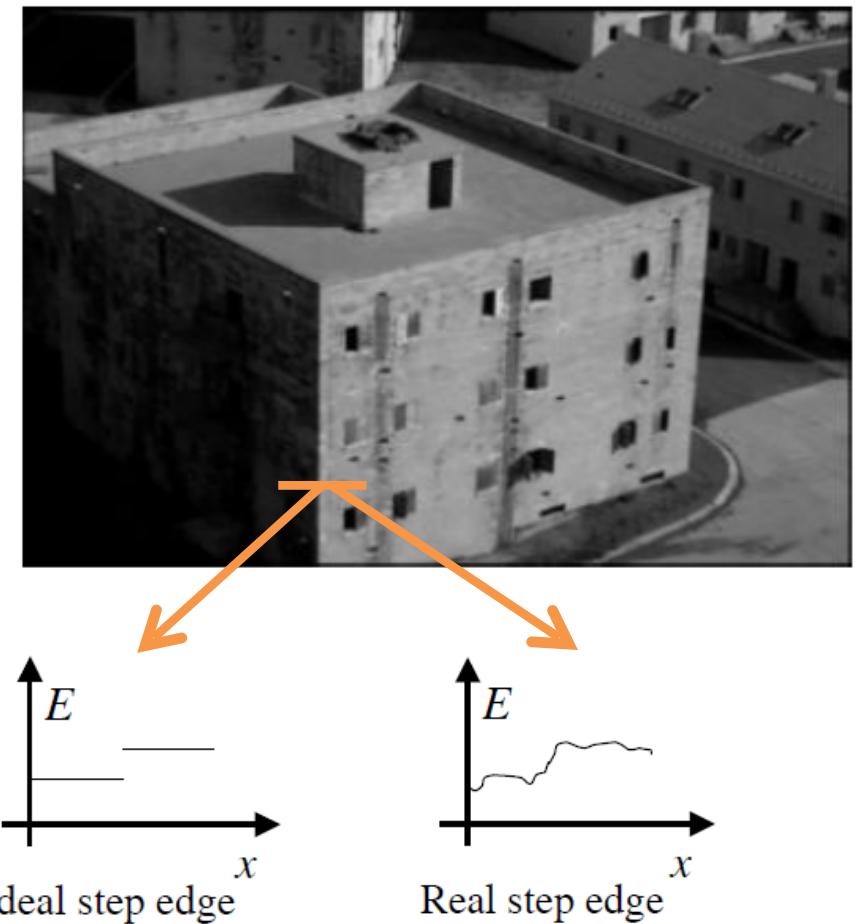
Edge detection

- Simple model
 - Consider a one-dimensional slice through an image in the vicinity of an abrupt change in image brightness.
 - Ideally, we might expect to find a step change in a plot of brightness vs. position.



Edge detection

- Simple model
 - Consider a one-dimensional slice through an image in the vicinity of an abrupt change in image brightness.
 - Ideally, we might expect to find a step change in a plot of brightness vs. position.
 - In practice, we find a corrupted version of this ideal
 - Step transition smoothed
 - Noise transitions superimposed.



Source: Richard Wildes slides

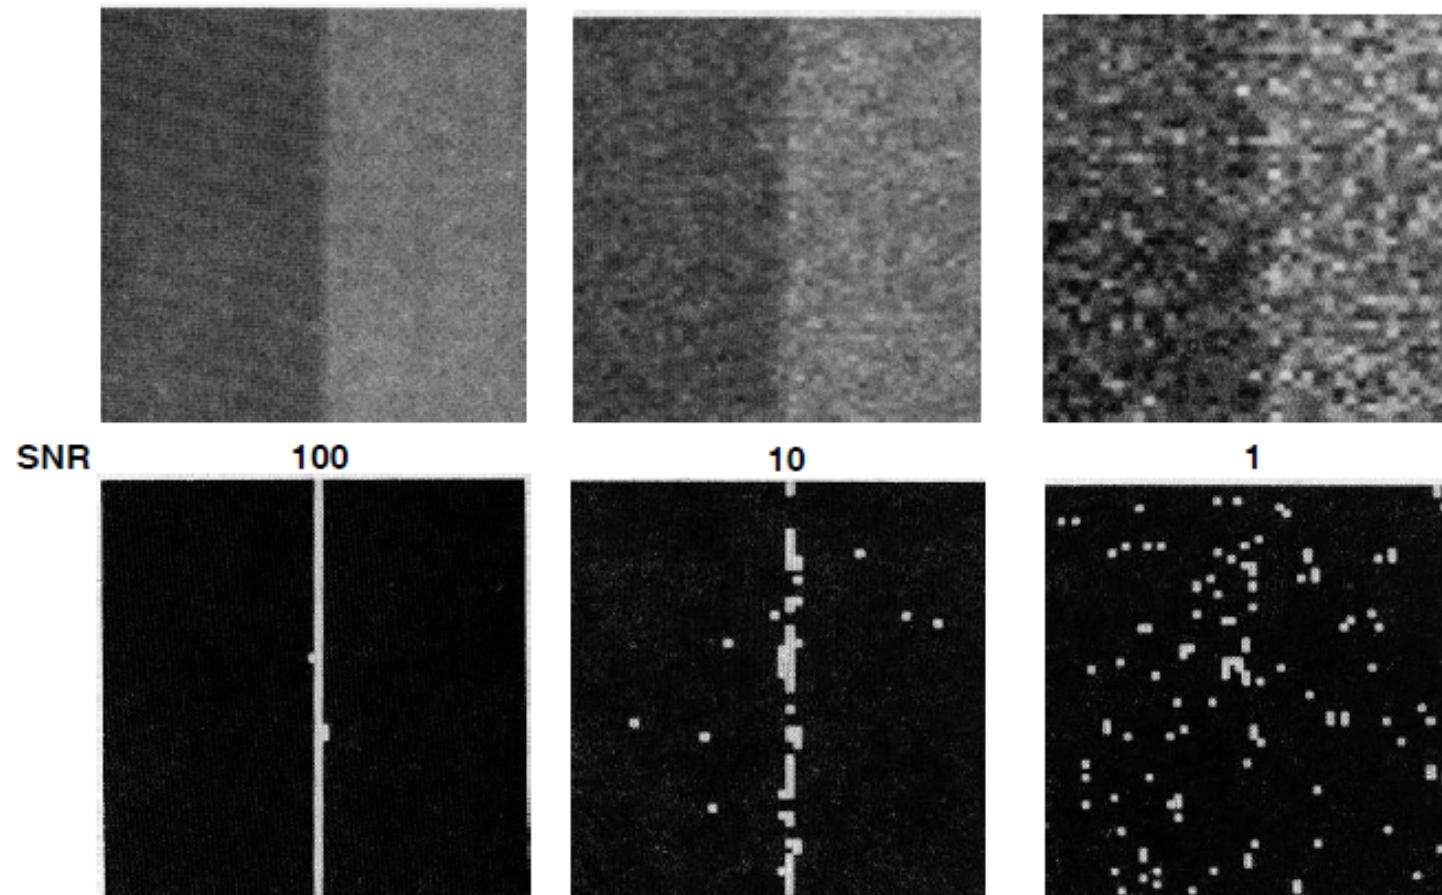
Edge detection

General approach

1. Suppress noise
2. Enhance edges
3. Locate edges

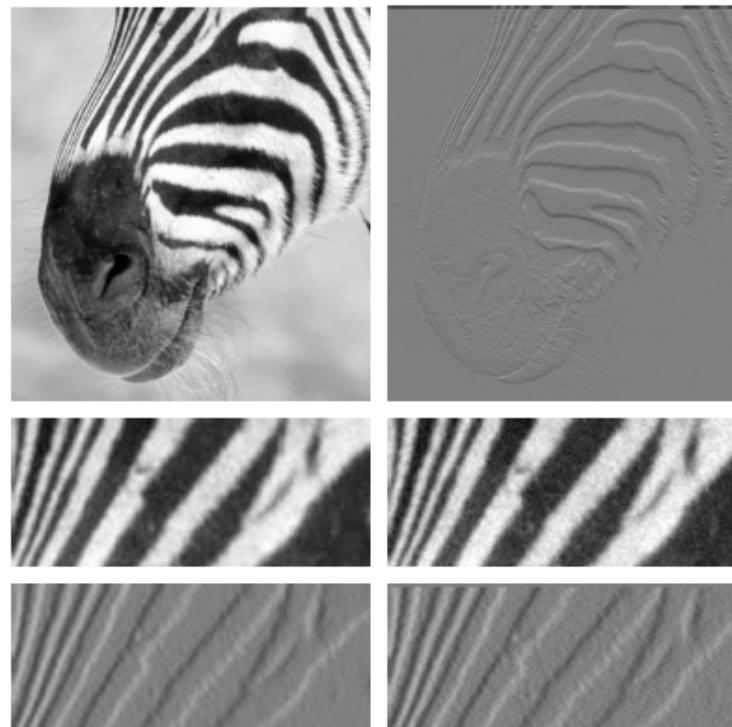
Edge detection: Noise

- Problem for edge detection.



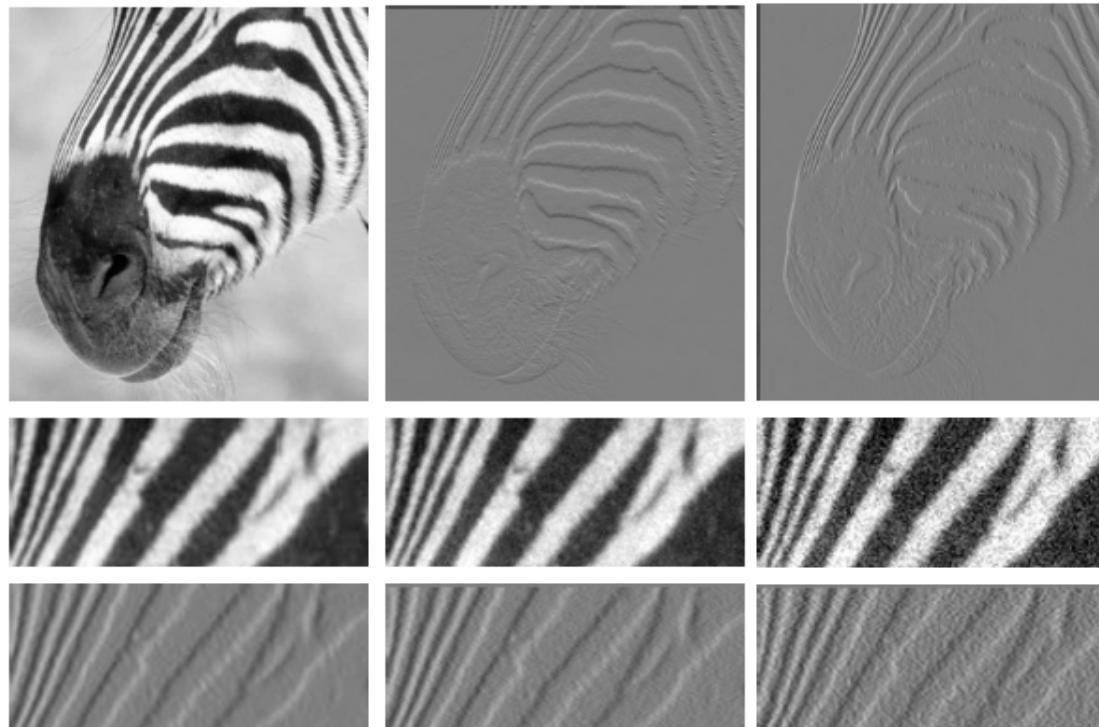
Edge detection: Noise

- Noise suppression
 - Many edge detection filters consist of finite difference computations as we will see next.
 - Differentiation accentuates high frequency components



Edge detection: Noise

- Noise suppression
 - Many edge detection filters consist of finite difference computations as we will see next.
 - Differentiation accentuates high frequency components



Edge detection: Noise

- Noise suppression
 - Many edge detection filters consist of finite difference computations as we will see next.
 - Differentiation accentuates high frequency components
 - We can assume that edges of interest will (typically) have frequency components across a wider range of frequencies (especially) lower frequencies than the noise.

Edge detection: Noise

- Noise suppression
 - In that case, a useful noise suppression is to convolve the image with a Gaussian.

$$h(x, y) = \frac{1}{2\pi\sigma^2} \exp\left[-\frac{1}{2}\left(\frac{x^2 + y^2}{\sigma^2}\right)\right]$$

Edge detection: Noise

- Noise suppression
 - In that case, a useful noise suppression is to convolve the image with a Gaussian.

$$h(x, y) = \frac{1}{2\pi\sigma^2} \exp\left[-\frac{1}{2}\left(\frac{x^2 + y^2}{\sigma^2}\right)\right]$$

- Usually noise suppression and enhancement are combined

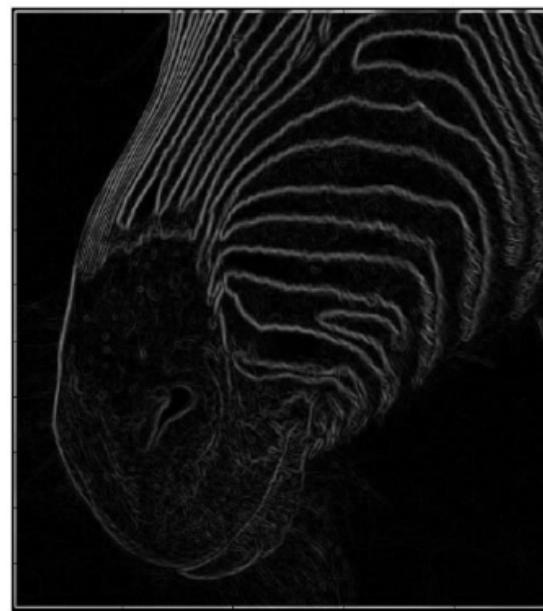
Edge detection: Noise

- Noise suppression

- The gradient magnitude can be estimated by smoothing an image and then differentiating it.
- This is equivalent to convolving with the derivative of a smoothing kernel.

$$h_x(x, y) = -\frac{x}{2\pi\sigma^4} \exp\left[-\frac{1}{2}\left(\frac{x^2 + y^2}{\sigma^2}\right)\right]$$

$$h_y(x, y) = -\frac{y}{2\pi\sigma^4} \exp\left[-\frac{1}{2}\left(\frac{x^2 + y^2}{\sigma^2}\right)\right]$$



Edge detection

General approach

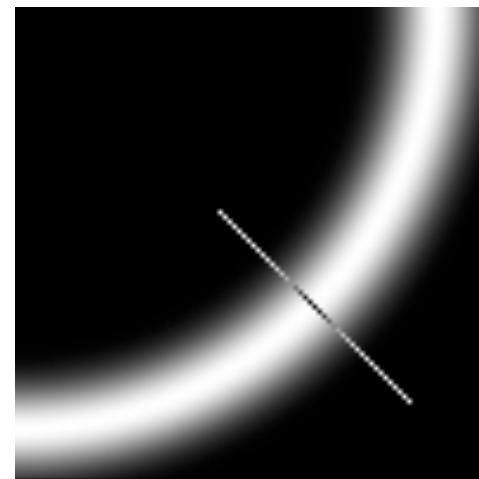
1. Suppress noise
2. Enhance edges
3. Locate edges

Edge detection: Enhancement

- Gradient-based edge detectors
 - We compute an estimate of the gradient magnitude
 - Look for points where the gradient magnitude is a maximum along the direction perpendicular to the edge
 - The direction perpendicular to the edge can be estimated using the direction of the gradient

Edge detection: Enhancement

- Gradient-based edge detectors
 - We compute an estimate of the gradient magnitude
 - Look for points where the gradient magnitude is a maximum along the direction perpendicular to the edge
 - The direction perpendicular to the edge can be estimated using the direction of the gradient



Edge detection: Enhancement

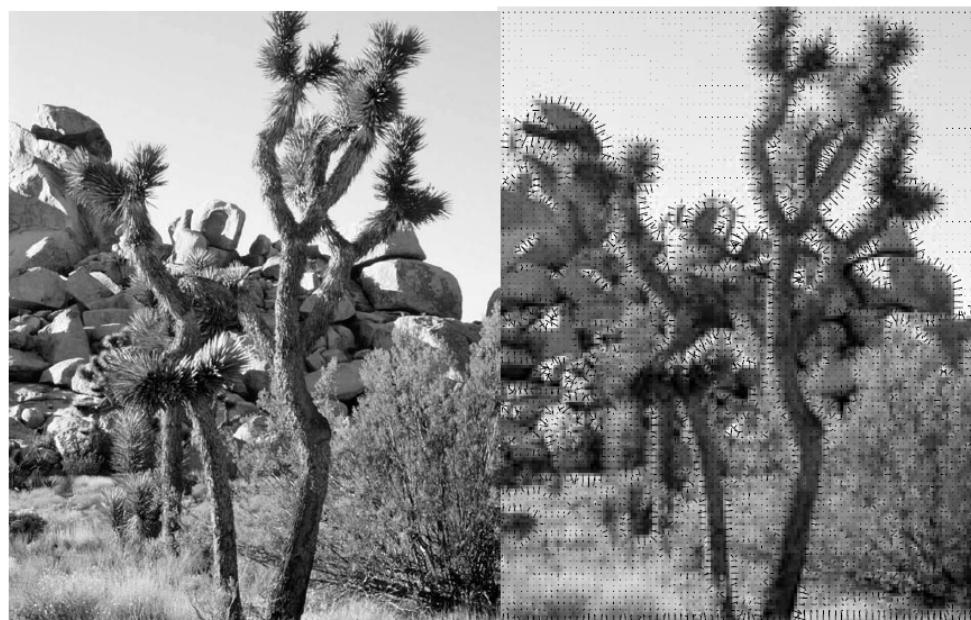
- Gradient-based edge detectors
 - We compute an estimate of the gradient magnitude
 - Look for points where the gradient **magnitude** is a maximum along the direction perpendicular to the edge
 - The direction perpendicular to the edge can be estimated using the direction of the gradient



Images credit: Christopher Rasmussen & Antonio Torralba

Edge detection: Enhancement

- Gradient-based edge detectors
 - We compute an estimate of the gradient magnitude
 - Look for points where the gradient magnitude is a maximum along the direction perpendicular to the edge
 - The **direction** perpendicular to the edge can be estimated using the direction of the gradient



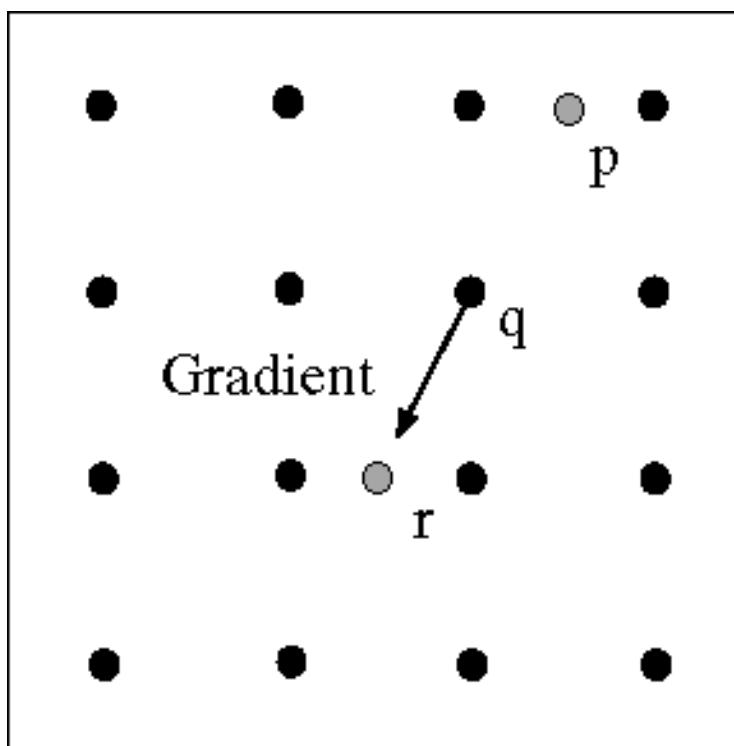
Edge detection: Localization

- Gradient-based edge detectors: Non-maximum suppression
 - Points where the gradient magnitude is at a maximum along the direction of the gradient.



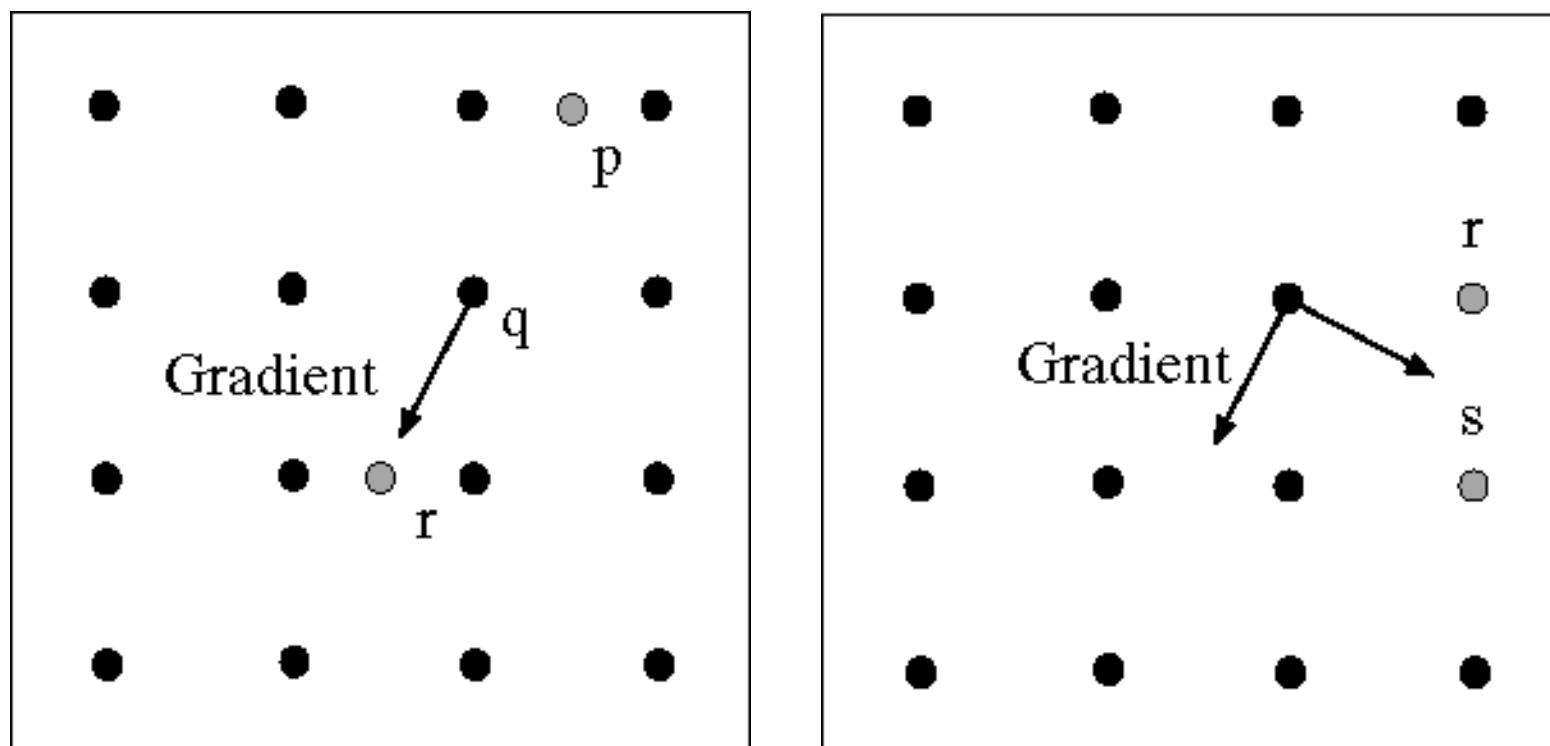
Edge detection: Localization

- Gradient-based edge detectors: Non-maximum suppression
 - Points where the gradient magnitude is at a maximum along the direction of the gradient.



Edge detection: Localization

- Gradient-based edge detectors: Non-maximum suppression
 - Points where the gradient magnitude is at a maximum along the direction of the gradient.



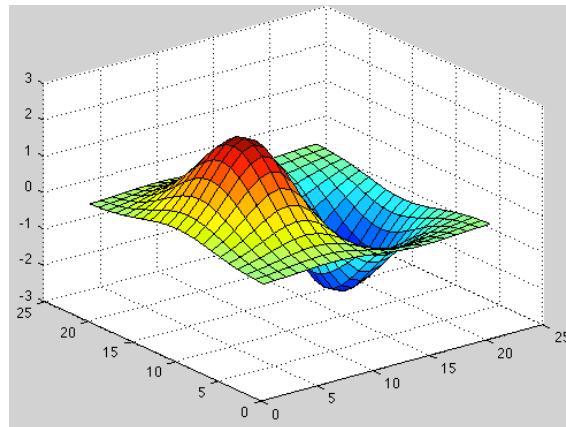
Edge detection

- Gradient-based edge detectors: Canny edge detector (1986)
 1. Noise reduction: Gaussian filter.

Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.

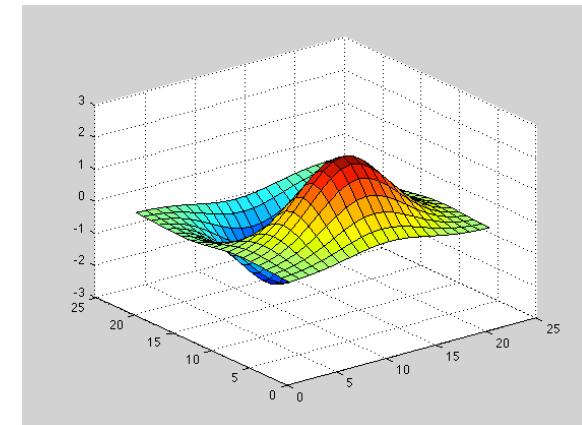
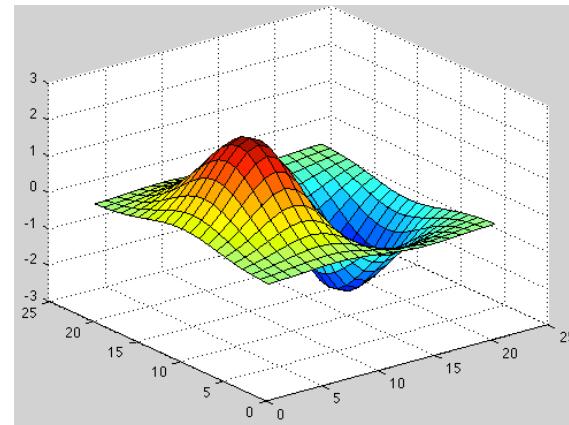
$$h_x(x,y) = \frac{\partial h(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.

$$h_x(x,y) = \frac{\partial h(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}} \quad h_y(x,y) = \frac{\partial h(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.

$$h_x(x,y) = \frac{\partial h(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}} \quad h_y(x,y) = \frac{\partial h(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Magnitude: $h_x(x,y)^2 + h_y(x,y)^2$ *Edge strength*

Angle: $\arctan\left(\frac{h_y(x,y)}{h_x(x,y)}\right)$ *Edge normal*

Edge detection

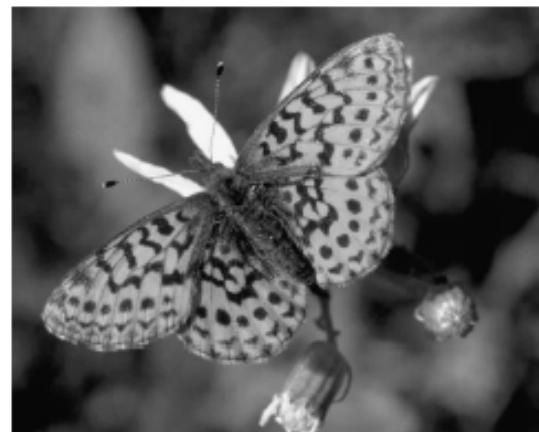
- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.

Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.
 4. Hysteresis.
 - By just applying non-maximum suppression there are too many of these curves.
 - Usually we apply a threshold, but ...

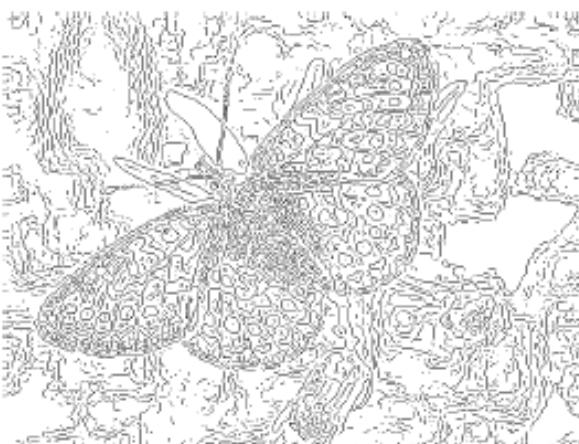
Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.
 4. Hysteresis.



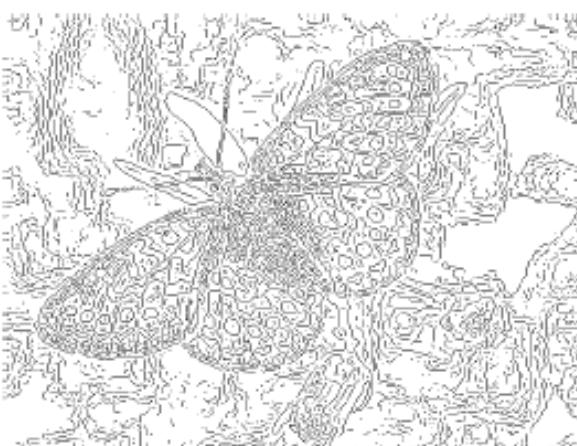
Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.
 4. Hysteresis.



Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.
 4. Hysteresis.

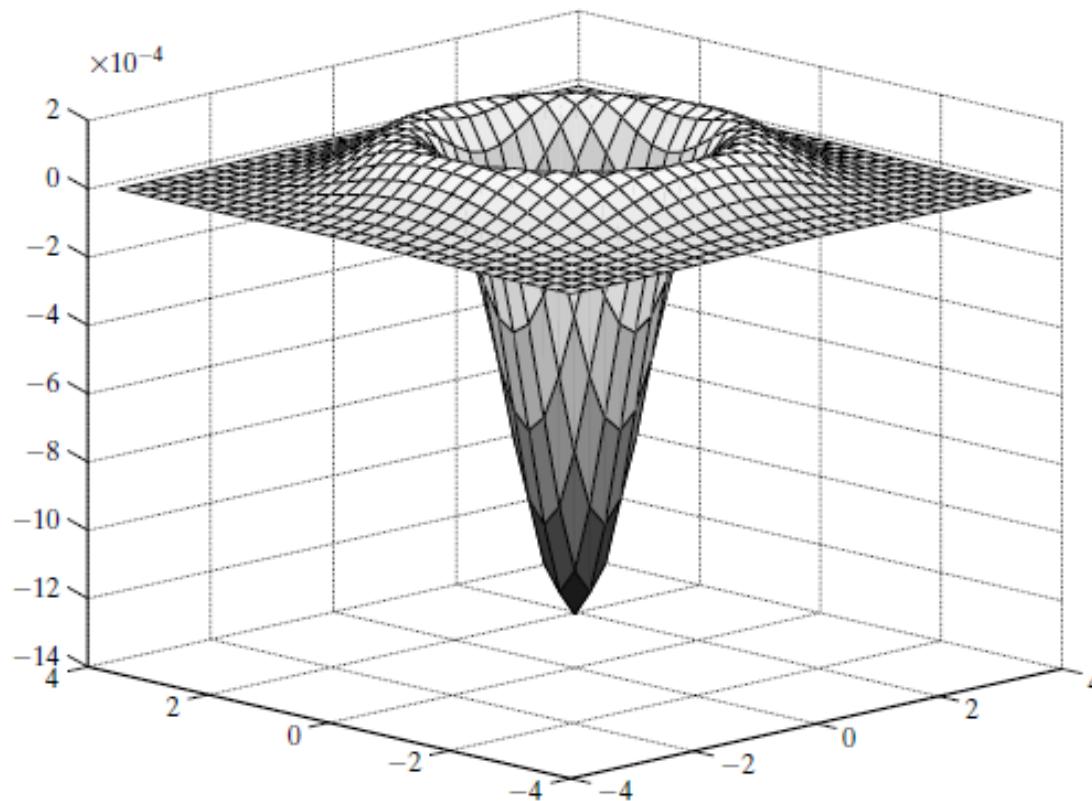


Edge detection

- Gradient-based edge detectors: Canny edge detector
 1. Noise reduction: Gaussian filter.
 2. Four filters to detect horizontal, vertical and diagonal edges.
 3. Non-maximum suppression.
 4. Hysteresis.
 - Solution: Have two thresholds.
 - Use the larger when starting an edge chain.
 - Use the smaller while following it.

Edge detection: Enhancement

- Laplacian
 - Remember: Second partial derivatives



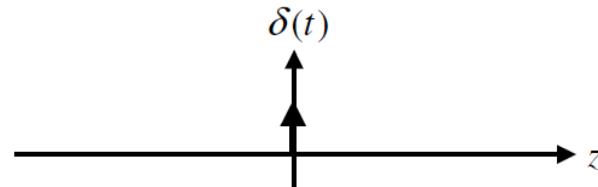
Edge detection: Enhancement

- Laplacian
 - Remember: Second partial derivatives

$$\frac{\partial^2 E}{\partial x^2} = \sin^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

$$\frac{\partial^2 E}{\partial y^2} = \cos^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

- Note
 - Remember δ is the impulse



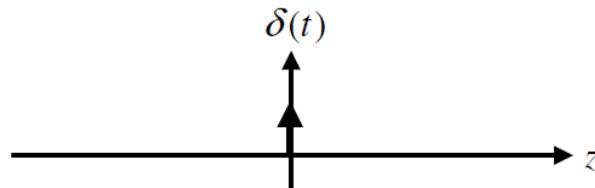
Edge detection: Enhancement

- Laplacian
 - Remember: Second partial derivatives

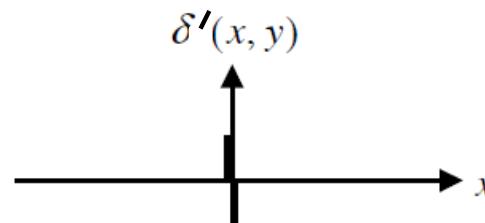
$$\frac{\partial^2 E}{\partial x^2} = \sin^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

$$\frac{\partial^2 E}{\partial y^2} = \cos^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

- Note
 - Remember δ is the impulse



- δ' is the impulse derivative



Edge detection: Enhancement

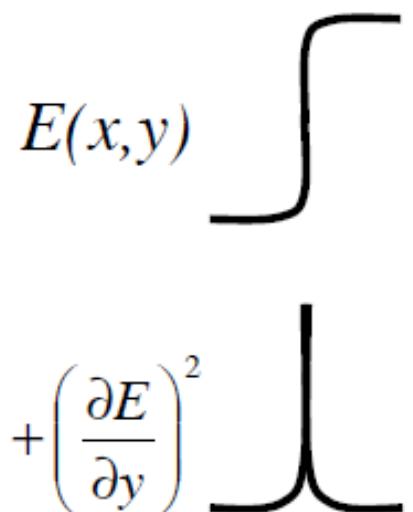
- Laplacian
 - Remember: Second partial derivatives

$$\frac{\partial^2 E}{\partial x^2} = \sin^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

$$\frac{\partial^2 E}{\partial y^2} = \cos^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

- We define the Laplacian as

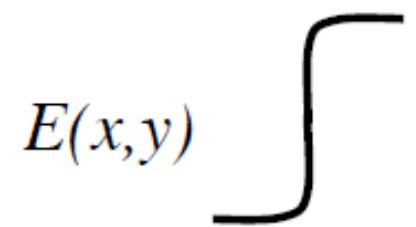
$$\nabla^2 E = \frac{\partial^2 E}{\partial x^2} + \frac{\partial^2 E}{\partial y^2} = (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$



Edge detection: Enhancement

- Laplacian
 - Remember: Second partial derivatives

$$\frac{\partial^2 E}{\partial x^2} = \sin^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$



$$\frac{\partial^2 E}{\partial y^2} = \cos^2 \theta (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$

- We define the Laplacian as

$$\left(\frac{\partial E}{\partial x} \right)^2 + \left(\frac{\partial E}{\partial y} \right)^2$$

$$\nabla^2 E = \frac{\partial^2 E}{\partial x^2} + \frac{\partial^2 E}{\partial y^2} = (B2 - B1) \delta'(x \sin \theta - y \cos \theta + r)$$



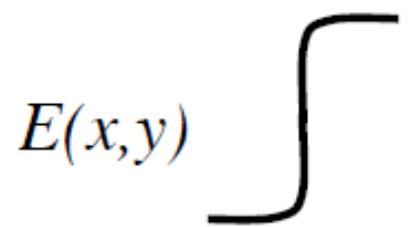
- This operation will show a “zero-crossing” as we cross the edge.



$$\nabla^2 E$$

Edge detection: Enhancement

- Laplacian
 - This operation will show a “zero-crossing” as we cross the edge.
 - Similarity with the squared gradient
 - The response of the operator is independent of the edge orientation.



$$\left(\frac{\partial E}{\partial x}\right)^2 + \left(\frac{\partial E}{\partial y}\right)^2$$

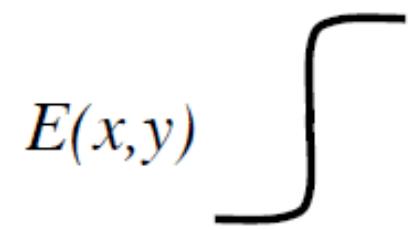
A graph of the squared gradient magnitude $\left(\frac{\partial E}{\partial x}\right)^2 + \left(\frac{\partial E}{\partial y}\right)^2$ plotted against position. The curve is zero everywhere except at the edge, where it has a very sharp, narrow peak reaching up to a maximum value.

$$\nabla^2 E$$

A graph of the Laplacian of the edge response $\nabla^2 E$ plotted against position. The curve is zero everywhere except at the edge, where it has a very sharp, narrow peak reaching up to a maximum value.

Edge detection: Enhancement

- Laplacian
 - This operation will show a “zero-crossing” as we cross the edge.
 - Similarity with the squared gradient
 - The response of the operator is independent of the edge orientation.
 - But unlike the squared gradient
 - The Laplacian is linear
 - The Laplacian the sign of intensity change across the edge.



$$\left(\frac{\partial E}{\partial x}\right)^2 + \left(\frac{\partial E}{\partial y}\right)^2$$

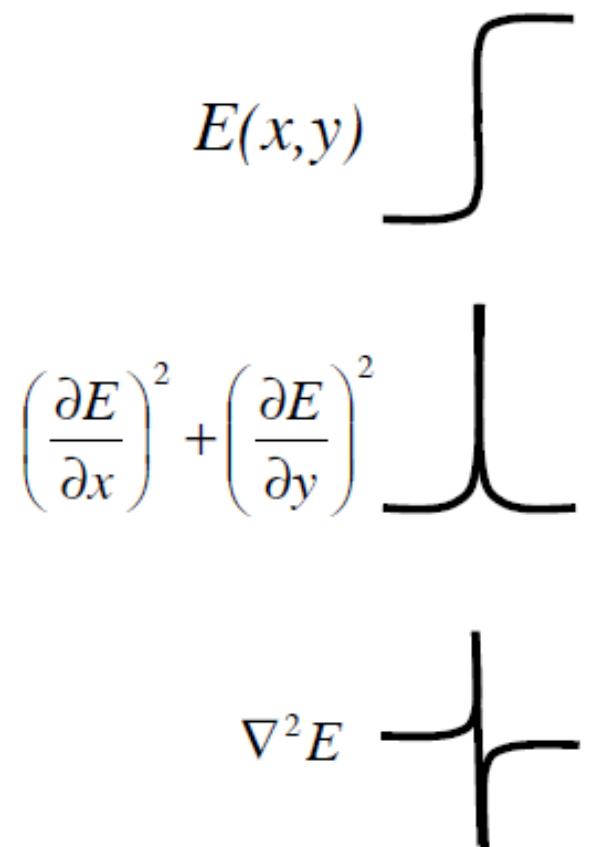
A graph of the squared gradient of $E(x,y)$ plotted against position. The curve is zero everywhere except at the edge, where it has a very sharp, narrow peak reaching up to a maximum value. This represents the squared gradient response across the edge.

$$\nabla^2 E$$

A graph of the Laplacian of $E(x,y)$ plotted against position. The curve is zero everywhere except at the edge, where it crosses the zero line from positive to negative values or vice versa. This represents the sign change in the intensity gradient across the edge.

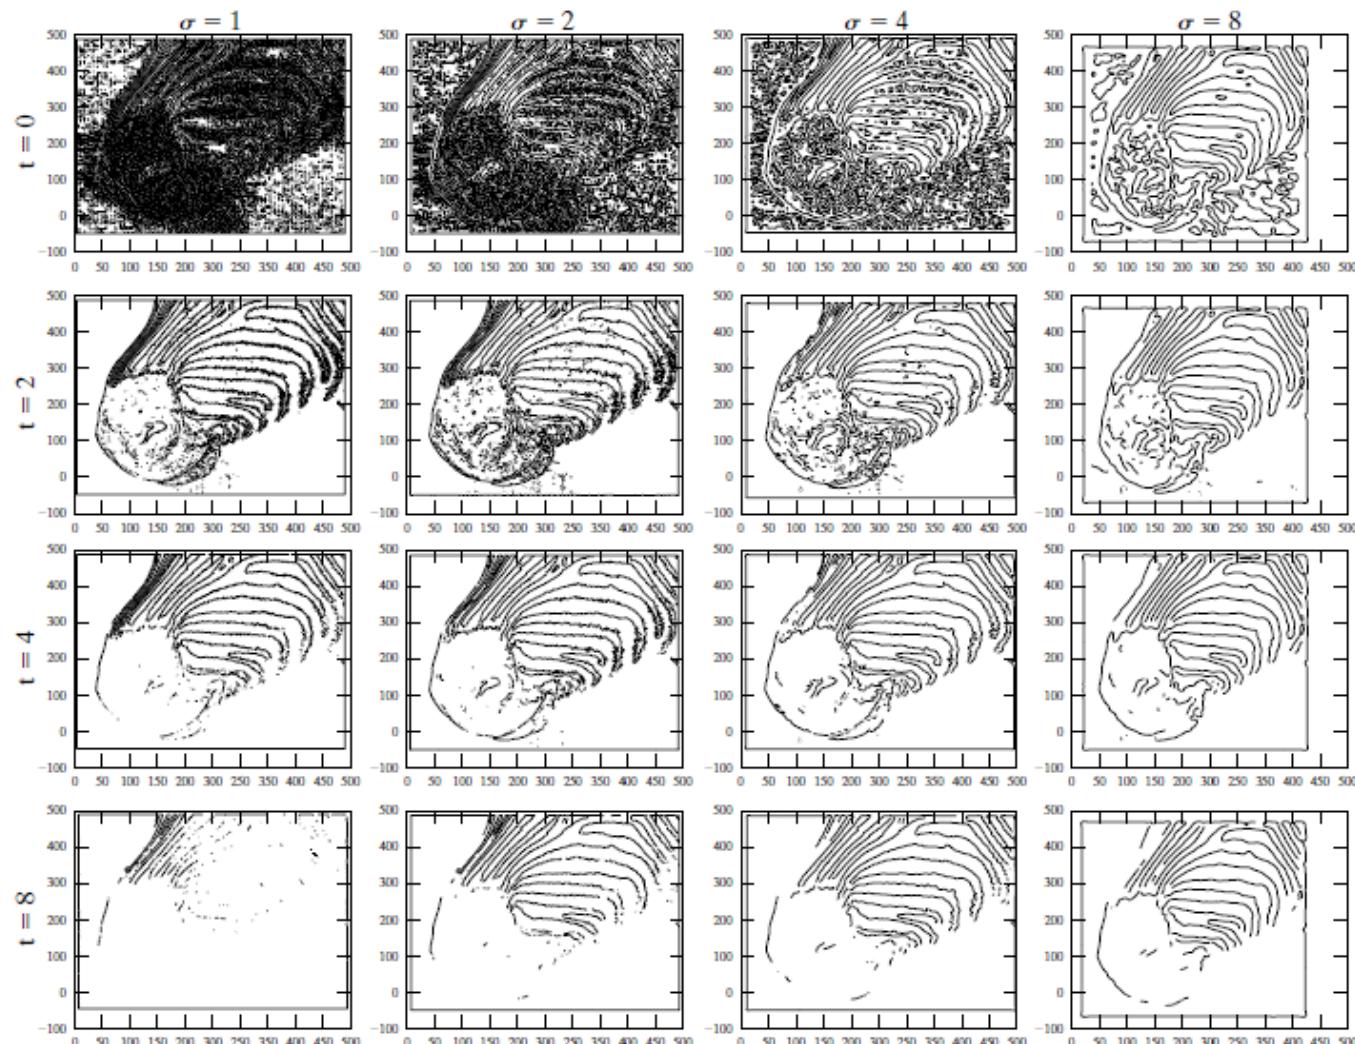
Edge detection: Localization

- Laplacian
 - Remember for the gradient we would use a threshold and non-maxima suppression
 - For the Laplacian we seek a transition magnitude across the zero-crossing above which we will declare a value as marking an edge location.



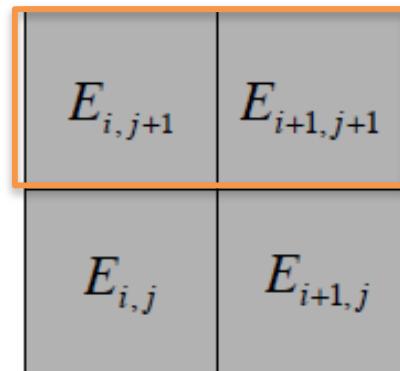
Edge detection: Localization

- Laplacian



Edge detection

- Discrete approximations: The gradient



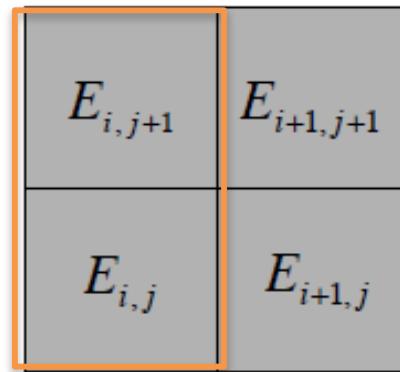
- We can make the following approximation

$$\frac{\partial E}{\partial x} \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i,j+1}) + (E_{i+1,j} - E_{i,j})]$$

$$\frac{\partial E}{\partial y} \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i+1,j}) + (E_{i,j+1} - E_{i,j})]$$

Edge detection

- Discrete approximations: The gradient



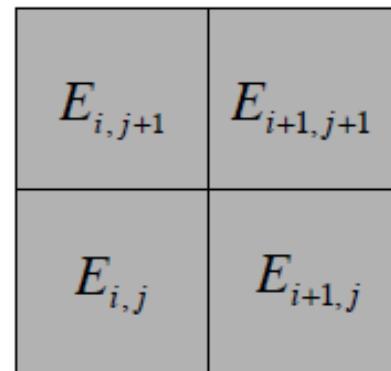
- We can make the following approximation

$$\frac{\partial E}{\partial x} \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i,j+1}) + (E_{i+1,j} - E_{i,j})]$$

$$\frac{\partial E}{\partial y} \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i+1,j}) + (E_{i,j+1} - E_{i,j})]$$

Edge detection

- Discrete approximations: The gradient



- The Gradient is

$$\left(\frac{\partial E}{\partial x}\right)^2 + \left(\frac{\partial E}{\partial y}\right)^2 \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i,j})^2 + (E_{i,j+1} - E_{i+1,j})^2]$$

Edge detection

- Discrete approximations: The gradient

$E_{i,j+1}$	$E_{i+1,j+1}$
$E_{i,j}$	$E_{i+1,j}$

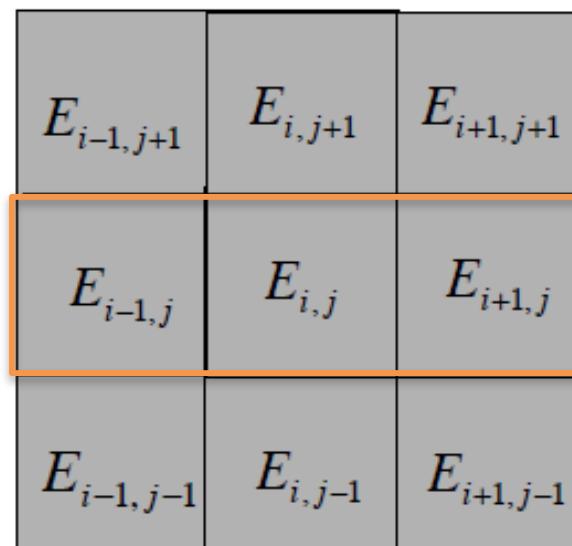
- The Gradient is

☞
$$\left(\frac{\partial E}{\partial x} \right)^2 + \left(\frac{\partial E}{\partial y} \right)^2 \approx \frac{1}{2\varepsilon} [(E_{i+1,j+1} - E_{i,j})^2 + (E_{i,j+1} - E_{i+1,j})^2]$$

- We obtain large values where the image brightness is changing rapidly.
- We obtain a new image with edges strongly emphasized.

Edge detection

- Discrete approximations: The Laplacian



- The approximation is

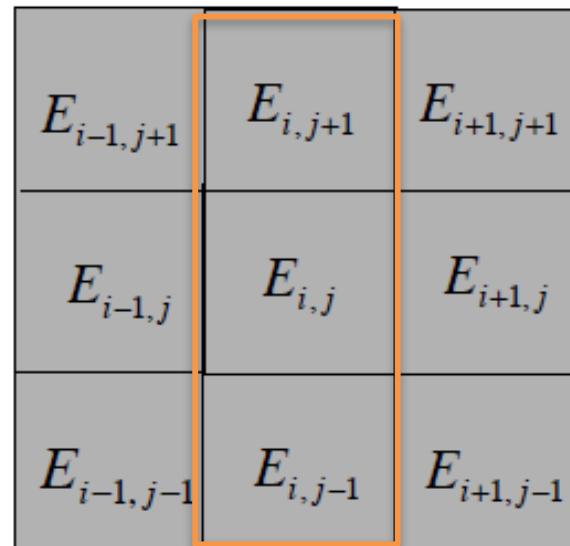
$$\frac{\partial^2 E}{\partial x^2} \approx \frac{1}{\varepsilon^2} (E_{i-1,j} - 2E_{i,j} + E_{i+1,j})$$

$$\frac{\partial^2 E}{\partial y^2} \approx \frac{1}{\varepsilon^2} (E_{i,j-1} - 2E_{i,j} + E_{i,j+1})$$

Source: Richard Wildes slides

Edge detection

- Discrete approximations: The Laplacian



- The approximation is

$$\frac{\partial^2 E}{\partial x^2} \approx \frac{1}{\varepsilon^2} (E_{i-1,j} - 2E_{i,j} + E_{i+1,j})$$

$$\frac{\partial^2 E}{\partial y^2} \approx \frac{1}{\varepsilon^2} (E_{i,j-1} - 2E_{i,j} + E_{i,j+1})$$

Source: Richard Wildes slides

Edge detection

- Discrete approximations: The Laplacian

$E_{i-1,j+1}$	$E_{i,j+1}$	$E_{i+1,j+1}$
$E_{i-1,j}$	$E_{i,j}$	$E_{i+1,j}$
$E_{i-1,j-1}$	$E_{i,j-1}$	$E_{i+1,j-1}$

- The Laplacian is then

$$\frac{\partial^2 E}{\partial x^2} + \frac{\partial^2 E}{\partial y^2} \approx \frac{4}{\varepsilon^2} \left[\frac{1}{4} (E_{i-1,j} + E_{i,j-1} + E_{i+1,j} + E_{i,j+1}) - E_{i,j} \right]$$

Edge detection

- Discrete approximations: The Laplacian

$E_{i-1,j+1}$	$E_{i,j+1}$	$E_{i+1,j+1}$
$E_{i-1,j}$	$E_{i,j}$	$E_{i+1,j}$
$E_{i-1,j-1}$	$E_{i,j-1}$	$E_{i+1,j-1}$

- The Laplacian is then

$$\frac{\partial^2 E}{\partial x^2} + \frac{\partial^2 E}{\partial y^2} \approx \frac{4}{\varepsilon^2} \left[\frac{1}{4} (E_{i-1,j} + E_{i,j-1} + E_{i+1,j} + E_{i,j+1}) - E_{i,j} \right]$$

- This is the average of the surrounding values and subtracts that of the center

Edge detection

- Discrete approximations: The Laplacian
 - It can be mapped onto this mask.

$$\frac{1}{\varepsilon^2} \begin{matrix} & \begin{matrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{matrix} \end{matrix}$$

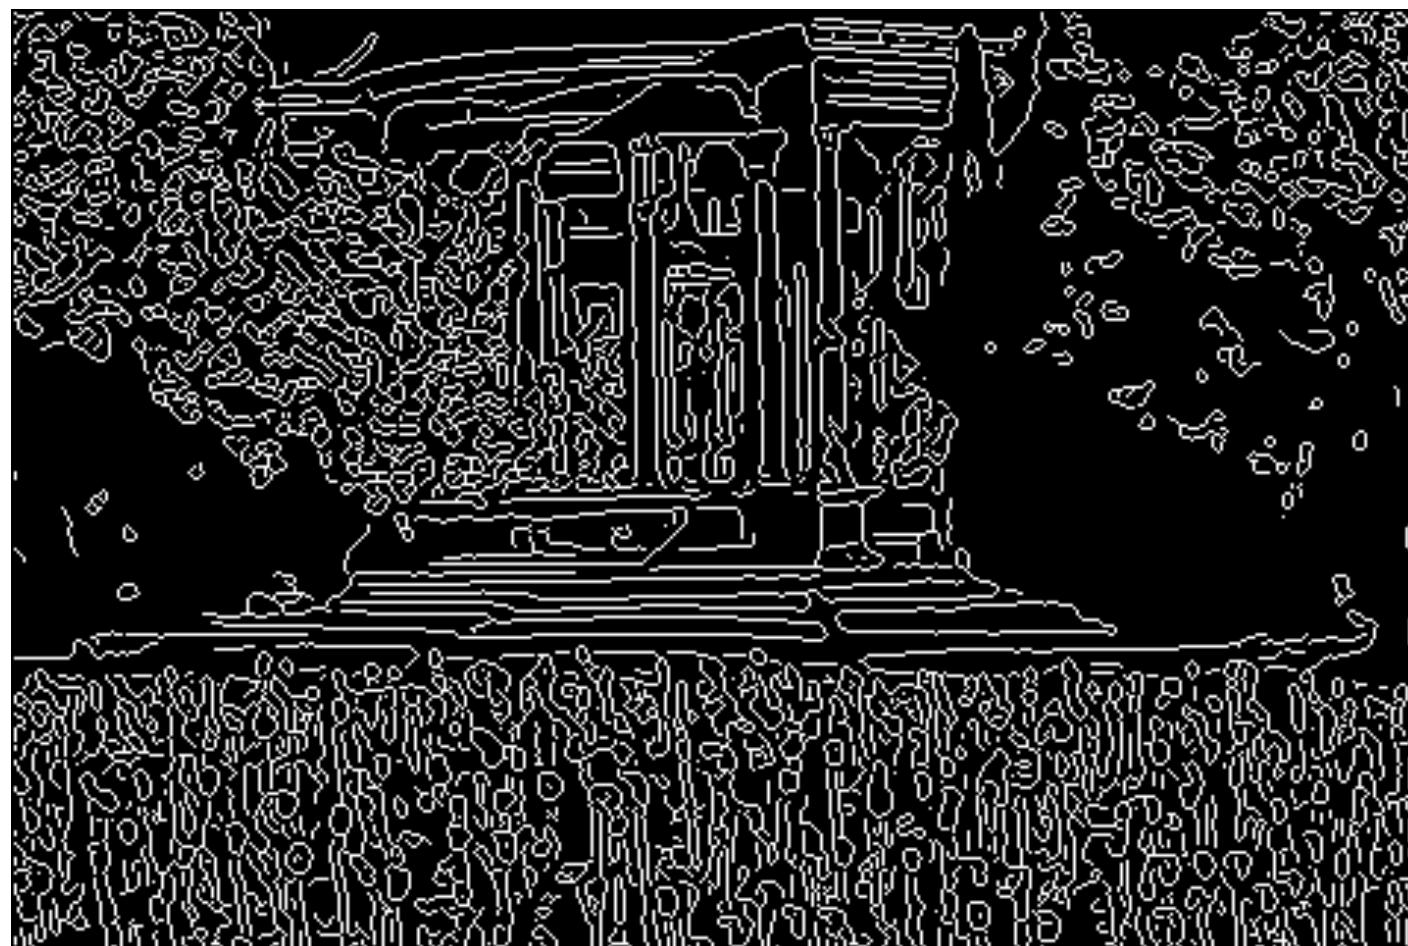
- Which can be used for convolutions.

Edge detection

- Seeking human performance

Edge detection

- Seeking human performance
 - Reason



Edge detection

- Seeking human performance
 - Reason



Edge detection

- Seeking human performance



Source: Antonio Torralba slides

Gaussian gradient
boundaries

Human boundaries

Edge detection

- Seeking human performance



Gaussian gradient
boundaries

Source: Antonio Torralba slides

Human boundaries

138

Edge detection

- Seeking human performance



Gaussian gradient
boundaries



Human boundaries

Edge detection

- Seeking human performance: Example



Source: Jitendra Malik (2004)

Edge detection

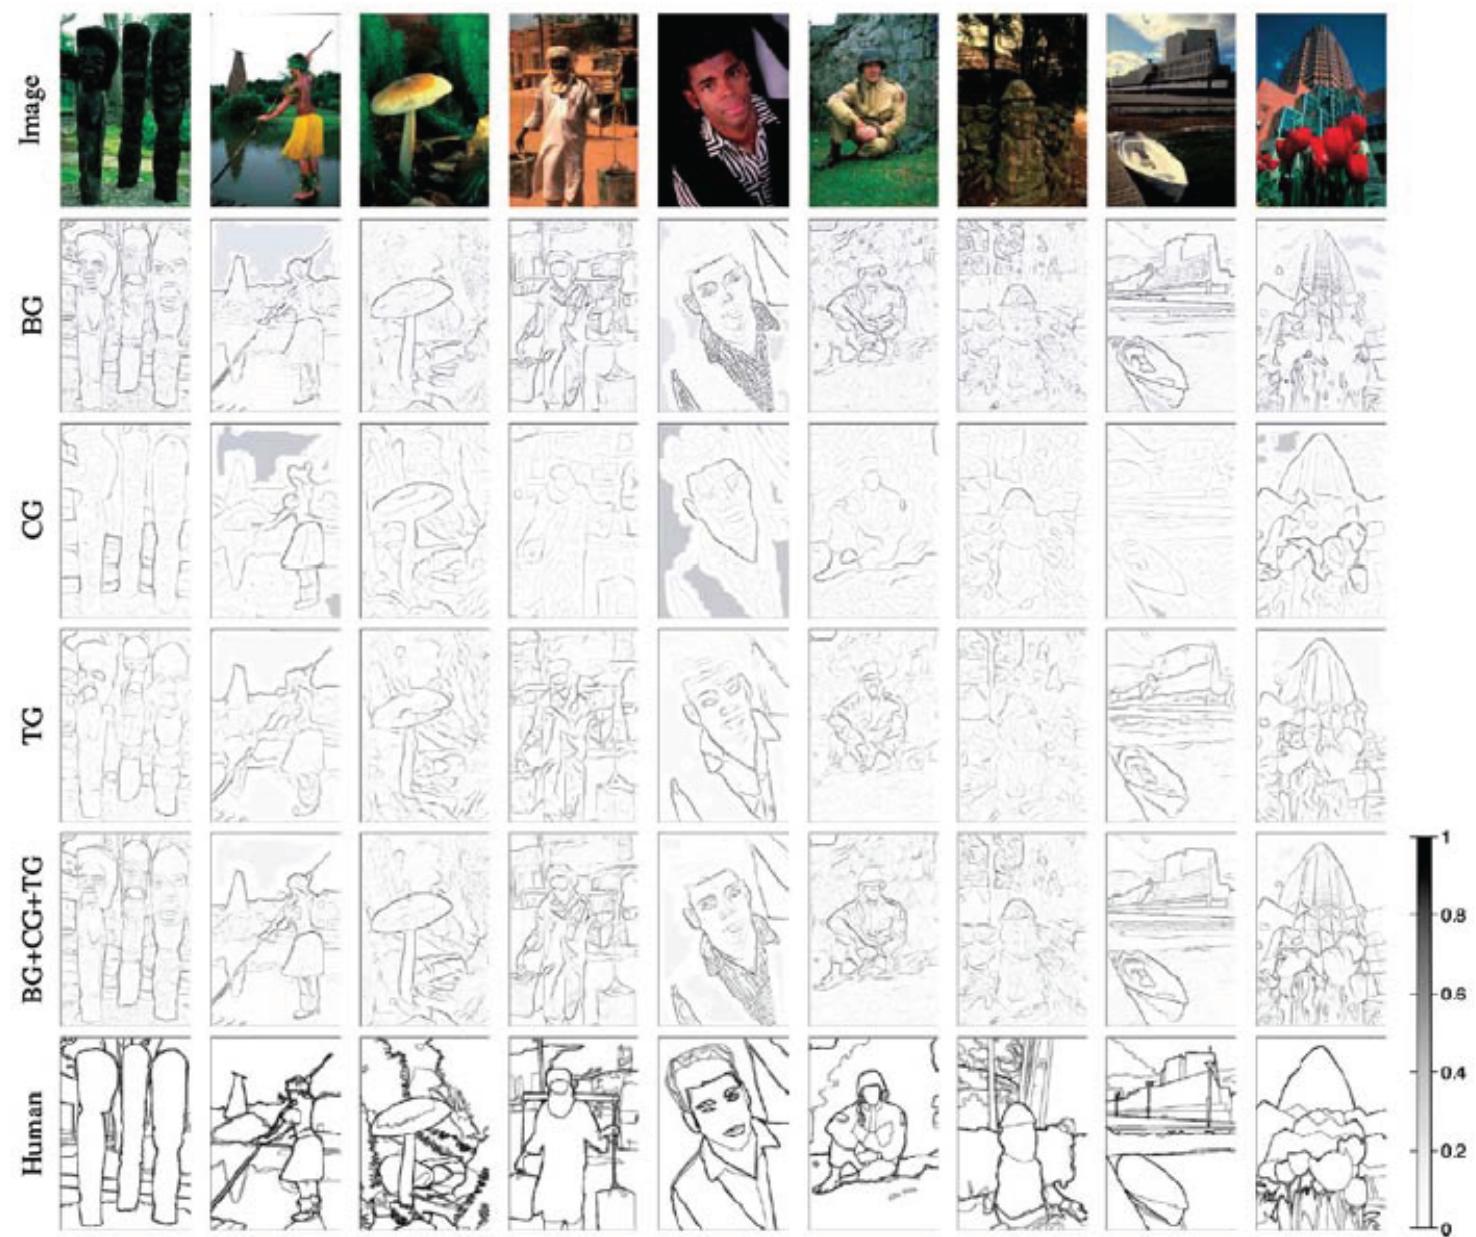


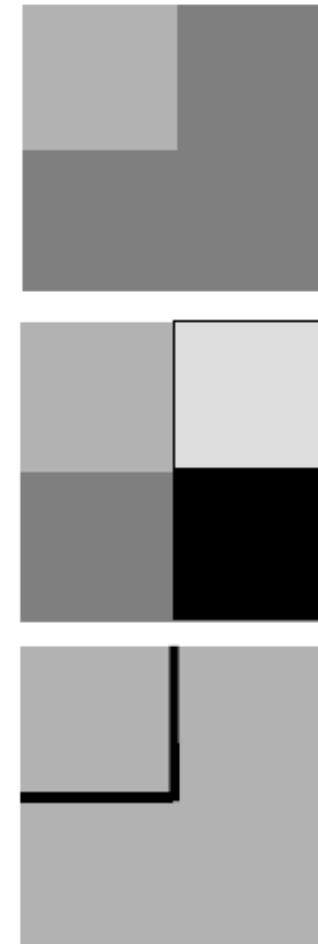
Figure 4.33 Combined brightness, color, texture boundary detector (Martin, Fowlkes, and Malik 2004) © 2004 IEEE. Successive rows show the outputs of the brightness gradient (BG), color gradient (CG), texture gradient (TG), and combined (BG+CG+TG) detectors. The final row shows human-labeled boundaries derived from a database of hand-segmented images (Martin, Fowlkes, Tal *et al.* 2001).

Outline

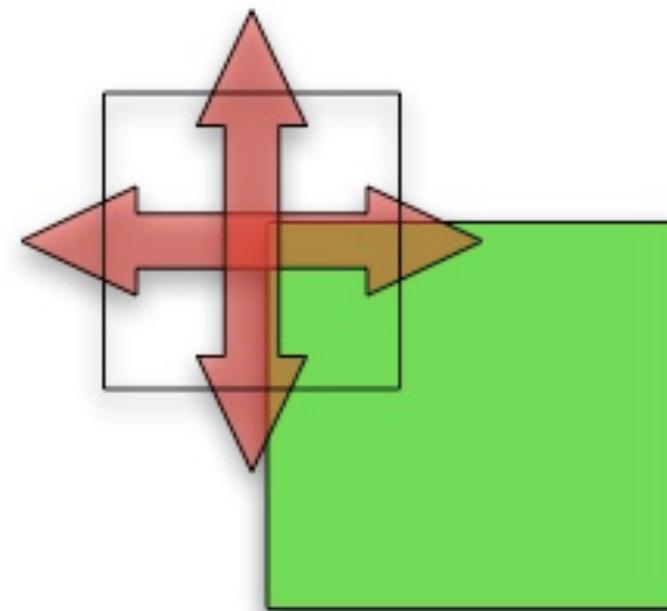
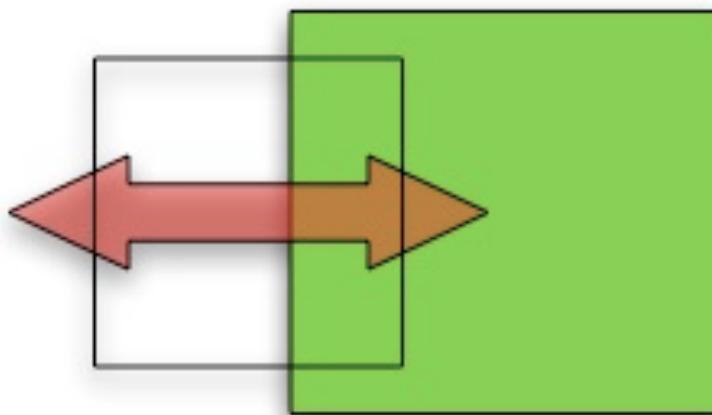
- Introduction
- Point and patches
- Edges
- Corners

Corners

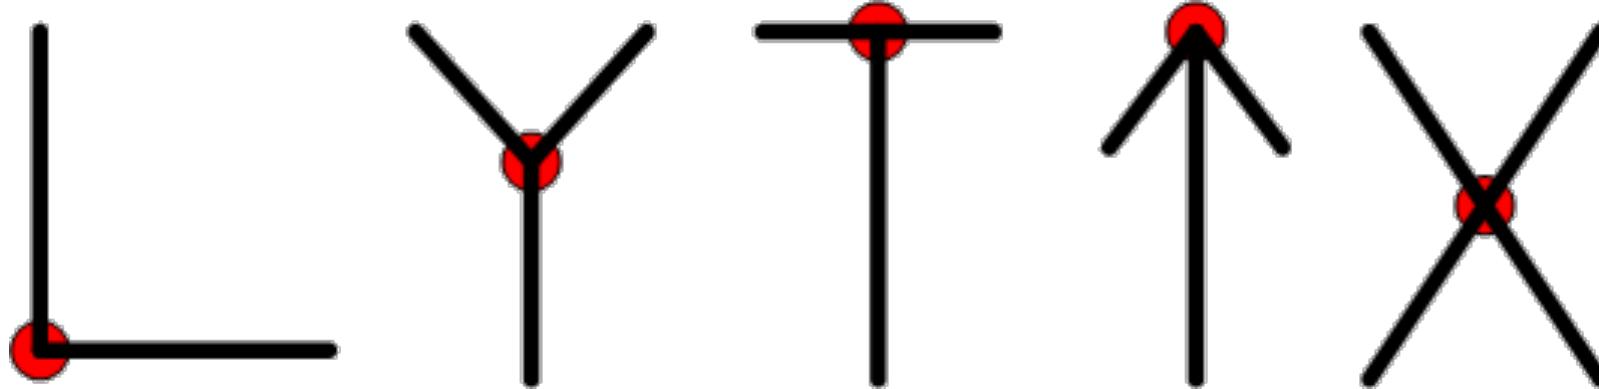
- What is a corner
 - A corner is an image location where two distinct image orientations occur in a local region.
 - Physically, image corners tend to arise for similar reasons as edges (e.g., changes of reflectance, surface orientation).
 - Corners are of interest for two main reasons
 1. Corners provide constrain 2 degrees of freedom in a pattern's location.
 2. Corners tend to persist across changes in viewpoint.



Corners vs Edges



Corners types



Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - Given that we are concerned with local measures of orientation, one approach is to calculate the local spatial derivatives E_x and E_y , using subscript notation for partial derivatives.
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = \begin{pmatrix} \sum E_x^2 & \sum E_x E_y \\ \sum E_x E_y & \sum E_y^2 \end{pmatrix}$$



Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = G(\sigma) * (\nabla I)(\nabla I)^T = G(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Because this matrix is symmetric, this can be simplified into

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = G(\sigma) * (\nabla I)(\nabla I)^T = G(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Because this matrix is symmetric, this can be simplified into

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- To provide three cases
 - If the region of interest is perfectly uniform, then the gradients are identically zero $\lambda_1 = \lambda_2 = 0$



Corners

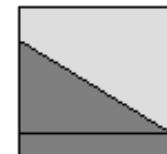
- We can use also the gradient (Harris & Stephens, 1988)
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = G(\sigma) * (\nabla I)(\nabla I)^T = G(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Because this matrix is symmetric, this can be simplified into

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- To provide three cases
 - If the region contains an ideal step edge, then there is only one gradient direction $\lambda_1 > 0, \lambda_2 = 0$



Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = G(\sigma) * (\nabla I)(\nabla I)^T = G(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Because this matrix is symmetric, this can be simplified into

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- To provide three cases
 - If the region contains two orientations, then there are two gradient directions $\lambda_1 > 0, \lambda_2 > 0$



Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - We choose to accumulate these measures over a neighborhood via summation and construct the matrix

$$M = G(\sigma) * (\nabla I)(\nabla I)^T = G(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Because this matrix is symmetric, this can be simplified into

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- Magnitude of λ_1 and λ_2 are called eigenvalues and capture the edge strength
- The direction is captured by the eigenvectors

Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - Because this matrix is symmetric, this can be simplified into
 - Magnitude of λ_1 and λ_2 are called eigenvalues and capture the edge strength
 - The direction is captured by the eigenvectors

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - Because this matrix is symmetric, this can be simplified into
 - Magnitude of λ_1 and λ_2 are called eigenvalues and capture the edge strength
 - The direction is captured by the eigenvectors

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

$$\lambda_1 + \lambda_2 = \text{Tr}(M) = M_{1,1} + M_{2,2}$$

$$\lambda_1 \lambda_2 = \text{Det}(M) = M_{1,1}M_{2,2} - M_{1,2}M_{2,1}$$



Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - Because this matrix is symmetric, this can be simplified into
 - Magnitude of λ_1 and λ_2 are called eigenvalues and capture the edge strength
 - The direction is captured by the eigenvectors

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

$$\lambda_1 + \lambda_2 = \text{Tr}(M) = M_{1,1} + M_{2,2}$$

$$\lambda_1 \lambda_2 = \text{Det}(M) = M_{1,1}M_{2,2} - M_{1,2}M_{2,1}$$

- The corner response is $R = \text{Det}(M) - k\text{Tr}^2(M)$

Corners

- We can use also the gradient (Harris & Stephens, 1988)
 - Because this matrix is symmetric, this can be simplified into
 - Magnitude of λ_1 and λ_2 are called eigenvalues and capture the edge strength
 - The direction is captured by the eigenvectors

$$\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

$$\lambda_1 + \lambda_2 = \text{Tr}(M) = M_{1,1} + M_{2,2}$$

$$\lambda_1 \lambda_2 = \text{Det}(M) = M_{1,1} M_{2,2} - M_{1,2} M_{2,1}$$

- The corner response is $R = \text{Det}(M) - k \text{Tr}^2(M)$
- Corners are all local maxima of R with $R > 0$ and $\text{Tr} > t$

Corners

- The discrete case

	1	-2	1
$\frac{1}{4}$	-2	4	-2
	1	-2	1

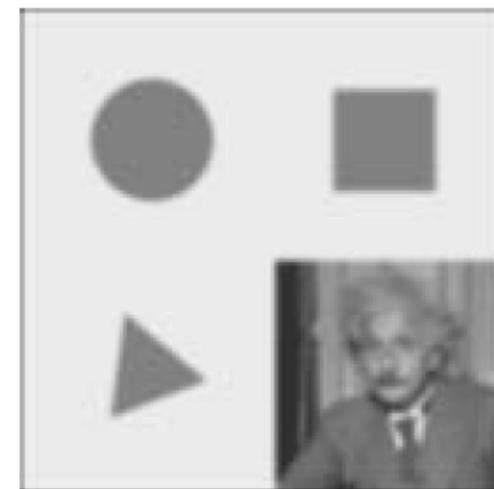
$\frac{1}{2}$	1	-2	1
---------------	---	----	---

Corners

- The discrete case

$$\frac{1}{4} \begin{array}{|c|c|c|}\hline 1 & -2 & 1 \\ \hline -2 & 4 & -2 \\ \hline 1 & -2 & 1 \\ \hline \end{array}$$

$$\frac{1}{2} \begin{array}{|c|c|c|}\hline 1 & -2 & 1 \\ \hline \end{array}$$



Summary

- Introduction
- Point and patches
- Edges
- Corners



universität
innsbruck



OpenCV. Feature extraction

Assoz.Prof. Antonio Rodríguez-Sánchez, PhD.

Sobel operator

- Remember from steerable filters
 - Consider a 2D circularly symmetric Gaussian function

$$G(x, y) = e^{-(x^2 + y^2)}$$

- The first x derivative is
- That same function rotated 90 degrees

$$G_1^{0^\circ} = \frac{\partial}{\partial x} e^{-(x^2 + y^2)} = -2xe^{-(x^2 + y^2)}$$

Sobel operator

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * I \quad G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * I$$

Sobel operator

$$G_x = \begin{bmatrix} -3 & 0 & +3 \\ -10 & 0 & +10 \\ -3 & 0 & +3 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -3 & -10 & -3 \\ 0 & 0 & 0 \\ +3 & +10 & +3 \end{bmatrix}$$

Sobel operator

```
#include "opencv2/imgproc/imgproc.hpp"
#include "opencv2/highgui/highgui.hpp"
#include <stdlib.h>
#include <stdio.h>

using namespace cv;

/** @function main */
int main( int argc, char** argv )
{

    Mat src, src_gray;
    Mat grad;
    char* window_name = "Sobel Demo - Simple Edge Detector";
    int scale = 1;
    int delta = 0;
    int ddepth = CV_16S;
    int c;

    /// Load an image
    src = imread( argv[1] );

    if( !src.data )
    { return -1; }
```

Sobel operator

```
if( !src.data )
{ return -1; }

GaussianBlur( src, src, Size(3,3), 0, 0, BORDER_DEFAULT );

/// Convert it to gray
cvtColor( src, src_gray, CV_BGR2GRAY );

/// Create window
namedWindow( window_name, CV_WINDOW_AUTOSIZE );

/// Generate grad_x and grad_y
Mat grad_x, grad_y;
Mat abs_grad_x, abs_grad_y;

/// Gradient X
//Scharr( src_gray, grad_x, ddepth, 1, 0, scale, delta, BORDER_DEFAULT );
Sobel( src_gray, grad_x, ddepth, 1, 0, 3, scale, delta, BORDER_DEFAULT );
convertScaleAbs( grad_x, abs_grad_x );
```

Sobel operator

```
if( !src.data )
{ return -1; }

GaussianBlur( src, src, Size(3,3), 0, 0, BORDER_DEFAULT );

/// Convert it to gray
cvtColor( src, src_gray, CV_BGR2GRAY );

/// Create window
namedWindow( window_name, CV_WINDOW_AUTOSIZE );

/// Generate grad_x and grad_y
Mat grad_x, grad_y;
Mat abs_grad_x, abs_grad_y;

/// Gradient X
//Scharr( src_gray, grad_x, ddepth, 1, 0, scale, delta, BORDER_DEFAULT );
Sobel( src_gray, grad_x, ddepth, 1, 0, 3, scale, delta, BORDER_DEFAULT );
convertScaleAbs( grad_x, abs_grad_x );
```

Sobel operator

```
/// Gradient X
//Schar( src_gray, grad_x, ddepth, 1, 0, scale, delta, BORDER_DEFAULT );
Sobel( src_gray, grad_x, ddepth, 1, 0, 3, scale, delta, BORDER_DEFAULT );
convertScaleAbs( grad_x, abs_grad_x );

/// Gradient Y
//Schar( src_gray, grad_y, ddepth, 0, 1, scale, delta, BORDER_DEFAULT );
Sobel( src_gray, grad_y, ddepth, 0, 1, 3, scale, delta, BORDER_DEFAULT );
convertScaleAbs( grad_y, abs_grad_y );

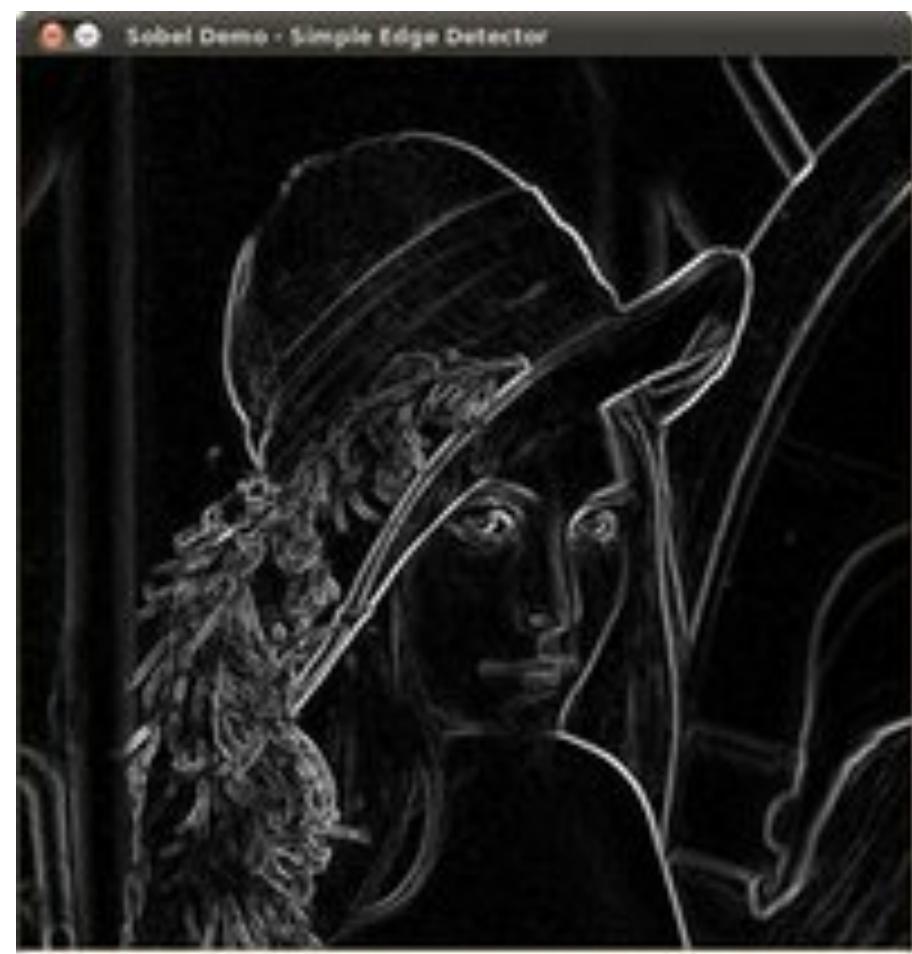
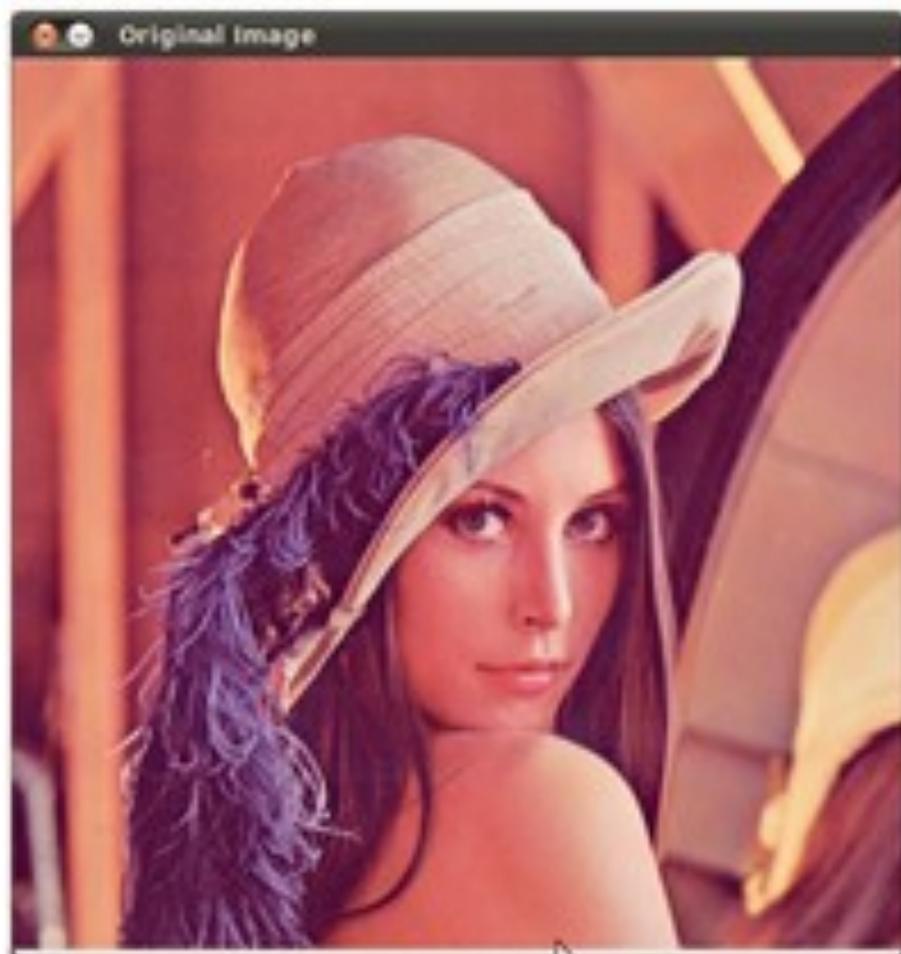
/// Total Gradient (approximate)
addWeighted( abs_grad_x, 0.5, abs_grad_y, 0.5, 0, grad );

imshow( window_name, grad );

waitKey(0);

return 0;
}
```

Sobel operator



Laplace operator

- Band-pass filtering
 - Filter low and high frequencies
 - Remember the Gaussian equation

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}},$$

- If we take the Laplacian

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Note: there is a typo in Szeliski book

- We obtain the Laplacian of Gaussian filter

$$\nabla^2 G(x, y; \sigma) = \left(\frac{x^2 + y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) G(x, y; \sigma),$$

Laplace operator

```
#include "opencv2/imgproc/imgproc.hpp"
#include "opencv2/highgui/highgui.hpp"
#include <stdlib.h>
#include <stdio.h>

using namespace cv;

/** @function main */
int main( int argc, char** argv )
{
    Mat src, src_gray, dst;
    int kernel_size = 3;
    int scale = 1;
    int delta = 0;
    int ddepth = CV_16S;
    char* window_name = "Laplace Demo";

    int c;

    /// Load an image
    src = imread( argv[1] );

    if( !src.data )
        { return -1; }
```

Laplace operator

```
if( !src.data )
{ return -1; }

/// Remove noise by blurring with a Gaussian filter
GaussianBlur( src, src, Size(3,3), 0, 0, BORDER_DEFAULT );

/// Convert the image to grayscale
cvtColor( src, src_gray, CV_BGR2GRAY );

/// Create window
namedWindow( window_name, CV_WINDOW_AUTOSIZE );

/// Apply Laplace function
Mat abs_dst;

Laplacian( src_gray, dst, ddepth, kernel_size, scale, delta, BORDER_DEFAULT );
convertScaleAbs( dst, abs_dst );

/// Show what you got
imshow( window_name, abs_dst );

waitKey(0);

return 0;
```

Laplace operator

```
if( !src.data )
{ return -1; }

/// Remove noise by blurring with a Gaussian filter
GaussianBlur( src, src, Size(3,3), 0, 0, BORDER_DEFAULT );

/// Convert the image to grayscale
cvtColor( src, src_gray, CV_BGR2GRAY );

/// Create window
namedWindow( window_name, CV_WINDOW_AUTOSIZE );

/// Apply Laplace function
Mat abs_dst;

Laplacian( src_gray, dst, ddepth, kernel_size, scale, delta, BORDER_DEFAULT );
convertScaleAbs( dst, abs_dst );

/// Show what you got
imshow( window_name, abs_dst );

waitKey(0);

return 0;
```

Laplace operator



Harris corner detector

```
#include "opencv2/highgui/highgui.hpp"
#include "opencv2/imgproc/imgproc.hpp"
#include <iostream>
#include <stdio.h>
#include <stdlib.h>

using namespace cv;
using namespace std;

/// Global variables
Mat src, src_gray;
int thresh = 200;
int max_thresh = 255;

char* source_window = "Source image";
char* corners_window = "Corners detected";

/// Function header
void cornerHarris_demo( int, void* );
```

Harris corner detector

```
// Function header
void cornerHarris_demo( int, void* );

/** @function main */
int main( int argc, char** argv )
{
    // Load source image and convert it to gray
    src = imread( argv[1], 1 );
    cvtColor( src, src_gray, CV_BGR2GRAY );

    // Create a window and a trackbar
    namedWindow( source_window, CV_WINDOW_AUTOSIZE );
    createTrackbar( "Threshold: ", source_window, &thresh, max_thresh, cornerHarris_demo );
    imshow( source_window, src );

    cornerHarris_demo( 0, 0 );

    waitKey(0);
    return(0);
}
```

Harris corner detector

```
/** @function cornerHarris_demo */
void cornerHarris_demo( int, void* )
{
    Mat dst, dst_norm, dst_norm_scaled;
    dst = Mat::zeros( src.size(), CV_32FC1 );

    /// Detector parameters
    int blockSize = 2;
    int apertureSize = 3;
    double k = 0.04;

    /// Detecting corners
    cornerHarris( src_gray, dst, blockSize, apertureSize, k, BORDER_DEFAULT );

    /// Normalizing
    normalize( dst, dst_norm, 0, 255, NORM_MINMAX, CV_32FC1, Mat() );
    convertScaleAbs( dst_norm, dst_norm_scaled );

    /// Drawing a circle around corners
    for( int j = 0; j < dst_norm.rows ; j++ )
    { for( int i = 0; i < dst_norm.cols; i++ )
        {
            if( (int) dst_norm.at<float>(j,i) > thresh )
            {
                circle( dst_norm_scaled, Point( i, j ), 5, Scalar(0), 2, 8, 0 );
            }
        }
    }
}
```

Harris corner detector

```
/** @function cornerHarris_demo */
void cornerHarris_demo( int, void* )
{

    Mat dst, dst_norm, dst_norm_scaled;
    dst = Mat::zeros( src.size(), CV_32FC1 );

    /// Detector parameters
    int blockSize = 2;
    int apertureSize = 3;
    double k = 0.04;

    /// Detecting corners
    cornerHarris( src_gray, dst, blockSize, apertureSize, k, BORDER_DEFAULT );

    /// Normalizing
    normalize( dst, dst_norm, 0, 255, NORM_MINMAX, CV_32FC1, Mat() );
    convertScaleAbs( dst_norm, dst_norm_scaled );

    /// Drawing a circle around corners
    for( int j = 0; j < dst_norm.rows ; j++ )
    { for( int i = 0; i < dst_norm.cols; i++ )
        {
            if( (int) dst_norm.at<float>(j,i) > thresh )
            {
                circle( dst_norm_scaled, Point( i, j ), 5, Scalar(0), 2, 8, 0 );
            }
        }
    }
}
```

Harris corner detector

```
/** @function cornerHarris_demo */
void cornerHarris_demo( int, void* )
{

    Mat dst, dst_norm, dst_norm_scaled;
    dst = Mat::zeros( src.size(), CV_32FC1 );

    /// Detector parameters
    int blockSize = 2;
    int apertureSize = 3;
    double k = 0.04;

    /// Detecting corners
    cornerHarris( src_gray, dst, blockSize, apertureSize, k, BORDER_DEFAULT );

    /// Normalizing
    normalize( dst, dst_norm, 0, 255, NORM_MINMAX, CV_32FC1, Mat() );
    convertScaleAbs( dst_norm, dst_norm_scaled );

    /// Drawing a circle around corners
    for( int j = 0; j < dst_norm.rows ; j++ )
    { for( int i = 0; i < dst_norm.cols; i++ )
        {
            if( (int) dst_norm.at<float>(j,i) > thresh )
            {
                circle( dst_norm_scaled, Point( i, j ), 5, Scalar(0), 2, 8, 0 );
            }
        }
    }
}
```

Harris corner detector

