



universität
innsbruck



Lecture 8. Stereopsis II

703142. Computer Vision

Assoz.Prof. Antonio Rodríguez-Sánchez, PhD.

Source: Hartley and Zisserman, 2004

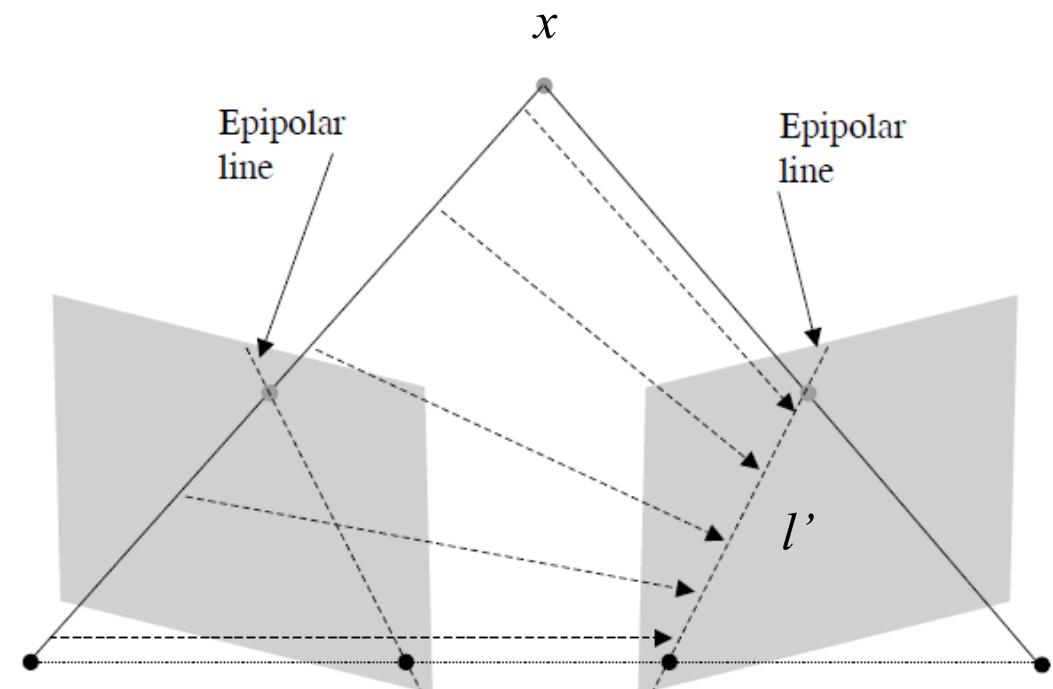
Outline

- Introduction
- 3D shapes from 2D images
- Stereo vision
- Correspondence
- Epipolar geometry
- The fundamental matrix
- The essential matrix
- RANSAC
- 3D reconstruction

The fundamental matrix

- Algebraic representation of epipolar geometry
- Extracted from the mapping between a point and its epipolar line
- There is a map

$$x \mapsto l'$$



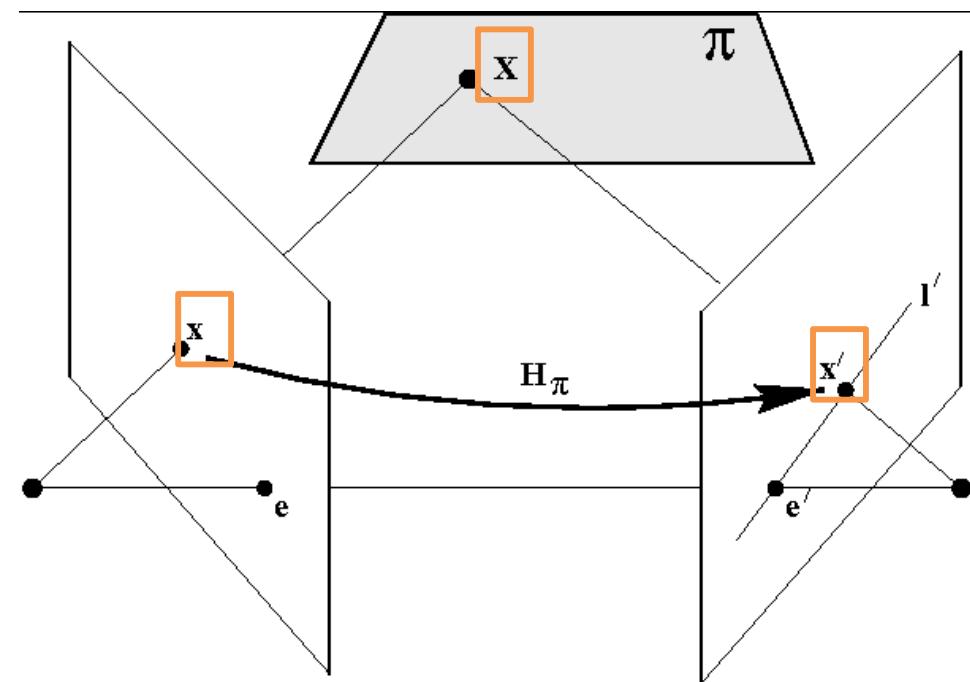
The fundamental matrix

- There is a map

$$x \mapsto l'$$

- Consider point \mathbf{X}

- Projection is \mathbf{x} on the left image plane
 - Projection is \mathbf{x}' on the right image plane



The fundamental matrix

- There is a map

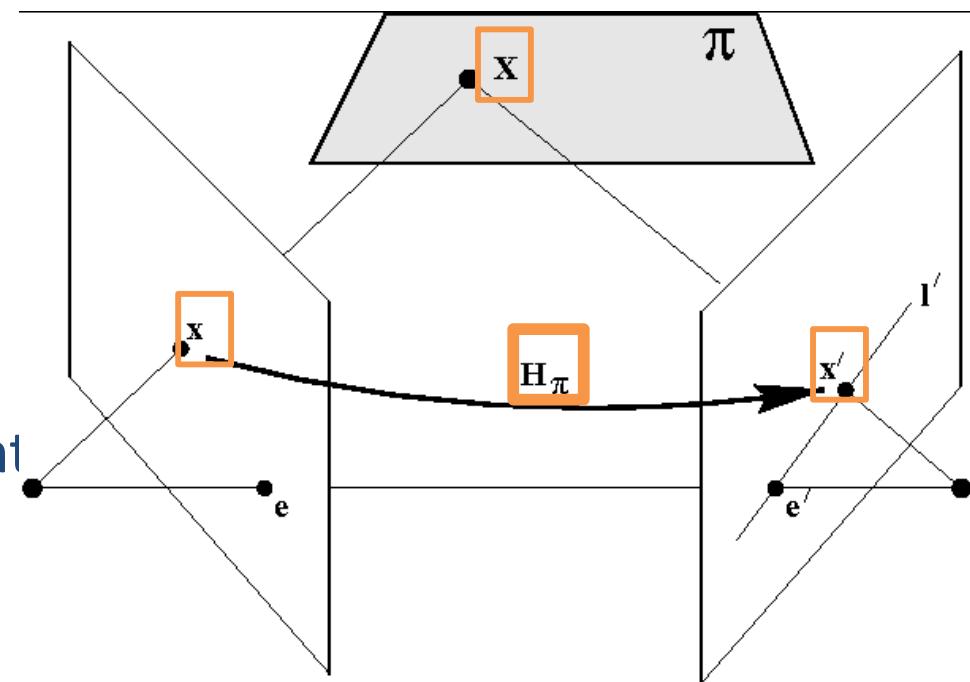
$$x \mapsto l'$$

- Consider point \mathbf{X}

- Projection is \mathbf{x} on the left image plane

- Projection is \mathbf{x}'_i on the right image plane

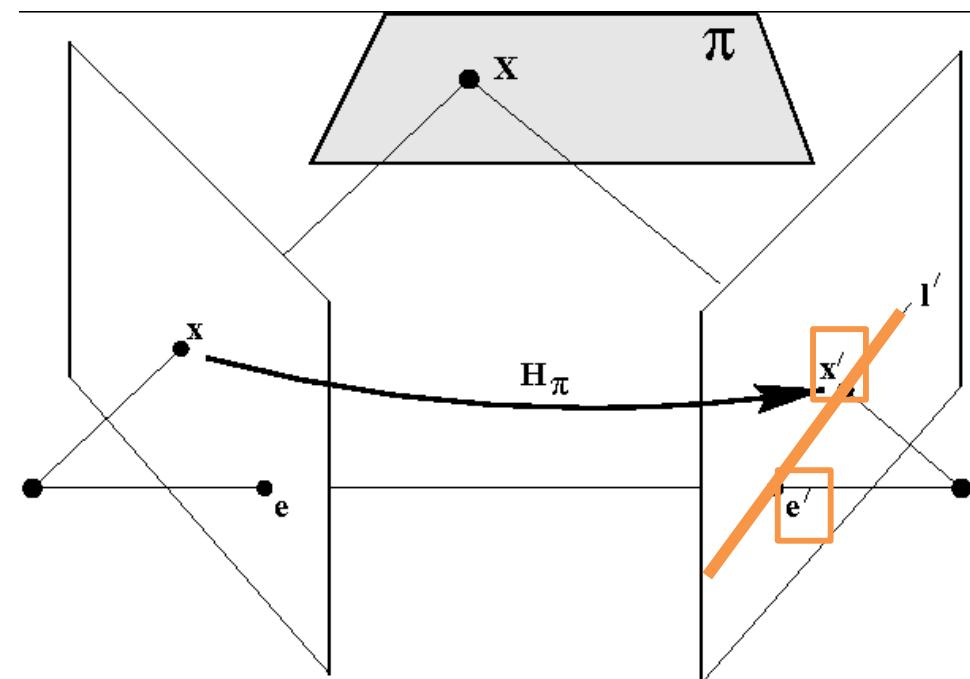
- Then there is a 2D homography \mathbf{H}_π mapping each \mathbf{x} into \mathbf{x}'_i



The fundamental matrix

- There is a 2D homography H_π mapping each \mathbf{x} into \mathbf{x}'_i
- The epipolar line is the line that goes through \mathbf{x}' and \mathbf{e}'

$$l' = e' \times x'$$



crossprod



Note: Skew symmetric matrices and cross products

- A vector $\mathbf{a}=(a_1, a_2, a_3)^T$ can be written as a skew matrix as

$$[\mathbf{a}]_{\times} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}$$



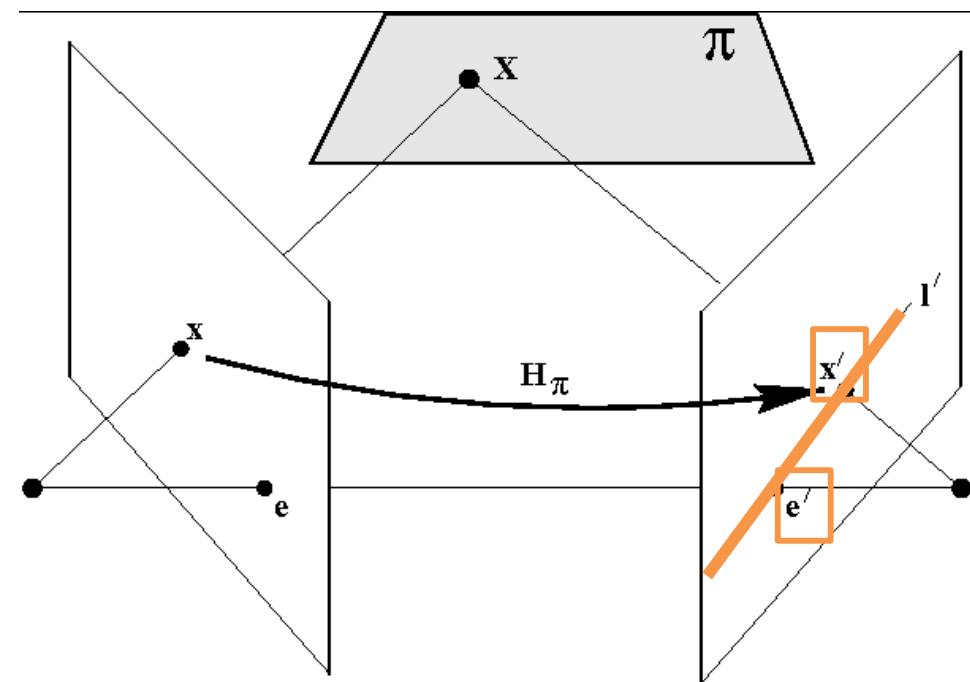
- The cross product is related to skew matrices according to

$$\mathbf{a} \times \mathbf{b} = [\mathbf{a}]_{\times} \mathbf{b} = (\mathbf{a}^T [\mathbf{b}])^T$$

The fundamental matrix

- There is a 2D homography H_π mapping each \mathbf{x} into \mathbf{x}'_i
- The epipolar line is the line that goes through \mathbf{x}' and \mathbf{e}'

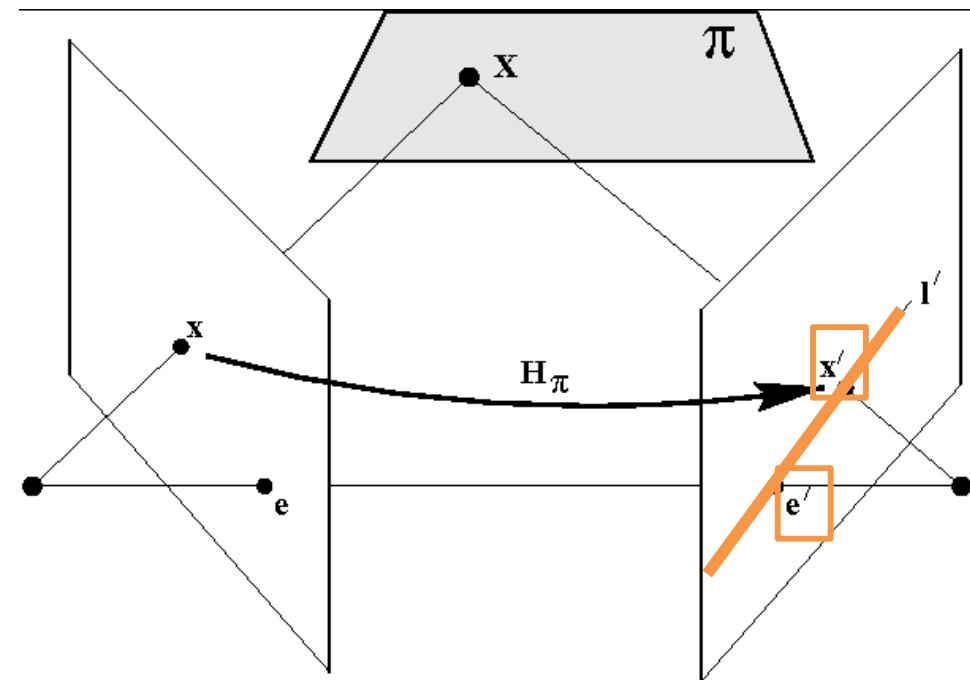
$$l' \equiv e' \times x'$$



The fundamental matrix

- There is a 2D homography H_π mapping each \mathbf{x} into \mathbf{x}'
- The epipolar line is the line that goes through \mathbf{x}' and \mathbf{e}'

$$l' = \mathbf{e}' \times \mathbf{x}'$$



- But $\mathbf{x}' = H_\pi \mathbf{x}$, so we have

$$l' = \mathbf{e}' \times \mathbf{x}' = [e']_x H_\pi x = Fx$$

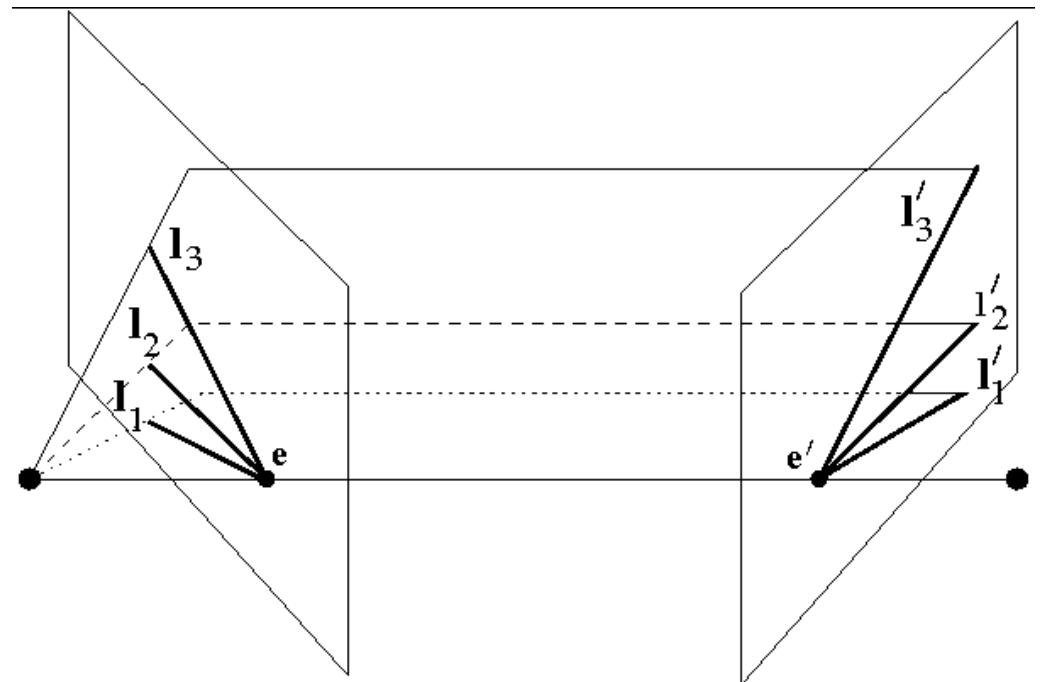
F .. Fundamental Mat

The fundamental matrix

- The fundamental matrix is

$$F = [e']_x H_\pi$$

- F represents a mapping from the 2-dimensional projective plane \mathbf{P}^2 of the first image to the pencil of epipolar lines through e'
 - That is, it is a mapping from 2D into 1D



The fundamental matrix

- The fundamental matrix is $F \in [e']_x H_\pi$
- For all corresponding points $x \leftrightarrow x'$, The fundamental matrix satisfies

$$x'^T F x = 0$$

The fundamental matrix

- The fundamental matrix is $F = [e']_x H_\pi$
- For all corresponding points $x \leftrightarrow x'$, The fundamental matrix satisfies

$$x'^T F x \neq 0$$

- The essential matrix is the specific case where it is assumed that cameras are calibrated
 - Thus, the fundamental matrix is the generalized essential matrix

Outline

- Introduction
- 3D shapes from 2D images
- Stereo vision
- Correspondence
- Epipolar geometry
- The fundamental matrix
- The essential matrix
- RANSAC
- 3D reconstruction

Camera models

- Intrinsic parameters
 - Relate camera's coordinate system to idealized coordinate system
 - That is, link the pixel coordinates of an image point with corresponding coordinates in the camera reference frame.
- Extrinsic parameters
 - Relate the camera's coordinate system to a fixed world coordinate system and specify its position and orientation in space
- The problem of estimating the intrinsic and extrinsic parameters of a camera is known as geometric camera calibration

Geometry of image formation

- Intrinsic parameters
 - Camera to pixel transformation
 - Typically given by 2 sets of parameters
 1. The focal length, f , serving to capture the (perspective) projection.
 2. The pixel coordinates of the principle point (image center), (o_x, o_y) , and the effective pixel horizontal and vertical dimensions, (s_x, s_y) , serving to capture the transformation between camera frame coordinates and pixel coordinates.
 - We already have considered how to incorporate the focal length.
 - Letting (x_i, y_i) be the pixel coordinates, we incorporate the second set of parameters via

$$x = -(x_i - o_x)s_x$$

$$y = -(y_i - o_y)s_y$$

- There are 5 total intrinsic parameters: (f, o_x, o_y, s_x, s_y)

Note

- Extrinsic parameters
 - Camera to world transformation
 - Typically given via two sets of parameters
 1. A 3D translation vector, \mathbf{T} , describing the relative positions of the two reference frames.
 2. A 3x3 rotation matrix, \mathbf{R} , that brings the corresponding axes of the two frames into alignment.
 - Letting \mathbf{P}_c and \mathbf{P}_w be the camera and world coordinates of the same point, we write:

$$\mathbf{P}_c = \mathbf{R}(\mathbf{P}_w - \mathbf{T})$$

- By definition, \mathbf{R} , is completely specified by 3 parameters (e.g., rotation about each of the coordinate axes); so, there are 6 extrinsic parameters in total (3 for \mathbf{T} ; 3 for \mathbf{R}).

Geometry of image formation

- Fundamental projections: **Perspective**

- The length of \mathbf{P} is $P \equiv -Z \sec a \equiv -(P \cdot \hat{\mathbf{Z}}) \sec a$

- The length of \mathbf{p} is $p \equiv f \sec a$

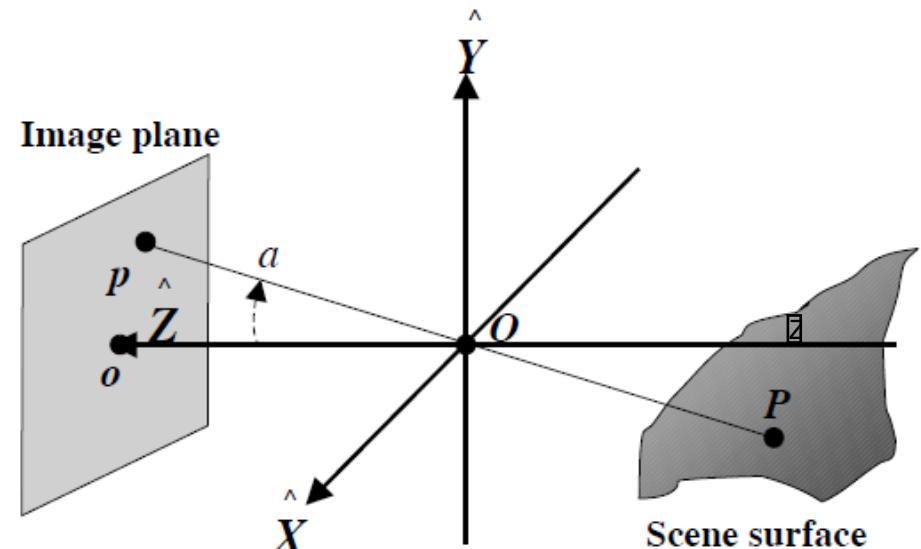
- So $\frac{1}{f} p = \frac{1}{P \cdot \hat{\mathbf{Z}}} P$

- In component form

$$\frac{x}{f} \equiv \frac{X}{Z}, \frac{y}{f} \equiv \frac{Y}{Z}$$

- Or

$$x \equiv f \frac{X}{Z}, y \equiv f \frac{Y}{Z}$$



Geometry of image formation

- World-image transformation
 - Recall the camera frame expression of perspective projection

$$x \equiv f \frac{X}{Z}, y \equiv f \frac{Y}{Z}$$

- We substitute the intrinsic parameterization on the lhs and the extrinsic parameterization on the rhs to find that (with R_i the i-th row of \mathbf{R})

$$(x_i - o_x)s_x \equiv f \frac{R_1^T(P_w + T)}{R_3^T(P_w + T)}$$

$$\boxed{P_c = R * (P_w - T)}$$

$$(y_i - o_y)s_y \equiv f \frac{R_2^T(P_w + T)}{R_3^T(P_w + T)}$$

- Which relates the 3D coordinates of a point in the world frame to its corresponding image coordinates.

Geometry of image formation

- World-image transformation
 - Rewrite in matrix formulation

$$-(x_i - o_x)s_x = f \frac{R_1^T(P_w - T)}{R_3^T(P_w - T)}$$

$$-(y_i - o_y)s_y = f \frac{R_2^T(P_w - T)}{R_3^T(P_w - T)}$$

- Define **matrices** that encapsulate intrinsic, \mathbf{M}_{int} , and extrinsic, \mathbf{M}_{ext} parameters

$$\mathbf{M}_{\text{int}} = \begin{pmatrix} -f/s_x & 0 & o_x \\ 0 & -f/s_y & o_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{M}_{\text{ext}} = \begin{pmatrix} \bullet & \bullet & \bullet & -\mathbf{R}_1^T T \\ \bullet & \mathbf{R} & \bullet & -\mathbf{R}_2^T T \\ \bullet & \bullet & \bullet & -\mathbf{R}_3^T T \end{pmatrix}$$

Geometry of image formation

- World-image transformation

$$\mathbf{M}_{int} = \begin{pmatrix} -f/s_x & 0 & o_x \\ 0 & -f/s_y & o_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{M}_{ext} = \begin{pmatrix} \bullet & \bullet & \bullet & -\mathbf{R}_1^T T \\ \bullet & \mathbf{R} & \bullet & -\mathbf{R}_2^T T \\ \bullet & \bullet & \bullet & -\mathbf{R}_3^T T \end{pmatrix}$$

- And concatenate them to write

$$\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \mathbf{M}_{int} \mathbf{M}_{ext} \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$

Remark: $(x_i, y_i) = (x_1/x_3, x_2/x_3)$

The essential matrix

- The fundamental matrix is $F \triangleq [e']_x H_\pi$
- For all corresponding points $x \leftrightarrow x'$, The fundamental matrix satisfies
$$x'^T F x = 0$$
- The essential matrix is the specific case where it is assumed that the cameras are normalized
 - Thus, the fundamental matrix is the generalized essential matrix

The essential matrix

- The fundamental matrix is $F = [e']_x H_\pi$
- For all corresponding points $\textcolor{brown}{x} \leftrightarrow \textcolor{brown}{x}'$, The fundamental matrix satisfies

$$x'^T F x = 0$$

- The essential matrix is the specific case where it is assumed that the cameras are normalized
 - Thus, the fundamental matrix is the generalized essential matrix
 - We assume that the world coordinate system is centered on the first camera, extrinsic parameters are $\{\mathbf{I}, \mathbf{0}\}$. $\mathbf{M}_{int1} = \mathbf{I}$
 - The second camera may be in any general position $\{\Omega, \tau\}$.

$$\mathbf{M}_{int2} = \mathbf{I}$$

The essential matrix

- For all corresponding points in homogeneous coordinates. The essential matrix satisfies

$$\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{x}}_1 = 0$$

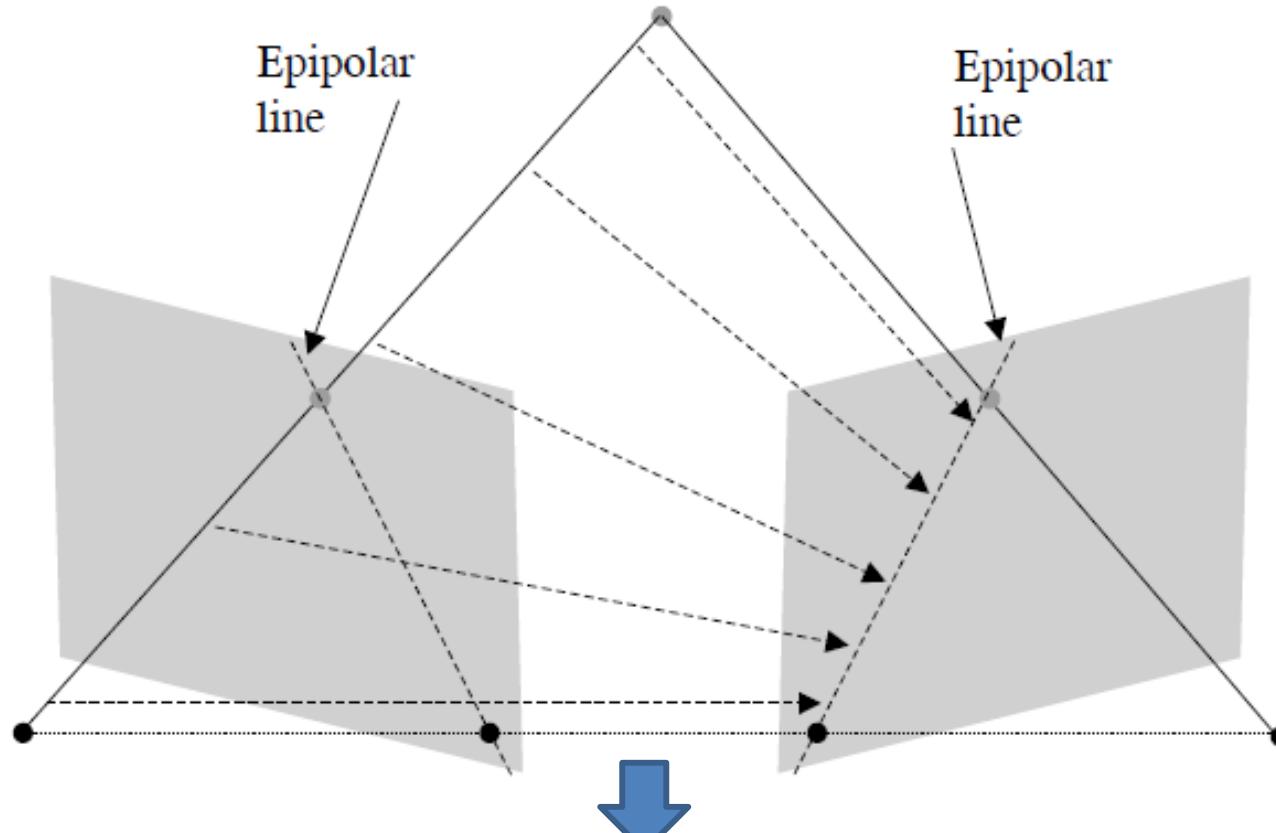


- The essential matrix is the specific case where it is assumed that the cameras are normalized
 - Thus, the fundamental matrix is the generalized essential matrix
 - We assume that the world coordinate system is centered on the first camera, extrinsic parameters are $\{\mathbf{I}, \mathbf{0}\}$. $\mathbf{M}_{int1} = \mathbf{I}$
 - The second camera may be in any general position $\{\boldsymbol{\Omega}, \boldsymbol{\tau}\}$.

$$\mathbf{M}_{int2} = \mathbf{I}$$

Epipolar geometry

- Given a stereo pair of cameras and a point in 3D space
 - There is a plane that goes through the point and the centres of projection of the cameras
 - We call this plane the epipolar plane
 - The lines where the plane intersects the images are called conjugate epipolar lines



Stereo correspondence has been reduced to a 1D search !

The essential matrix

- For all corresponding points in homogeneous coordinates. The essential matrix satisfies

$$\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{x}}_1 = 0$$

- The epipolar lines are easily retrieved from the essential matrix. The condition for a point being on a line

$$\begin{bmatrix} a & b & c \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = 0.$$

The essential matrix

- For all corresponding points in homogeneous coordinates. The essential matrix satisfies

$$\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{x}}_1 = 0$$

- The epipolar lines are easily retrieved from the essential matrix. The condition for a point being on a line

$$\begin{bmatrix} a & b & c \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = 0.$$

- The line $\mathbf{l}_1 = \tilde{\mathbf{x}}_2 \mathbf{E}$ is the epipolar line in image 1 due to the point \mathbf{x}_2 in image 2

The essential matrix

- For all corresponding points in homogeneous coordinates. The essential matrix satisfies

$$\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{x}}_1 = 0$$

- The epipolar lines are easily retrieved from the essential matrix. The condition for a point being on a line
- The line $\mathbf{l}_1 = \tilde{\mathbf{x}}_2^T \mathbf{E}$ is the epipolar line in image 1 due to the point \mathbf{x}_2 in image 2
- Similarly $\mathbf{l}_2 = \tilde{\mathbf{x}}_1^T \mathbf{E}^T$

The essential matrix

- For all corresponding points in homogeneous coordinates. The essential matrix satisfies

$$\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{x}}_1 = 0$$

- The epipoles can also be extracted from the essential matrix
 - Every epipolar line in image 1 passes through the epipole $\tilde{\mathbf{e}}_1$, such that $\tilde{\mathbf{x}}_2^T \mathbf{E} \tilde{\mathbf{e}}_1 = 0$
 - Epipoles can be retrieved by computing the singular value decomposition $\mathbf{E} = \mathbf{U} \mathbf{L} \mathbf{V}^T$ of the essential matrix
 - Setting $\tilde{\mathbf{e}}_1$ to the last column of \mathbf{V} and
 - $\tilde{\mathbf{e}}_2$ to the last row of \mathbf{U}

The essential matrix

- The essential matrix is defined

$$\mathbf{E} = \tau_x \boldsymbol{\Omega}$$

- Applying $\mathbf{E} = \mathbf{U}\mathbf{L}\mathbf{V}^T$

- $\tau_x = \mathbf{U}\mathbf{L}\mathbf{W}\mathbf{U}^T$

- $\boldsymbol{\Omega} = \mathbf{U}\mathbf{W}^{-1}\mathbf{V}^T$

$$\mathbf{W} = \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

3D reconstruction

- Goal
 - Given binocular stereo images and their correspondences.
 - Recover the 3D geometry of the imaged scene.
- What can be achieved
 - Given intrinsic and extrinsic geometry of the cameras:
Absolute Euclidean reconstruction.
 - Given only intrinsic camera geometries: Reconstruction up
to a scale factor.
 - Given no information on camera geometries:
Reconstruction up to a projective transformation.

Outline

- Introduction
- 3D shapes from 2D images
- Stereo vision
- Correspondence
- Epipolar geometry
- The fundamental matrix
- The essential matrix
- RANSAC
- 3D reconstruction

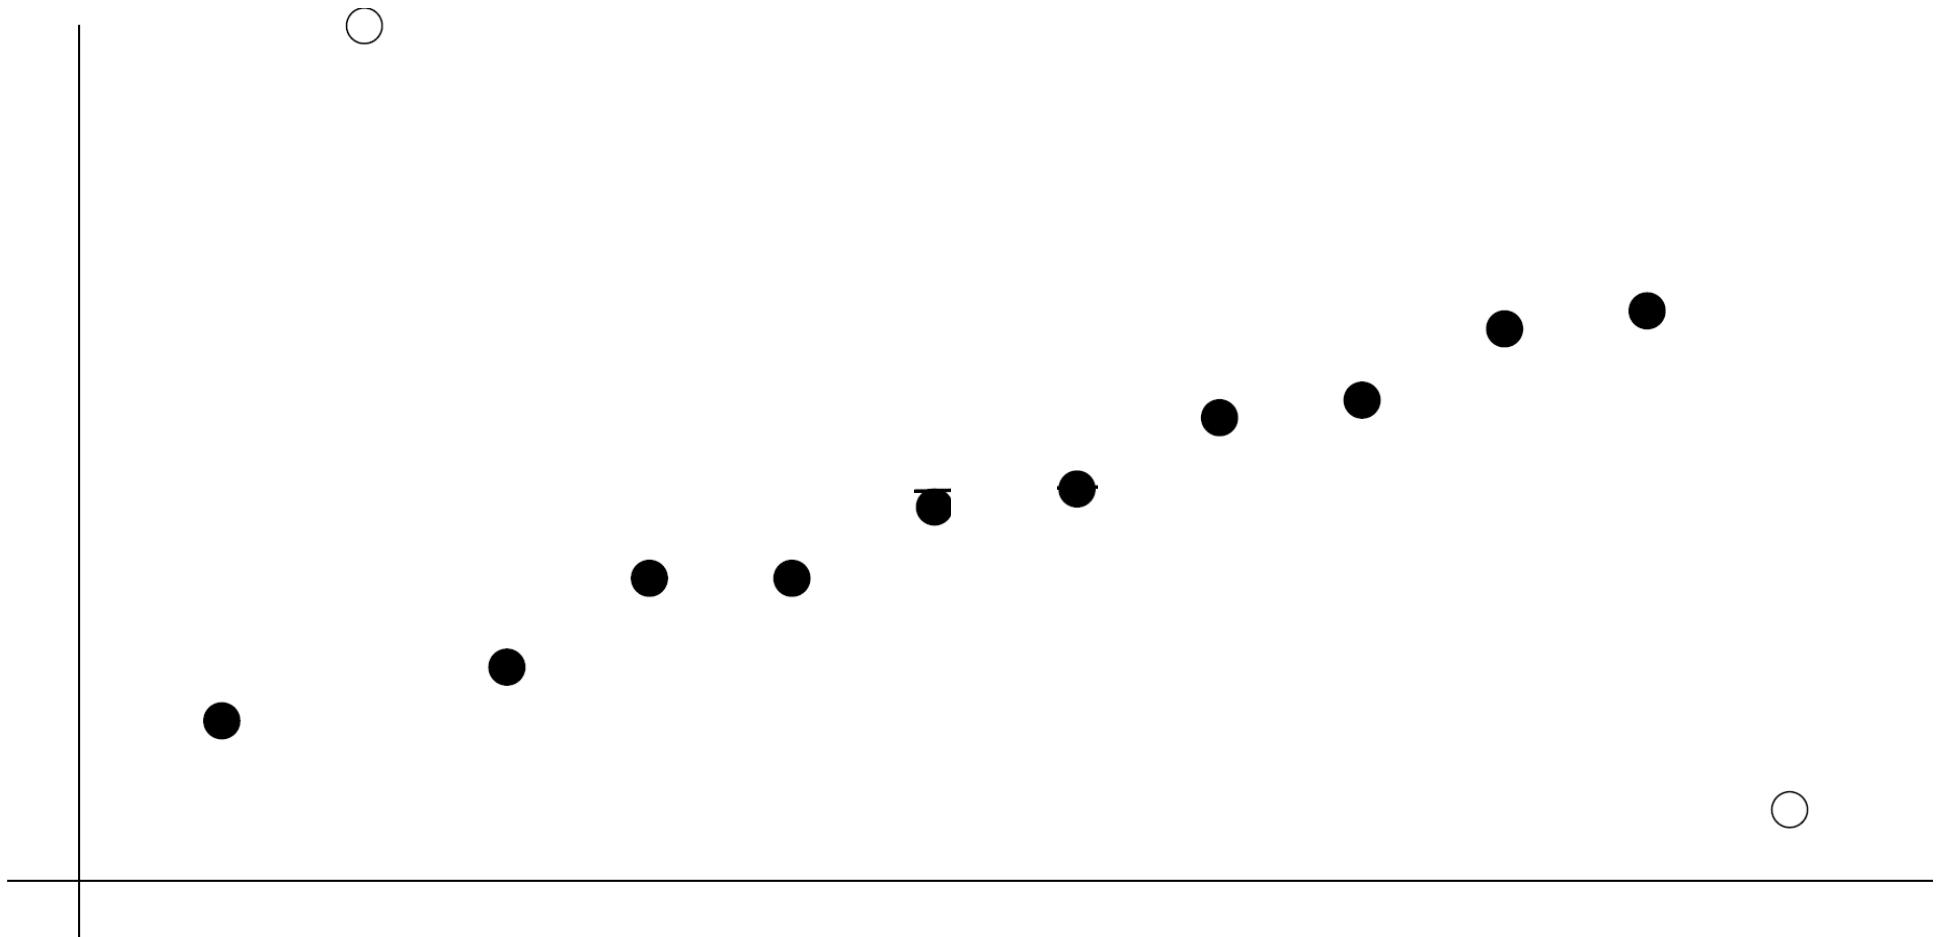
RANSAC

- RAndom SAmples Consensus
- Fischler and Bolles, 1981
 - Estimating a straight line
 - Fit of a set of 2-dimensional points
 - Equivalent to estimating a 1-dimensional affine transformation

$$x' = ax + b$$

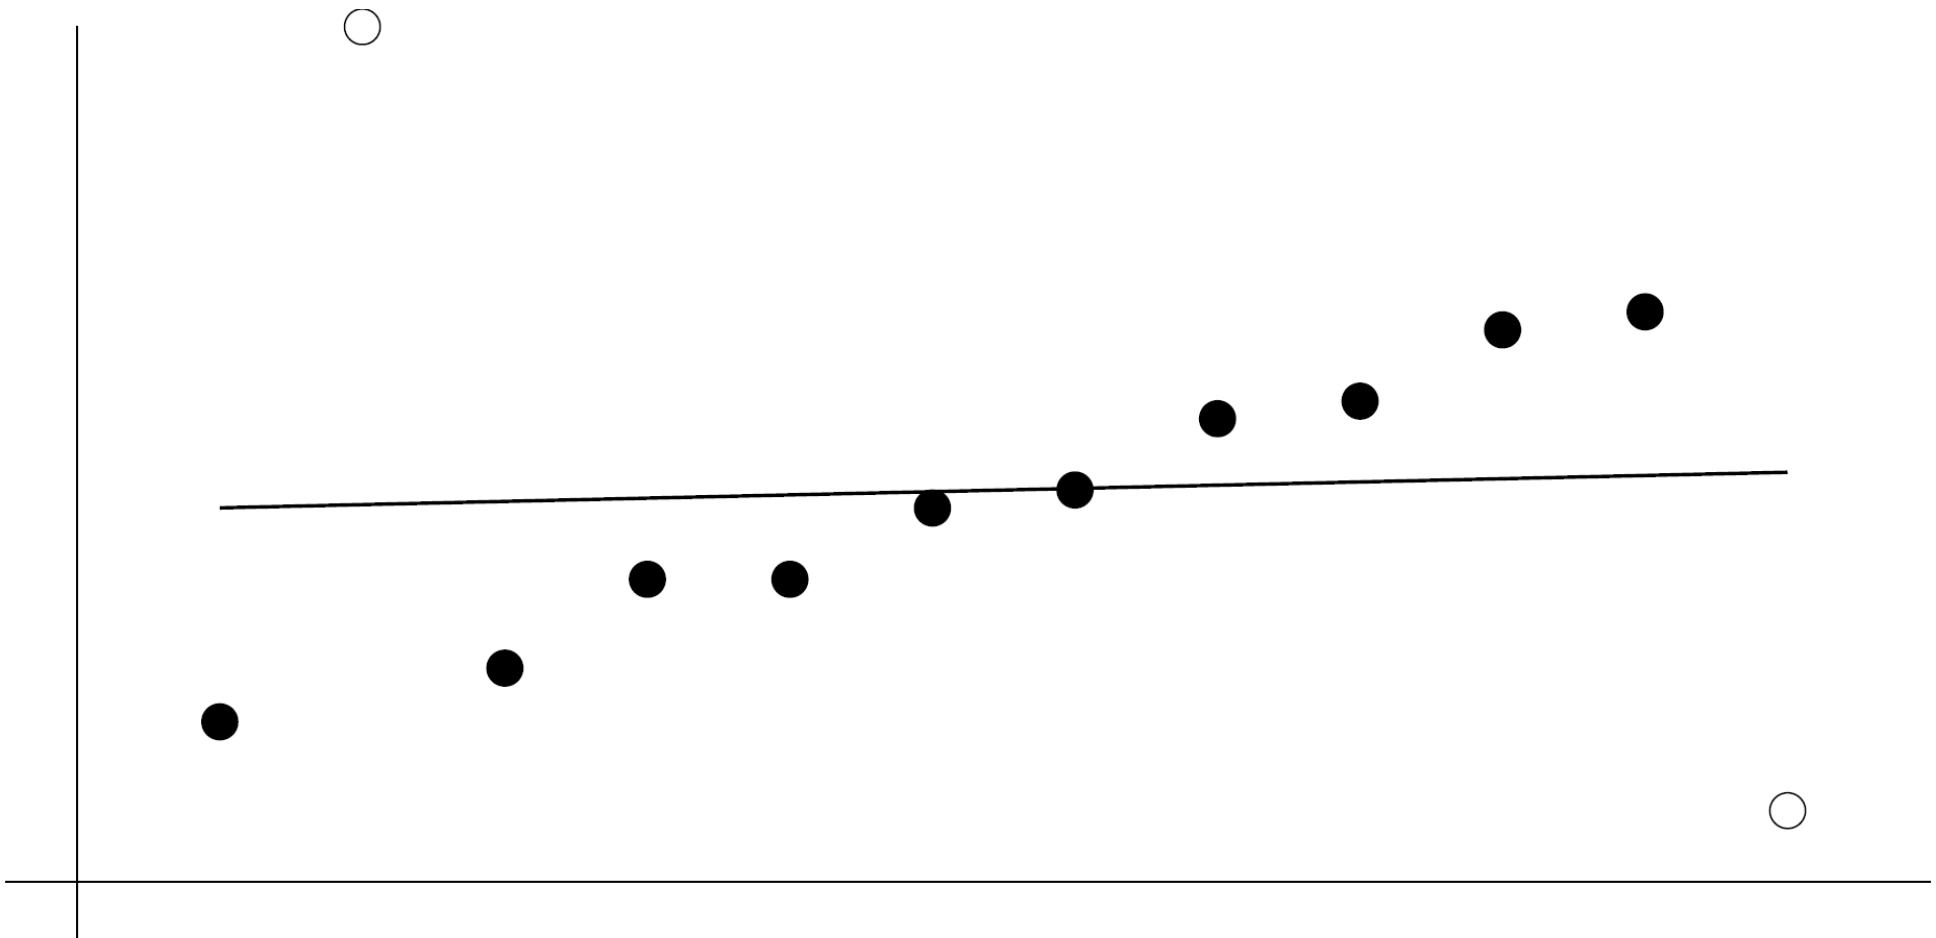
RANSAC

- RAndom SAmple Consensus



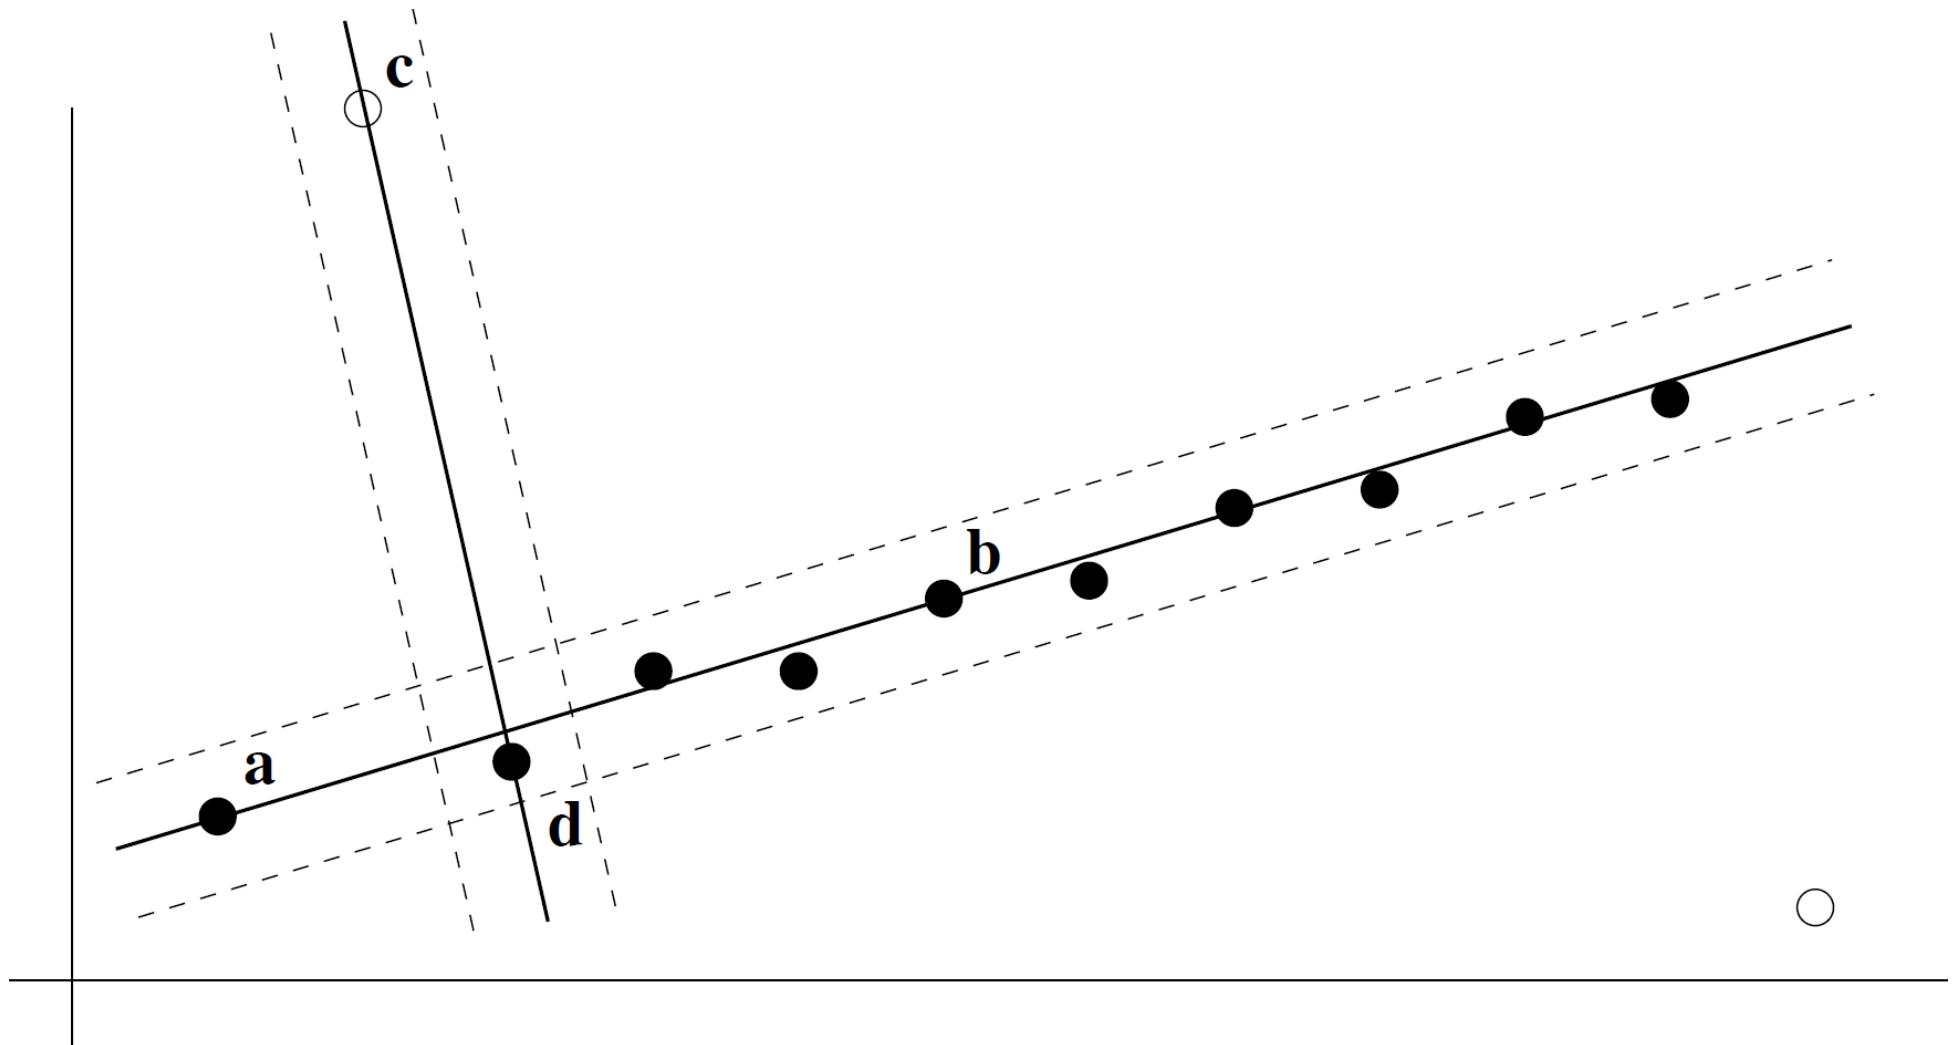
RANSAC

- RAndom SAmples Consensus



RANSAC

- RAndom SAmples Consensus



RANSAC

- RAndom SAmples Consensus

1. Repeat k times:

1. Draw a random sample of n data points. 
2. Determine the model parameters using these points.
3. For each data point outside the sample:
 1. If its distance to the model is less than d , then label it as a “good” point.
4. If there are at least t “good” data points, then declare this model “good” and quit this loop.

2. Refit the model with the most “good” points using all of them, including the n sample points.

RANSAC

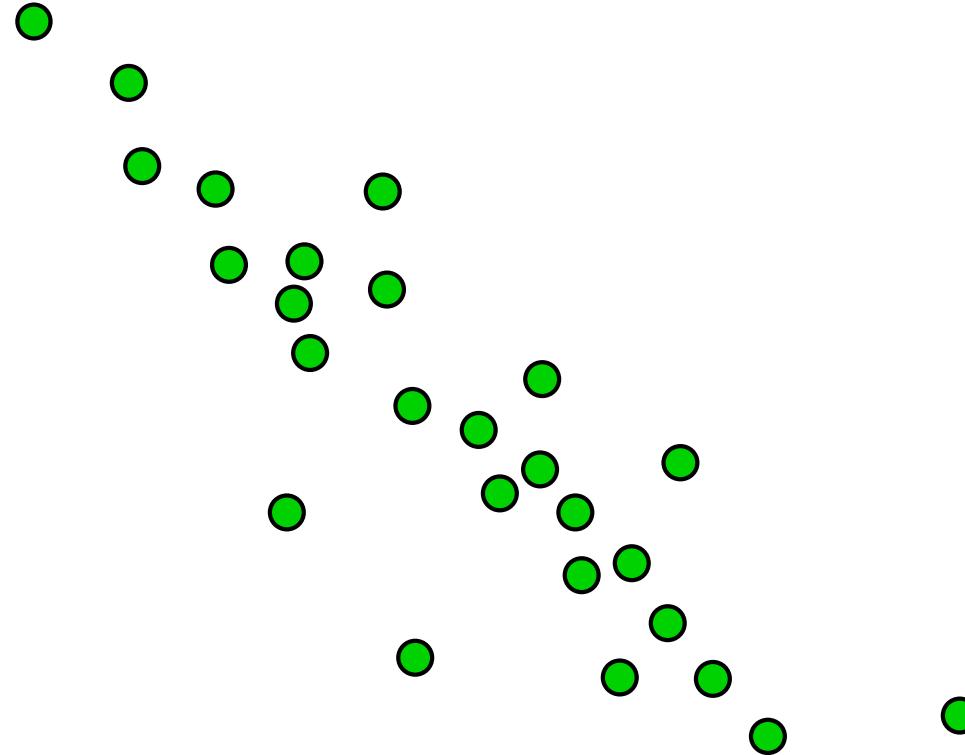
- RAndom SAmples Consensus

1. Repeat k times:

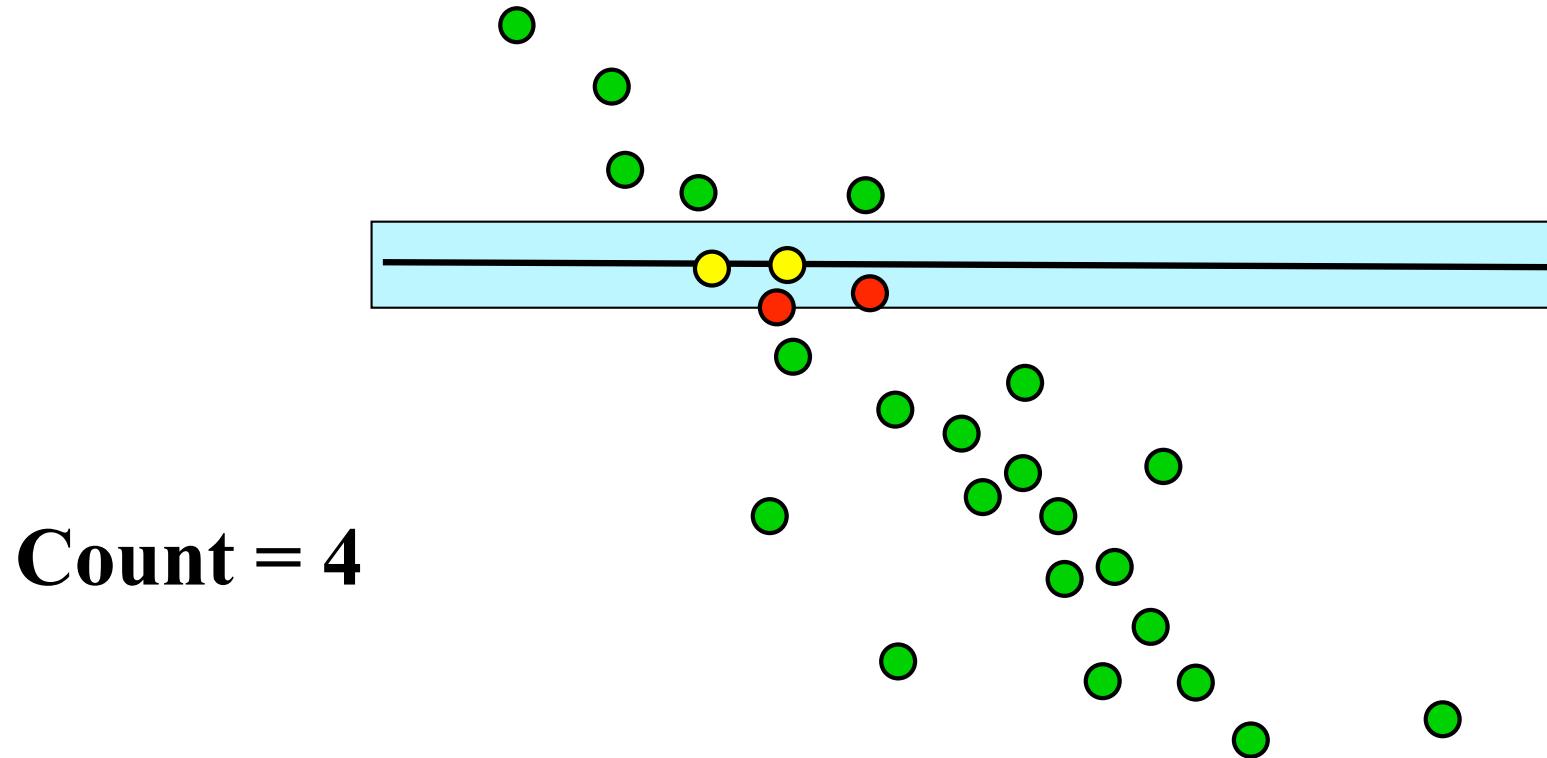
1. Draw a random sample of n data points.
2. Determine the model parameters using these points.
3. For each data point outside the sample:
 1. If its distance to the model is less than d , then label it as a “good” point.
4. If there are at least t “good” data points, then declare this model “good” and quit this loop.

2. Refit the model with the most “good” points using all of them, including the n sample points.

RANSAC

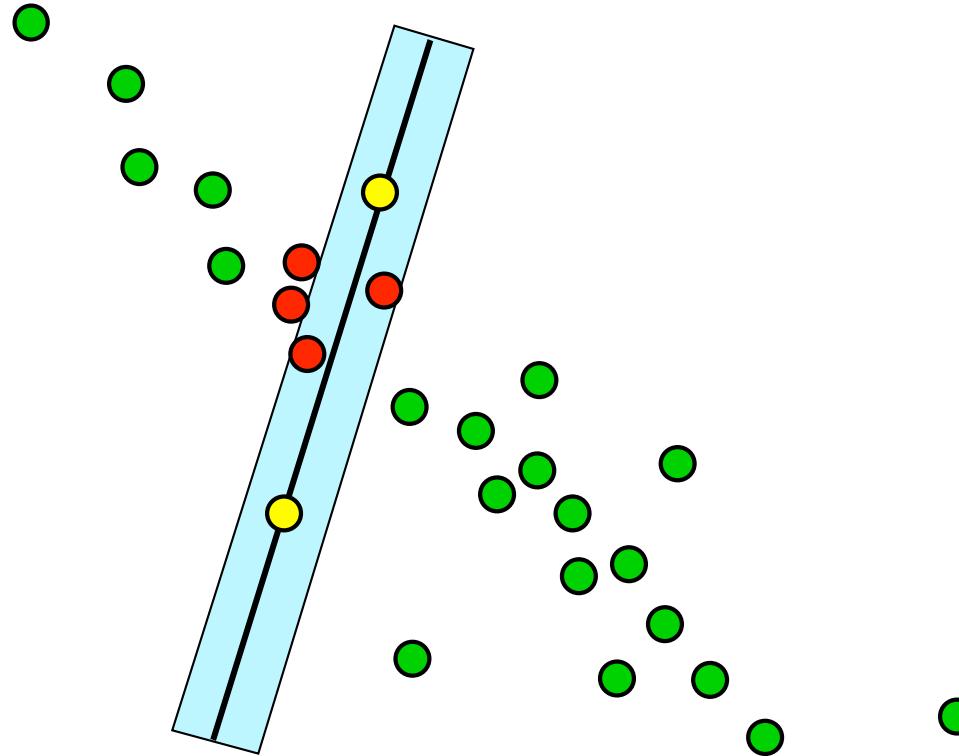


RANSAC

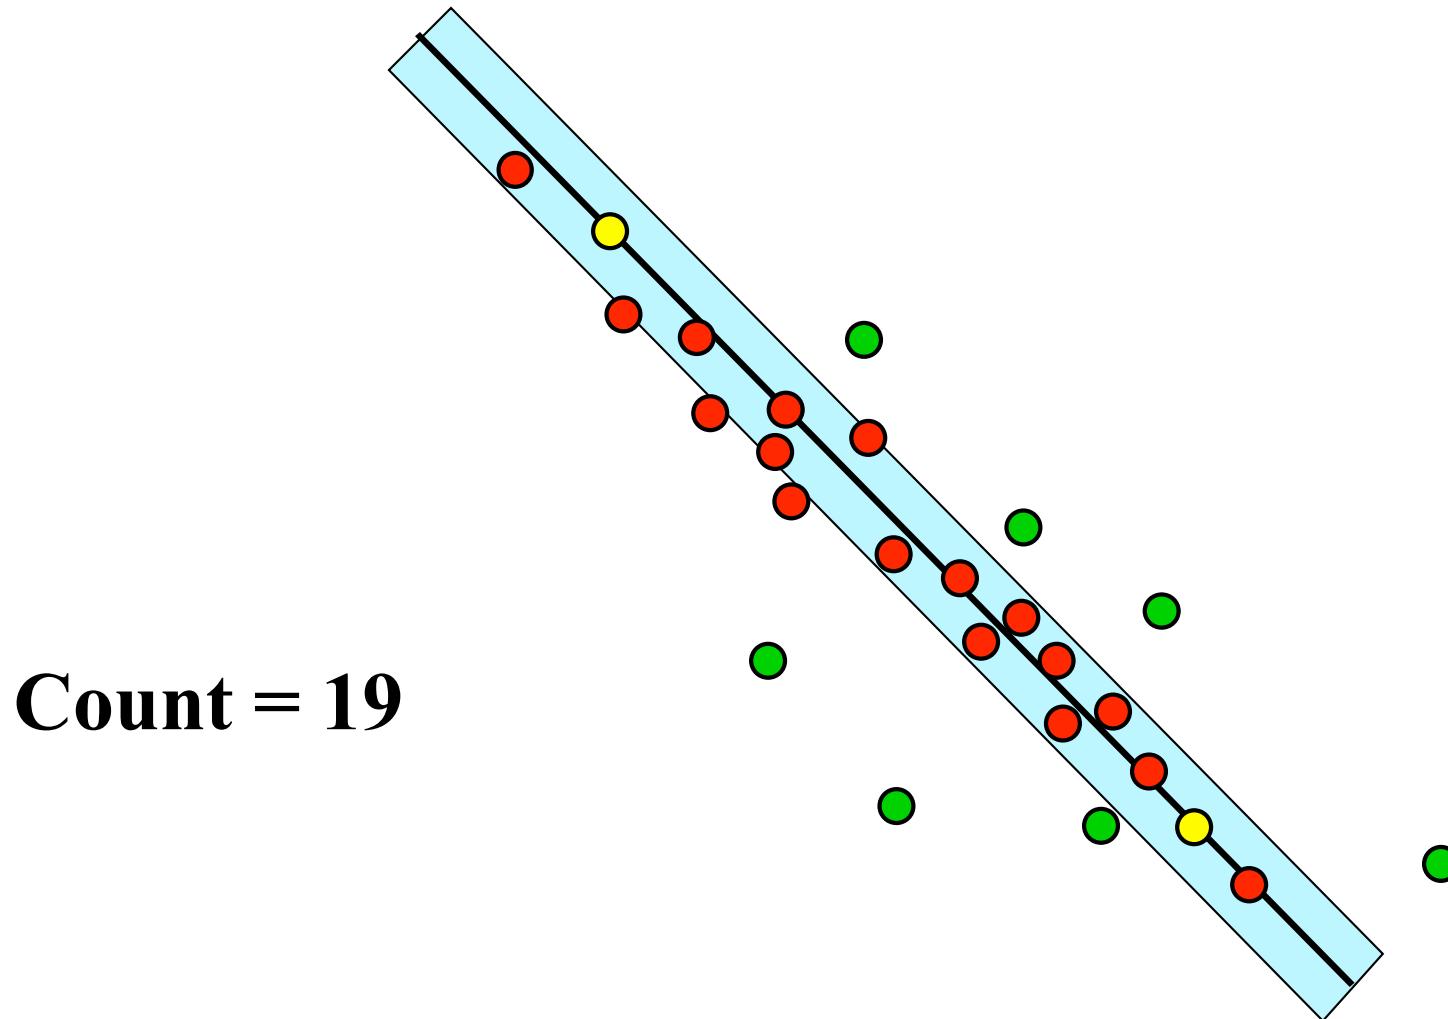


RANSAC

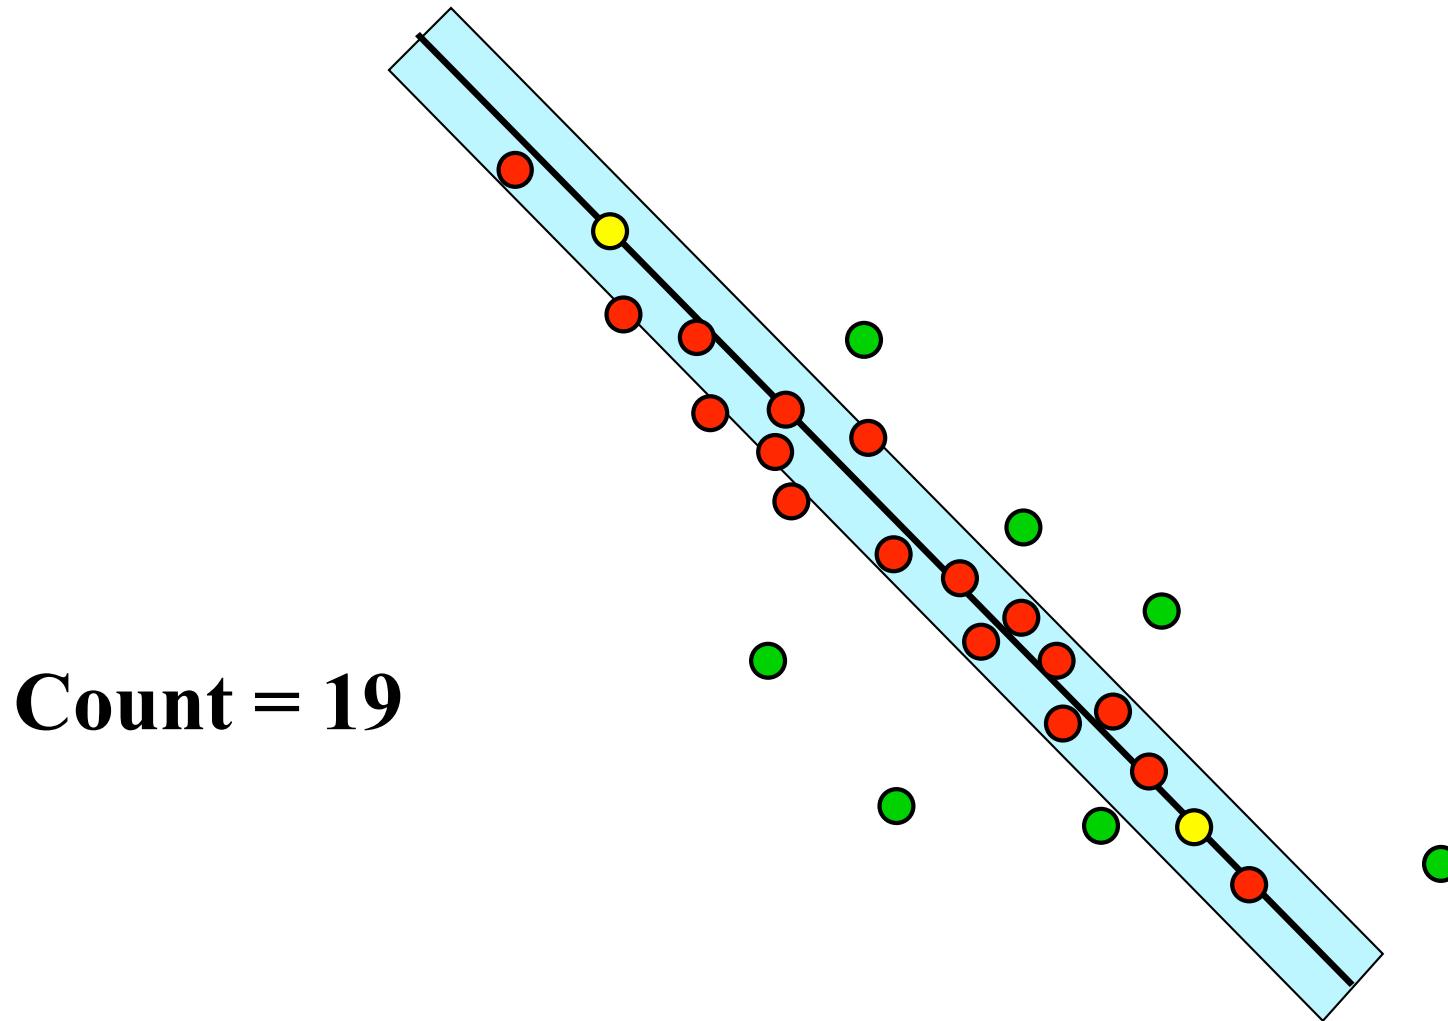
Count = 6



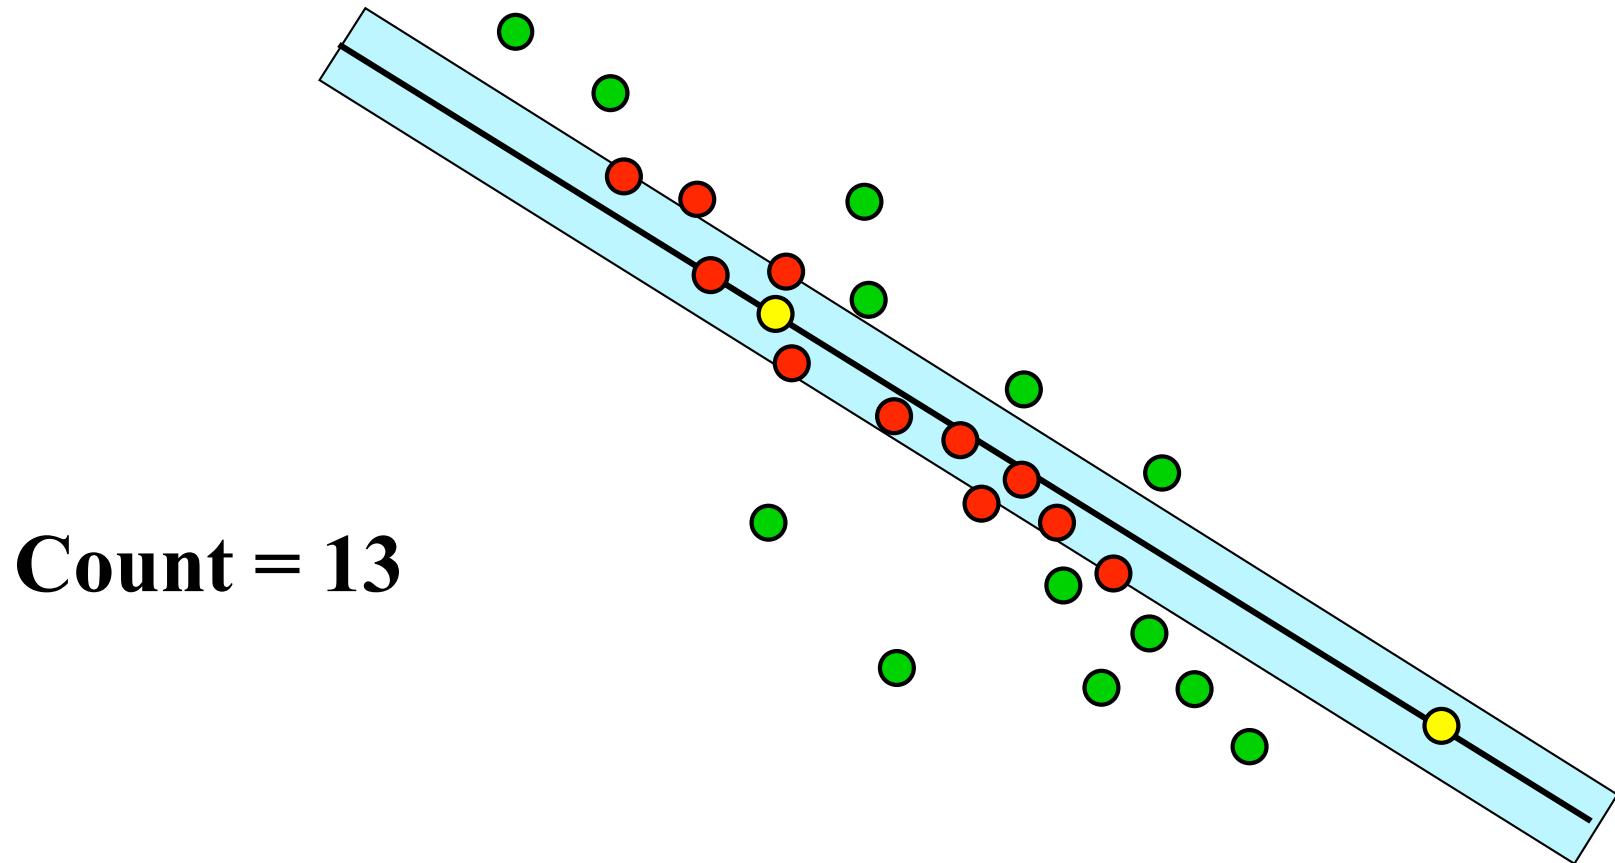
RANSAC



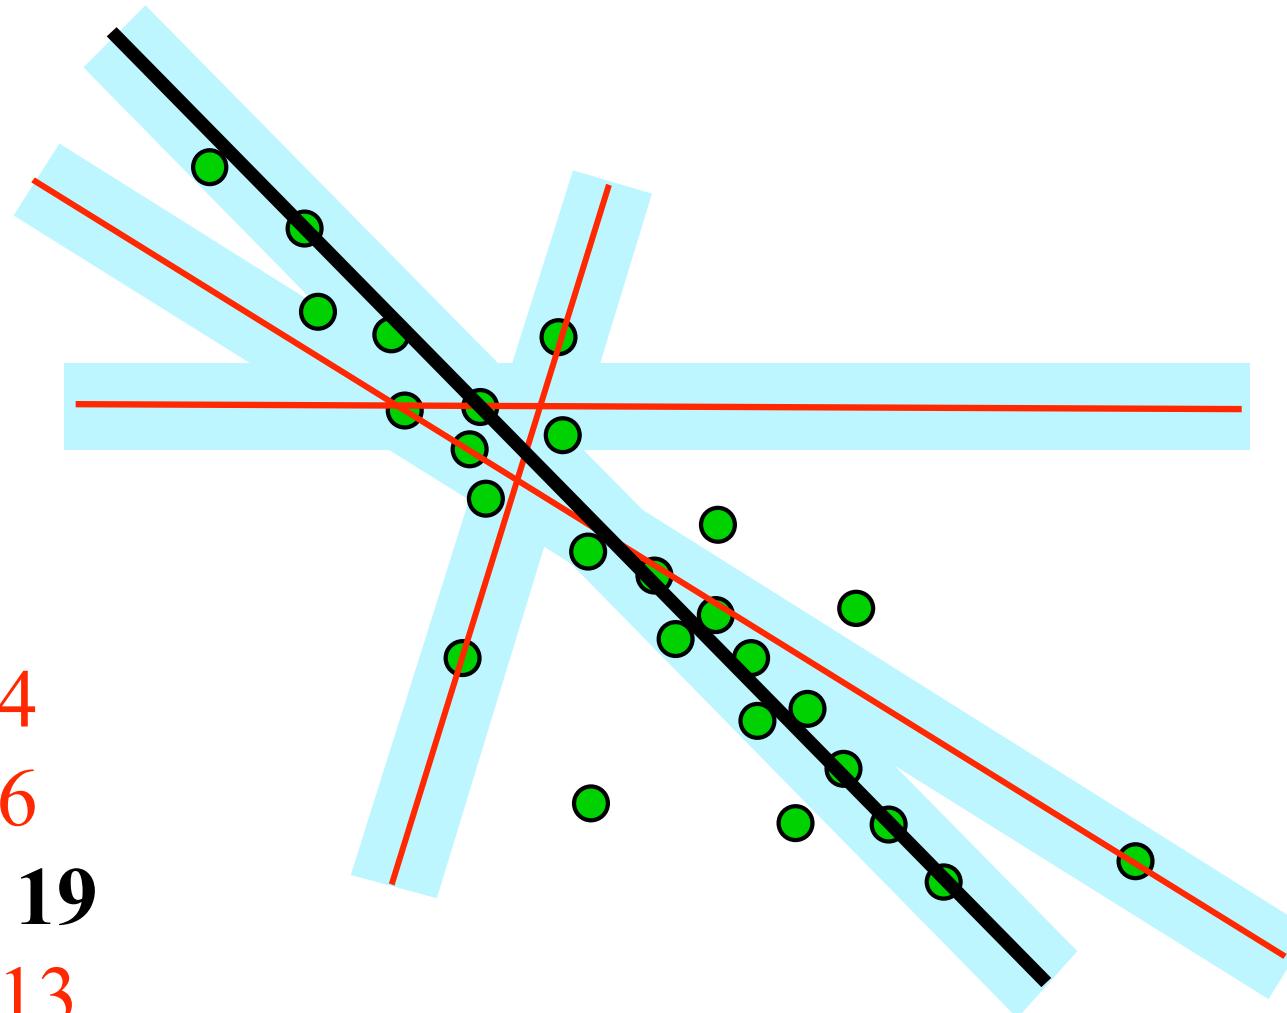
RANSAC



RANSAC



RANSAC



Count = 4

Count = 6

Count = 19

Count = 13

RANSAC

- RAndom SAmples CoNsensus
- Issues
 - Each sample comprises exactly n data points for an n -parameter model.
 - Measure of the distance of a data point to the model.
 - Choice of the three parameters k , d , and t .

RANSAC

- k : How many trials?
- It is unfeasible and unnecessary to try every possible trial
- The number of trials k must be sufficiently high to ensure with probability p that at least one of the random samples of n is free from outliers

RANSAC

- k : How many trials?
- The number of trials k must be sufficiently high to ensure with probability p that at least one of the random samples of n is free from outliers



$$\begin{aligned} E(k) &= P(\text{one_good_sample_in_one_draw}) \\ &+ P(\text{one_good_sample_in_two_draws}) + \dots \end{aligned}$$

RANSAC

- k : How many trials?
- The number of trials k must be sufficiently high to ensure with probability p that at least one of the random samples of n is free from outliers

$$\begin{aligned} E(k) &= P(\text{one_good_sample_in_one_draw}) \\ &+ P(\text{one_good_sample_in_two_draws}) + \dots \end{aligned}$$

- Let's call w the probability of inliers, that is the probability that a selected point is an inlier

$$E(k) = w^n + 2w^n(1-w^n) + 3w^n(1-w^n)^2 + \dots = w^{-n}$$

RANSAC

- k : How many trials?
- Let's call w the fraction of inliers, well, the probability that a selected point is an inlier

$$E(k) = w^n + 2w^n(1-w^n) + 3w^n(1-w^n)^2 + \dots = w^{-n}$$

- We need to increase confidence, so, select more than w^{-n}
- Add a few standard deviations to k

$$\sigma(k) = \frac{\sqrt{1 - w^n}}{w^n}$$

RANSAC

- k : How many trials?
- Let's call w the fraction of inliers, well, the probability that a selected point is an inlier

$$E(k) = w^n + 2w^n(1-w^n) + 3w^n(1-w^n)^2 + \dots = w^{-n}$$

- We need to increase confidence, so, select more than w^{-n}
- Add a few standard deviations to k

$$\sigma(k) = \frac{\sqrt{1-w^n}}{w^n}$$

- The probability z of only outliers $z = (1-w^n)^k$

RANSAC

- k : How many trials?
- Let's call w the fraction of inliers, well, the probability that a selected point is an inlier

$$E(k) = w^n + 2w^n(1-w^n) + 3w^n(1-w^n)^2 + \dots = w^{-n}$$

- We need to increase confidence, so, select more than w^{-n}
- Add a few standard deviations to k

$$\sigma(k) = \frac{\sqrt{1-w^n}}{w^n}$$

- The probability z of only outliers $z = (1-w^n)^k$

$$k = \frac{\log(z)}{\log(1-w^n)}$$

RANSAC

- k : How many trials?, $z=0.01$

Number of samples required

n	Proportion of Outliers $1-w$						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Table 1: From [Hartley and Zisserman, 2000]: The number k of samples required to ensure, with a probability $p = 0.99$, that at least one sample has no outliers for a given sample size s and a proportion of outliers ϵ .

RANSAC

- d : How close must a point to be good?
 - It is usually chosen empirically
 - Expected measurement error
 - Eyeballing
 - Trial-and-error

RANSAC

- d : How close must a point to be good?
 - It is usually chosen empirically
 - Expected measurement error
 - Eyeballing
 - Trial-and-error
 - But there are fancier methods too

RANSAC

- d : How close must a point to be good?
 - It is usually chosen empirically
 - Expected measurement error
 - Eyeballing
 - Trial-and-error
 - But there are fancier methods too
 - Using a cumulative chi-squared distribution
 - Assuming error is Gaussian with zero mean

RANSAC

- t : How many good points for a good fit? Number of points that must agree

RANSAC

- t : How many good points for a good fit? Number of points that must agree
 - A good rule of thumb is to require at least as many “good” points as we would expect from the proportion w of inliers present in the data

$$t = wN$$

- Where N is the total number of data points

RANSAC

- t : How many good points for a good fit? Number of points that must agree $t = wN$
 - Estimating the proportion of inliers w
 - Begin with $w=0$
 - After each RANSAC iteration
 - Choose a sample and count the number of inliers
 - Set $w=(\text{number of inliers})/(\text{total number of points})$
 - Recompute k
- $$k = \frac{\log(z)}{\log(1 - w^n)}$$

RANSAC

- Compute the 2D homography matrix between two images
 - Input: Two images
 - Output: Homography and correspondence of interesting points

RANSAC

- Compute the 2D homography matrix between two images
 - Input: Two images
 - Output: Homography and correspondence of interesting points
- RANSAC for Automatic Computation of a Fundamental Matrix

$$x'^T \boxed{F} x = 0$$

RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T \boxed{F} x = 0$
 1. Compute interest points in each image
 - E.g. Use Harris to detect 7 points in each image

RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T \boxed{F}x = 0$
 1. Compute interest points in each image
 - E.g. Use Harris to detect 7 points in each image
 2. Putative correspondences
 - Compute a set of interest point matches based on proximity and similarity of their intensity neighborhood

RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T \boxed{F}x = 0$
 1. Compute interest points in each image
 - E.g. Use Harris to detect 7 points in each image
 2. Putative correspondences
 - Compute a set of interest point matches based on proximity and similarity of their intensity neighborhood
 3. Putative F : Use the minimal algorithm, if there are three solutions, try them all

RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T \boxed{F} x = 0$
 1. Compute interest points in each image
 - E.g. Use Harris to detect 7 points in each image
 2. Putative correspondences
 - Compute a set of interest point matches based on proximity and similarity of their intensity neighborhood
 3. Putative F : Use the minimal algorithm, if there are three solutions, try them all
 4. Assess the error
 - Use the symmetric transfer error or the Sampson distance

$$\frac{(x'^T F x)^2}{(Fx)_1^2 + (Fx)_2^2 + (F^T x')_1^2 + (F^T x')_2^2}$$



RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix

$$x'^T \boxed{F} x = 0$$

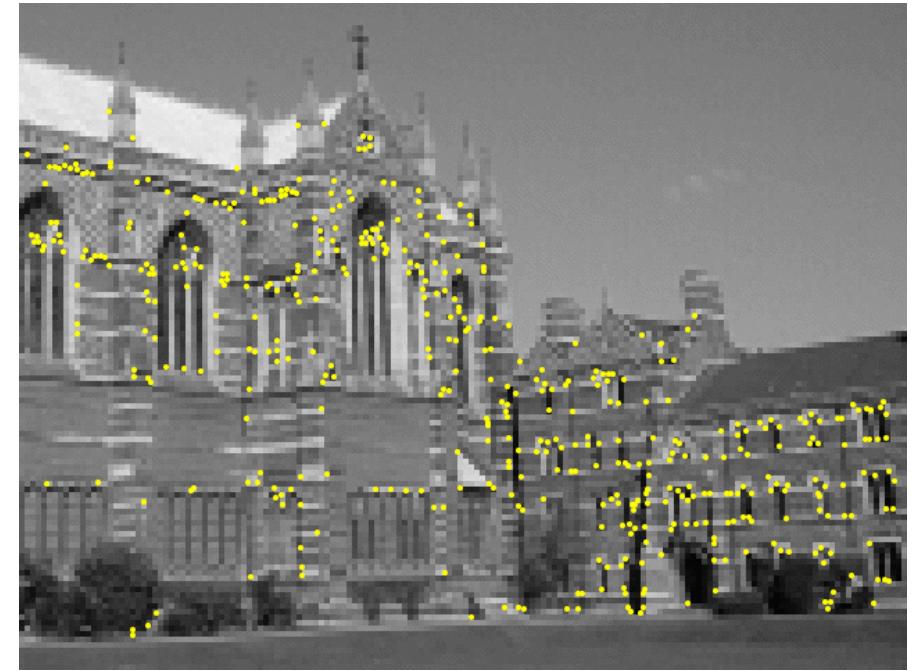
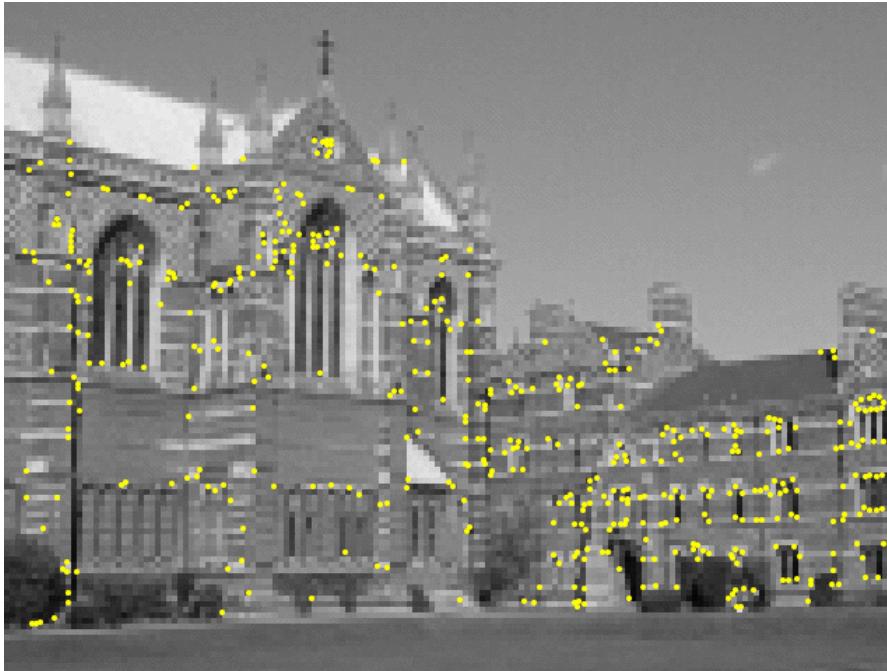
3. Putative F : Use the minimal algorithm, if there are three solutions, try them all
4. Assess the error
 - Use the symmetric transfer error or the Sampson distance

$$\frac{(x'^T F x)^2}{(Fx)_1^2 + (Fx)_2^2 + (F^T x')_1^2 + (F^T x')_2^2}$$

5. Re-estimate $\textcolor{brown}{F}$ from all inliers of the RANSAC result by optimizing the same error function as used as the distance measure. We may also add new correspondences by searching in a small strip around the epipolar lines.

RANSAC

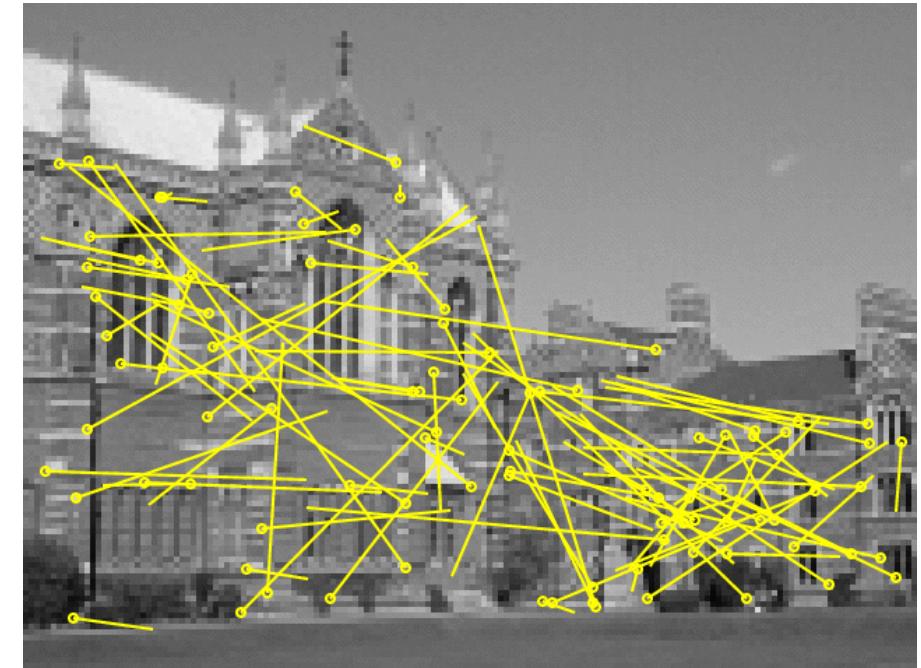
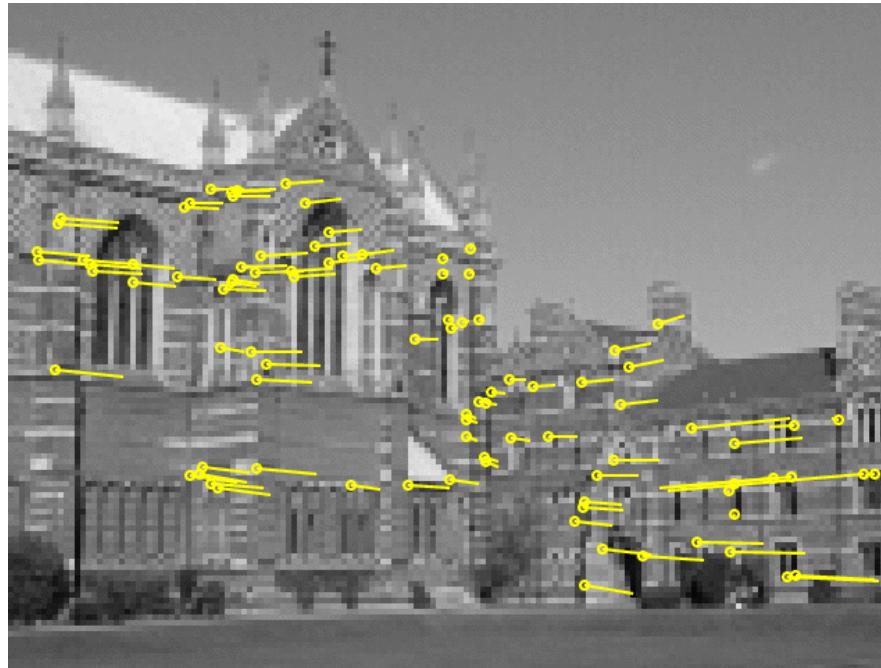
- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T \boxed{F}x = 0$



Step 1. ~500 interest points per image

RANSAC

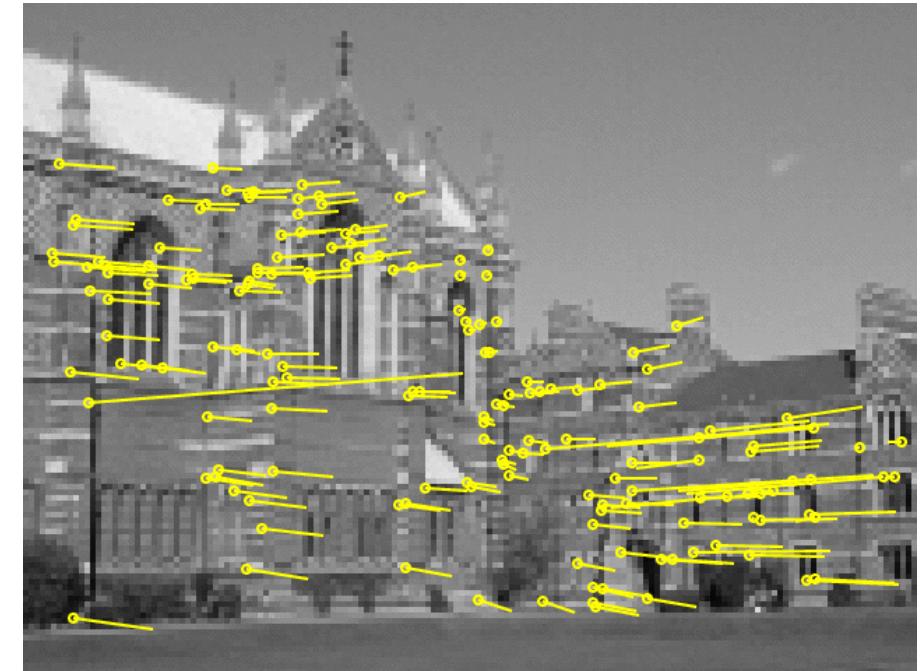
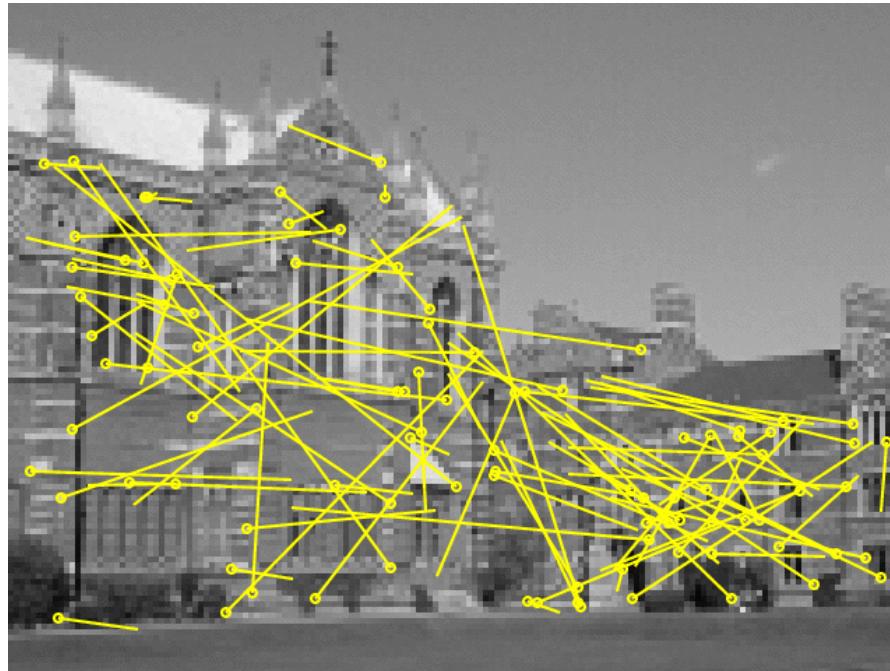
- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T F x = 0$



Step 2. Putative matches: 99 inliers, 89 outliers

RANSAC

- RANSAC for Automatic Computation of a Fundamental Matrix $x'^T F x = 0$



Step 5. All 188 putative matches: 157 revised matches

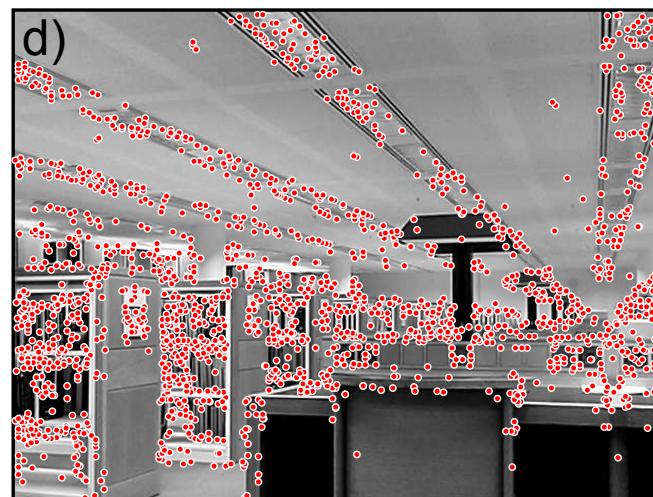
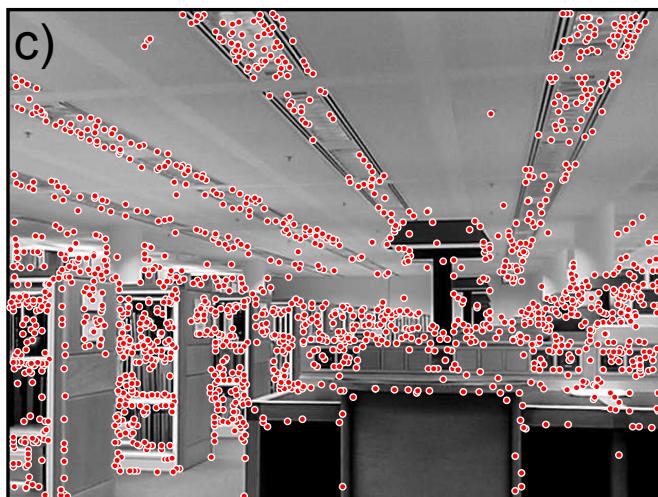
Summary

- Introduction
- 3D shapes from 2D images
- Stereo vision
- Correspondence
- Epipolar geometry
- The fundamental matrix
- The essential matrix
- RANSAC
- 3D reconstruction

Two-view reconstruction pipeline

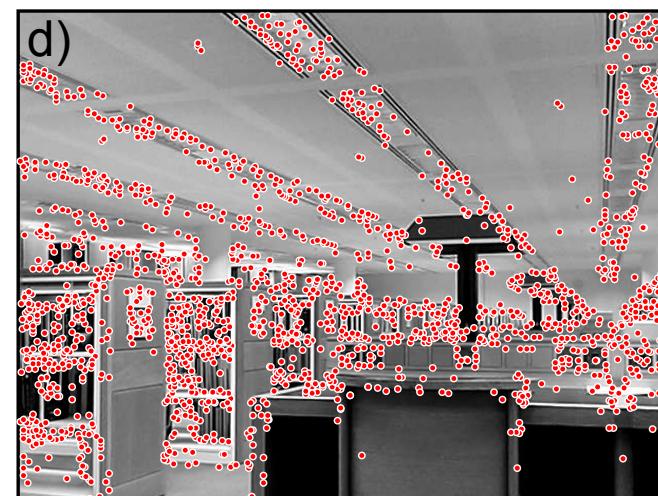
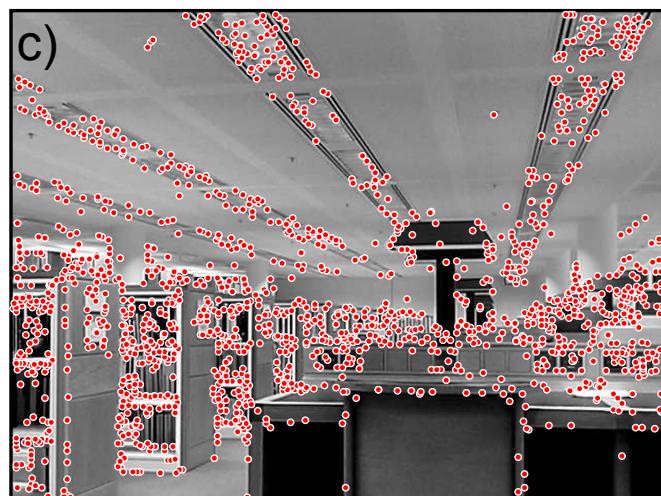
1. Compute image features

- Find salient points in each image an interest point detector such as the SIFT detector



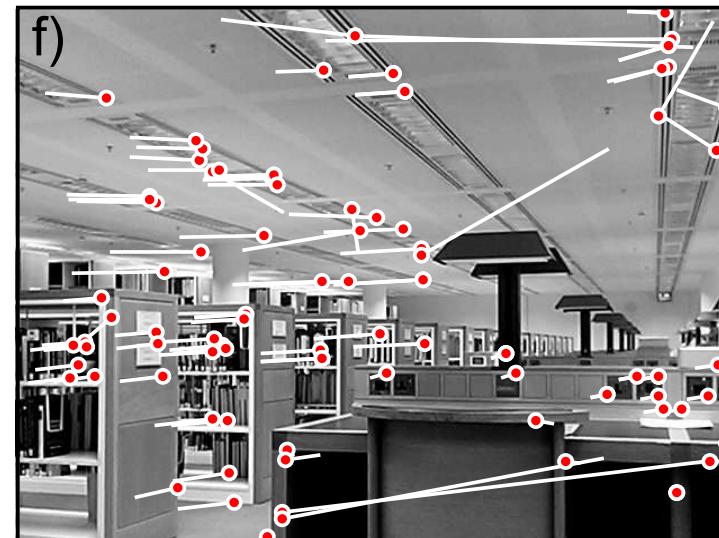
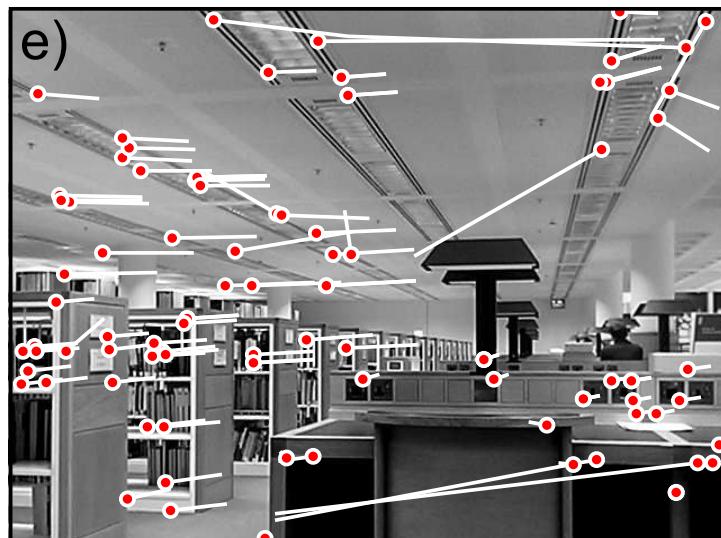
Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
 - Characterise the region around each feature in each image with a low dimensional vector, e.g. use the SIFT descriptor



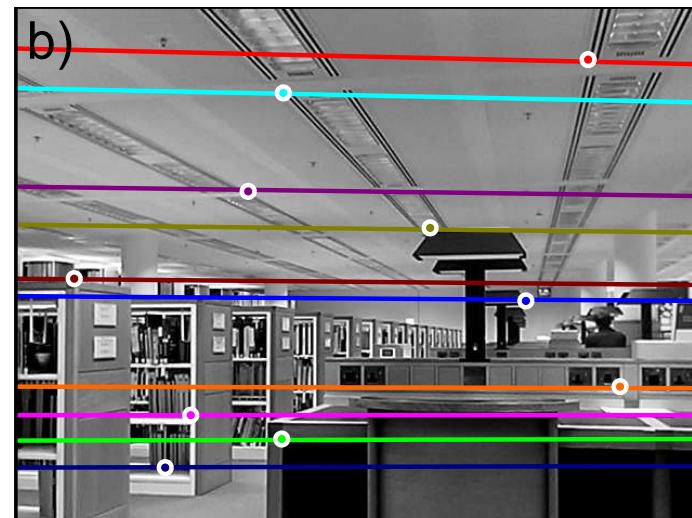
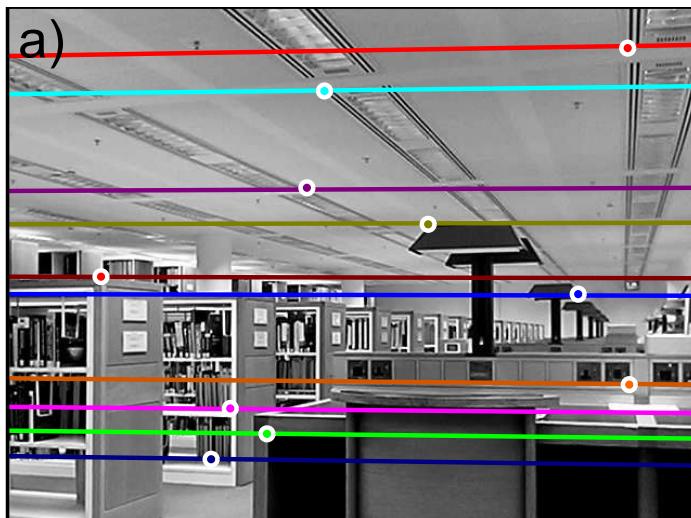
Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
 - Greedily match features between two features
 - E.g. Compute the squared distance between their region descriptors and stop when squared distance exceeds a pre-defined threshold.



Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
 - Compute the fundamental matrix using the eight-point algorithm or RANSAC

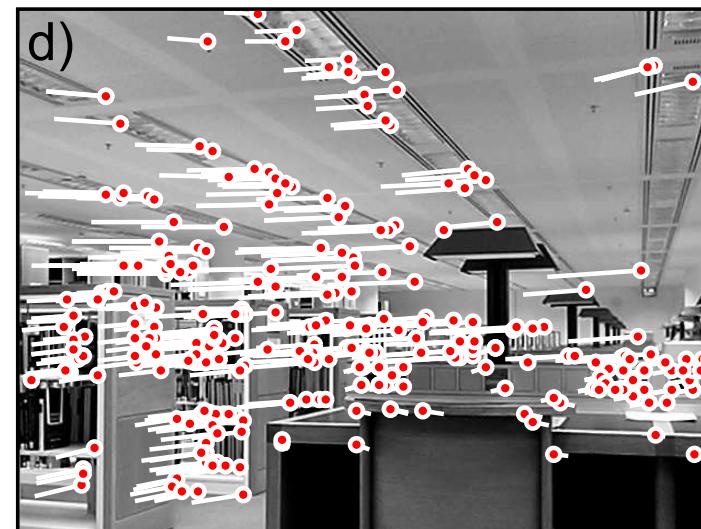
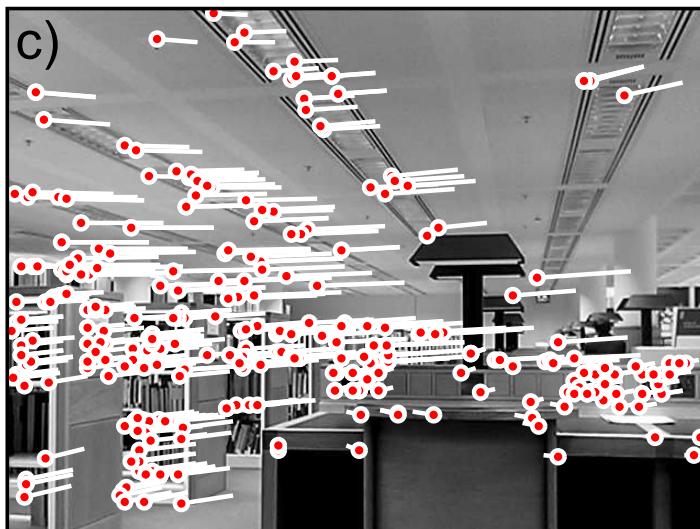


Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
5. Refine matches
 - Again greedy match features, but using epipolar geometry: If a putative match is not close to the induced epipolar line, it is rejected.
 - Recompute the fundamental matrix based on all the remaining point matches

Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
5. Refine matches



Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
5. Refine matches
6. Estimate essential matrix
 - Use intrinsic matrices from both cameras

$$\mathbf{E} = \mathbf{M}_{\text{int2}} \mathbf{F} \mathbf{M}_{\text{int1}}$$

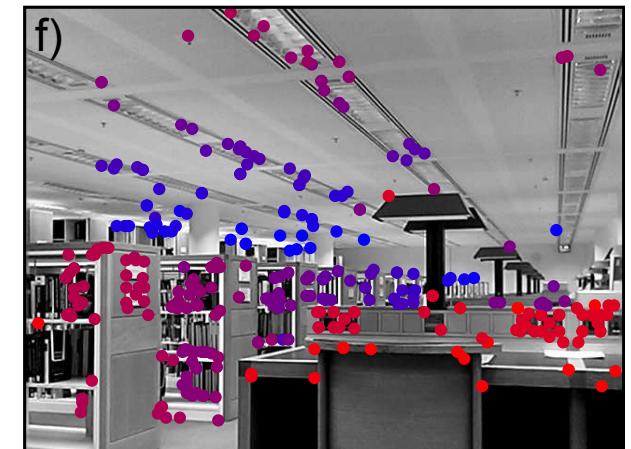
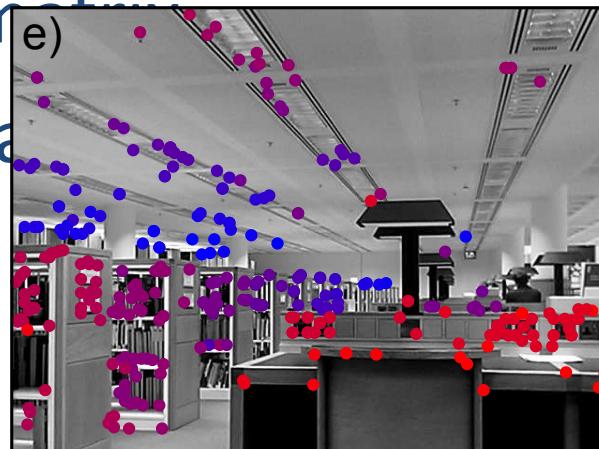
Two-view reconstruction pipeline

1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
5. Refine matches
6. Estimate essential matrix
7. Decompose essential matrix

$$\begin{aligned} \mathbf{E} &= \tau_x \boldsymbol{\Omega} & \tau_x &= \mathbf{U} \mathbf{L} \mathbf{W} \mathbf{U}^T \\ & & \boldsymbol{\Omega} &= \mathbf{U} \mathbf{W}^{-1} \mathbf{V}^T \end{aligned}$$

Two-view reconstruction pipeline

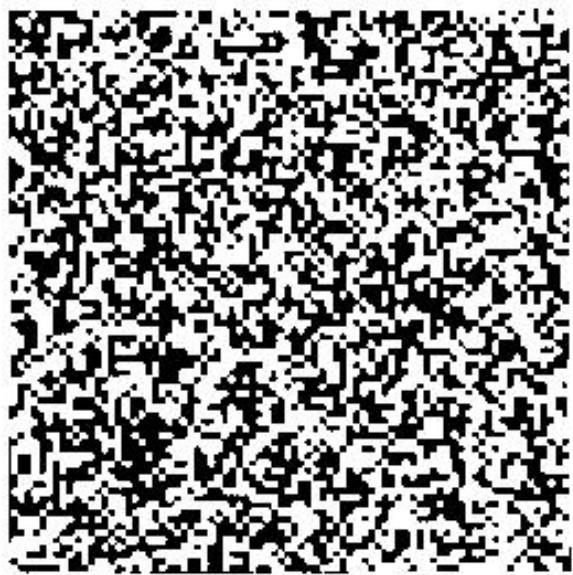
1. Compute image features
2. Compute feature descriptors
3. Find initial matches
4. Compute fundamental matrix
5. Refine matches
6. Estimate essential matrix
7. Decompose essential matrix
8. Estimate 3D points



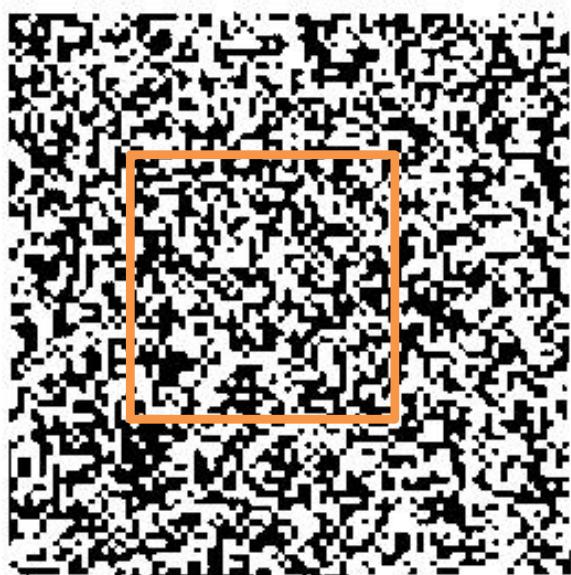
3D reconstruction

- Goal
 - Given binocular stereo images and their correspondences.
 - Recover the 3D geometry of the imaged scene.
- What can be achieved
 - Given intrinsic and extrinsic geometry of the cameras:
Absolute Euclidean reconstruction.
 - Given only intrinsic camera geometries: Reconstruction up
to a scale factor.
 - Given no information on camera geometries:
Reconstruction up to a projective transformation.

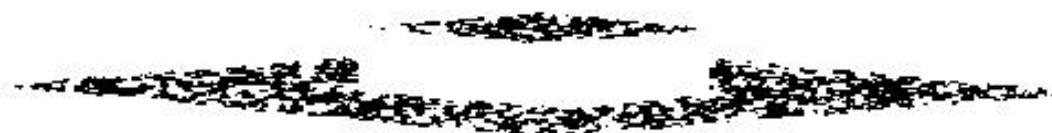
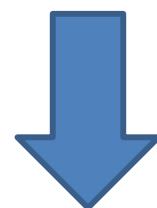
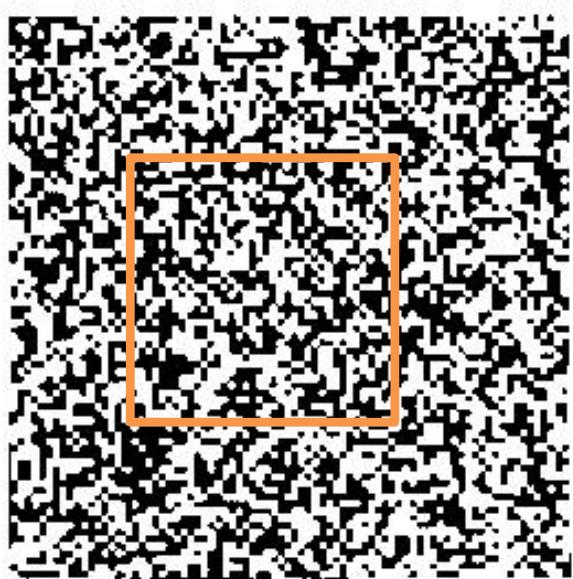
Random-dot stereogram



Random-dot stereogram



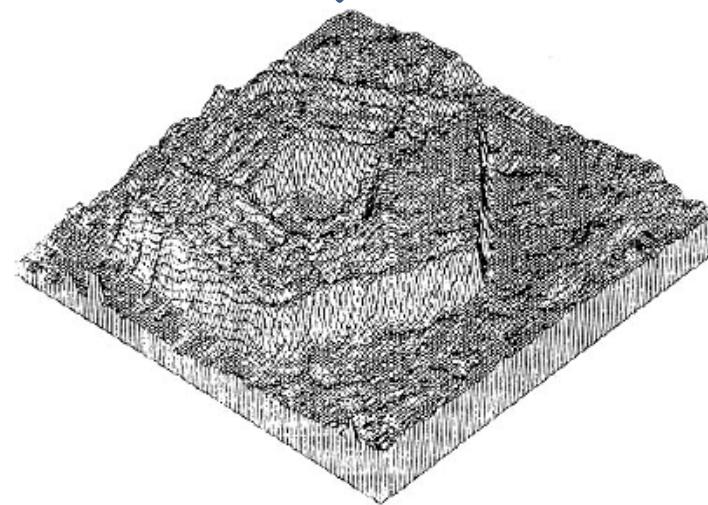
Random-dot stereogram



Real-world scene



Real-world scene

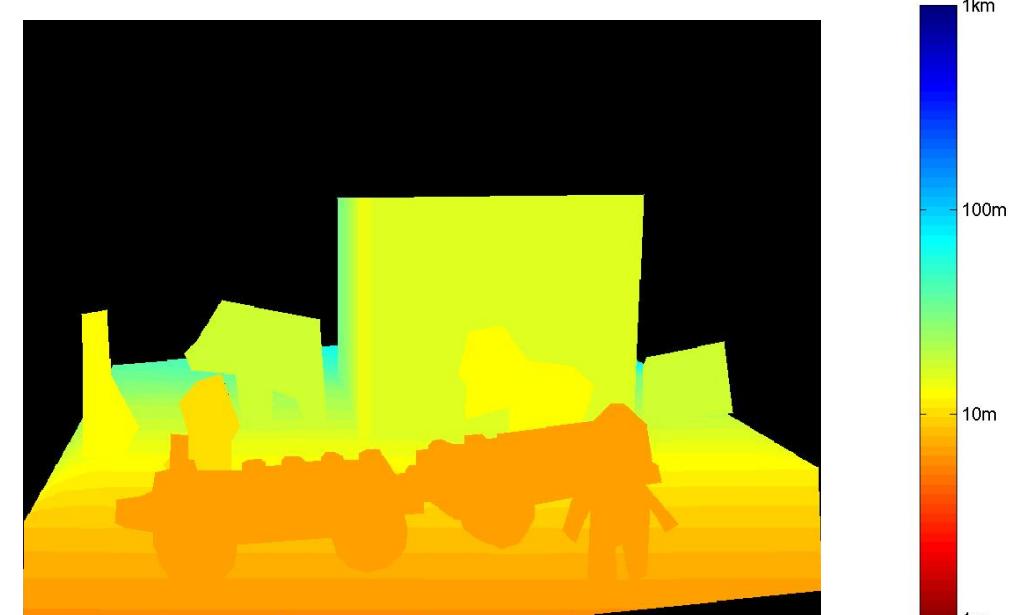


Laboratory-calibrated scene



Source: Antonio Torralba slides

Laboratory-calibrated scene

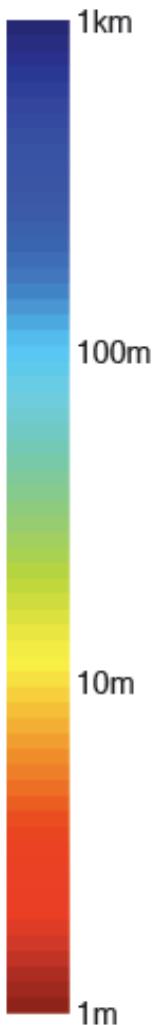
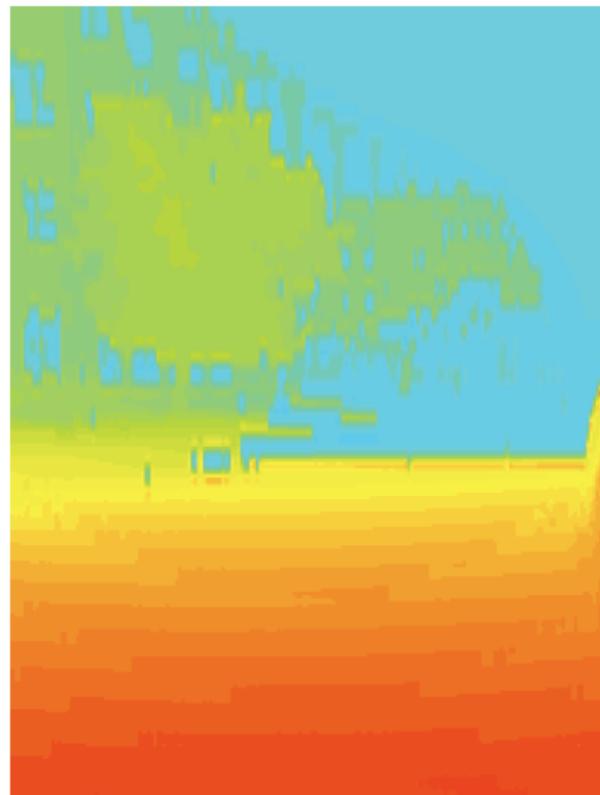


Street scene I



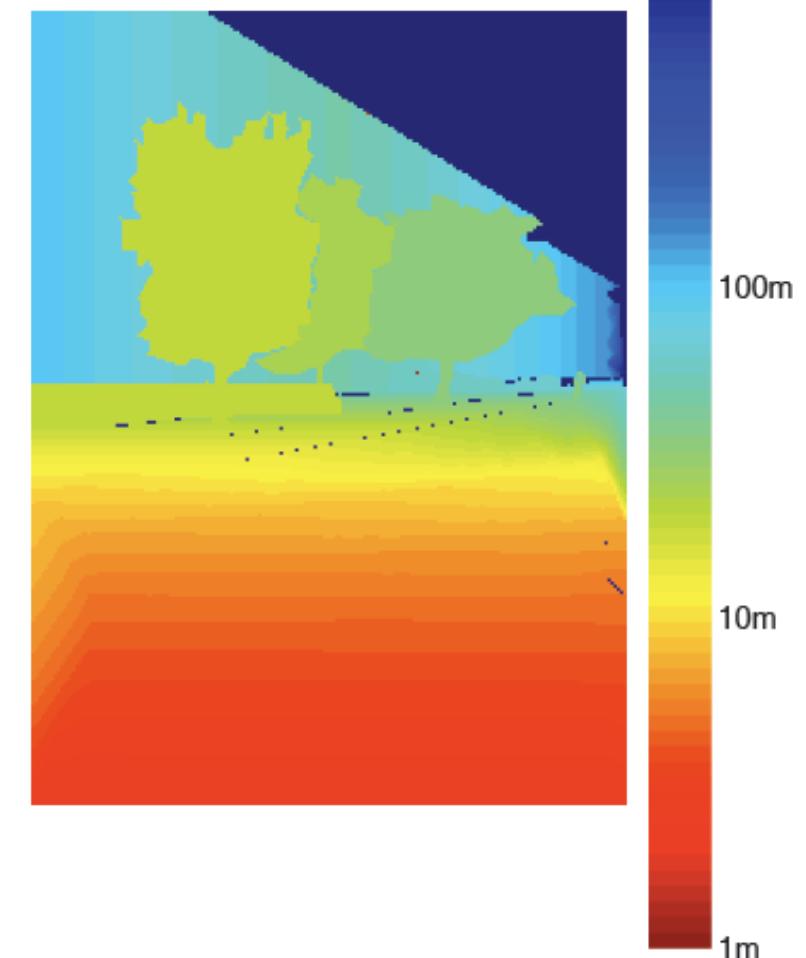
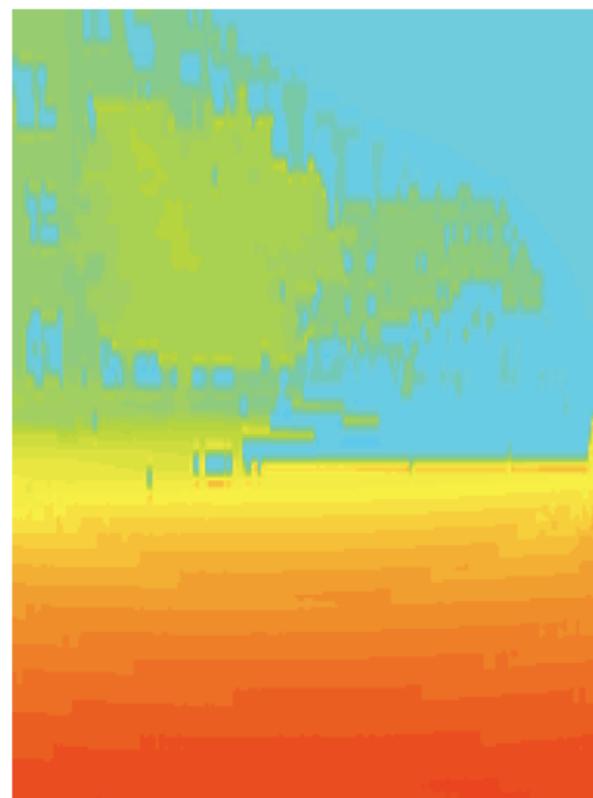
Source: Antonio Torralba slides

Street scene I



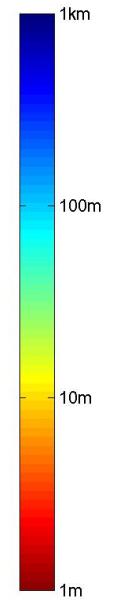
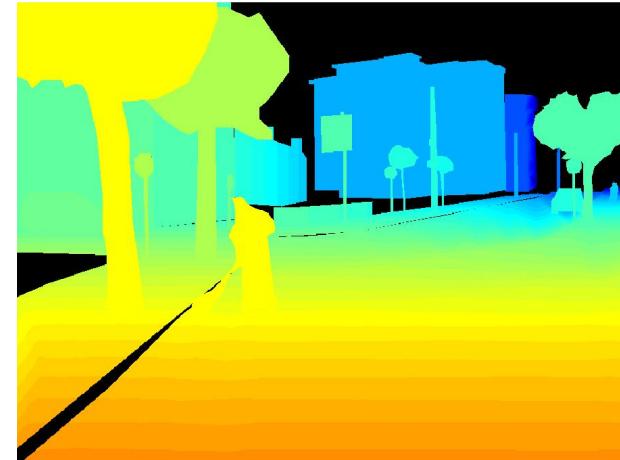
Source: Antonio Torralba slides

Street scene I

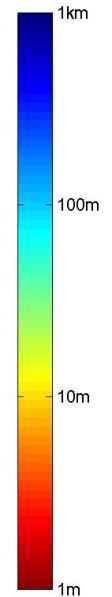
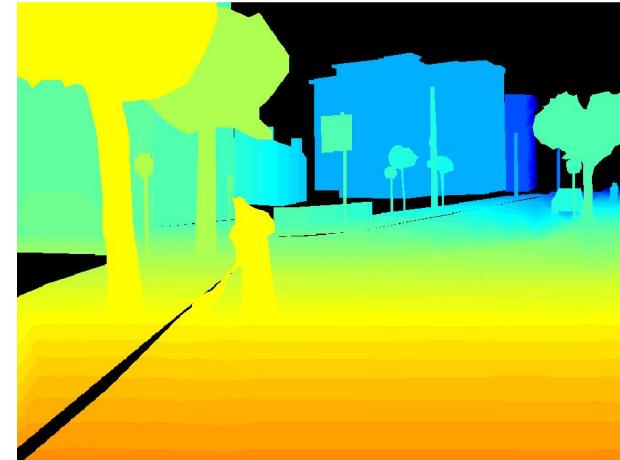


Source: Antonio Torralba slides

Street scene II

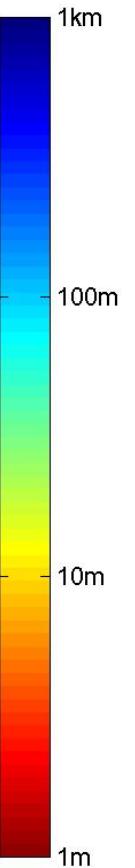
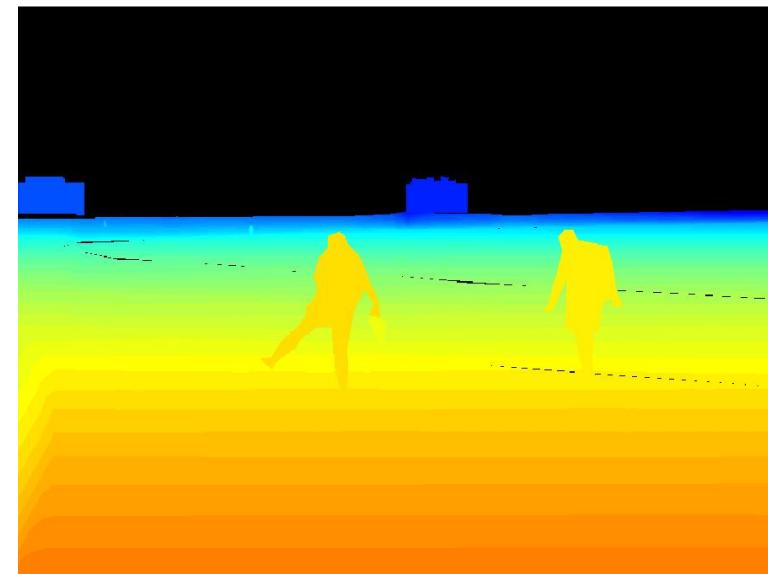


Street scene II

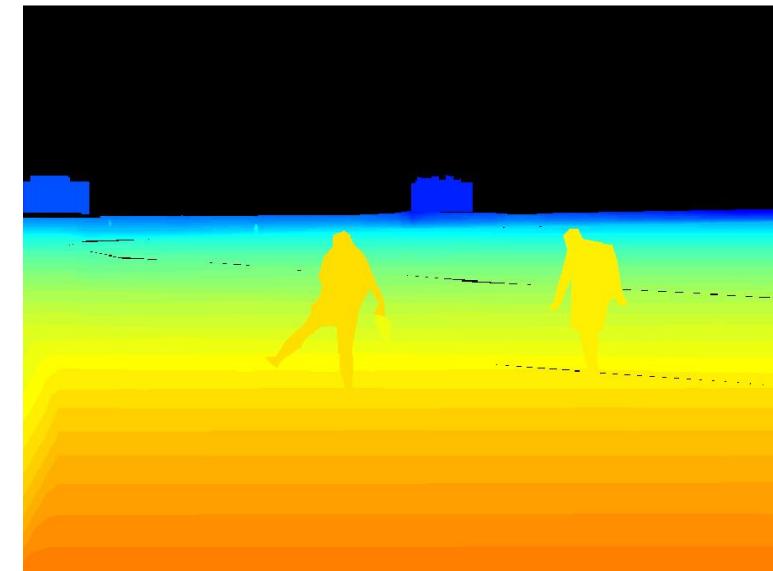


Source: Antonio Torralba slides

Beach scene

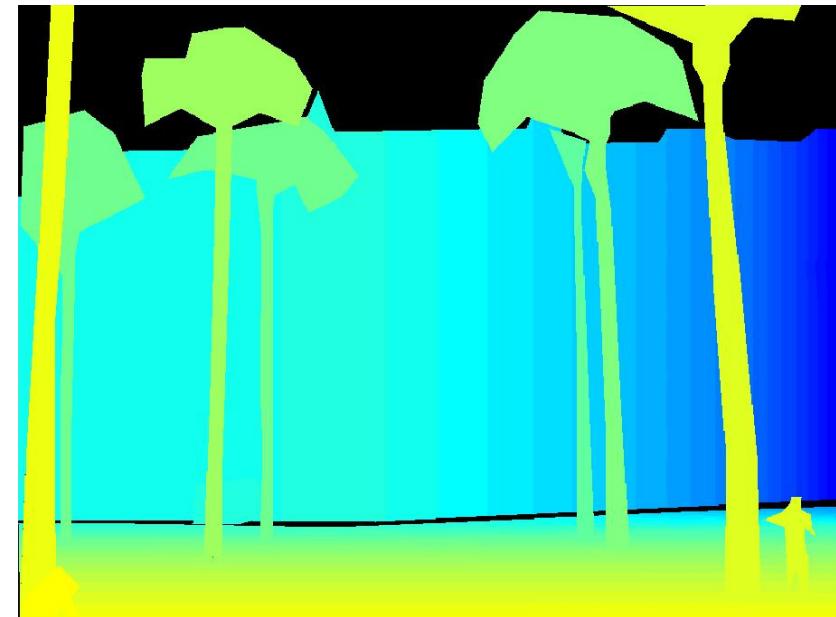


Beach scene



Source: Antonio Torralba slides

Building scene



Source: Antonio Torralba slides

Example: Multiview stereo



Summary

- Introduction
- 3D shapes from 2D images
- Stereo vision
- Correspondence
- Epipolar geometry
- The fundamental matrix
- The essential matrix
- RANSAC
- 3D reconstruction